
Nearly Optimal Differentially Private ReLU Regression

Meng Ding²

Mingxi Lei²

Shaowei Wang³

Tianhang Zheng⁴

Di Wang^{*5}

Jinhui Xu^{*1}

¹School of Information Science and Technology, University of Science and Technology of China

²Department of Computer Science and Engineering, State University of New York at Buffalo

³Institute of Artificial Intelligence and Blockchain, Guangzhou University

⁴The State Key Laboratory of Blockchain and Data Security, Zhejiang University

⁵Division of CEMSE, King Abdullah University of Science and Technology

Abstract

In this paper, we investigate one of the most fundamental non-convex learning problems—ReLU regression—in the Differential Privacy (DP) model. Previous studies on private ReLU regression heavily rely on stringent assumptions, such as constant-bounded norms for feature vectors and labels. We relax these assumptions to a more standard setting, where data can be i.i.d. sampled from $O(1)$ -sub-Gaussian distributions. We first show that when $\varepsilon = \tilde{O}(\sqrt{\frac{1}{N}})$ and there is some public data, it is possible to achieve an upper bound of $\tilde{O}(\frac{d^2}{N^2\varepsilon^2})$ for the excess population risk in (ε, δ) -DP, where d is the dimension and N is the number of data samples. Moreover, we relax the requirement of ε and public data by proposing and analyzing a one-pass mini-batch Generalized Linear Model Perceptron algorithm (DP-MBGLMtron). Additionally, using the tracing attack argument technique, we demonstrate that the minimax rate of the estimation error for (ε, δ) -DP algorithms is lower bounded by $\Omega(\frac{d^2}{N^2\varepsilon^2})$. This shows that DP-MBGLMtron achieves the optimal utility bound up to logarithmic factors. Experiments further support our theoretical results.

1 INTRODUCTION

Privacy preservation has become a critical consideration, posing a significant challenge for machine learning models that process sensitive data. To address this issue, Differential Privacy (DP) [Dwork et al., 2006] has emerged as a widely used approach, providing verifiable protection against identification and resistance to any auxiliary information that attackers might have.

Stochastic Optimization (SO) and its empirical counterpart, Empirical Risk Minimization (ERM), represent some of the most fundamental challenges in machine learning and statistics, which are especially susceptible to privacy leaks when involved with sensitive data. Therefore, significant efforts have been made to develop differentially private algorithms tailored to these challenges, specifically referred to as DP-SO and DP-ERM. Although there is an extensive body of research on DP-SO and DP-ERM [Bassily et al., 2014, Wang et al., 2017, 2023b, Feldman et al., 2020, Song et al., 2020, Su and Wang, 2021, Asi et al., 2021, Bassily et al., 2021b, Kulkarni et al., 2021, Hu et al., 2022, Zhang et al., 2025, Su et al., 2024, 2023], the majority of existing studies primarily focus on convex loss functions. This focus inadvertently neglects the crucial role of nonconvex optimization, which is essential for the development of advanced machine learning models. Recent progress has introduced algorithms for DP nonconvex optimization [Zhang et al., 2017, Wang et al., 2017, 2019, Wang and Xu, 2019a, Zhang et al., 2021, Bassily et al., 2021a, Wang et al., 2023c, Wang and Xu, 2024]. However, unlike the convex loss function, DP-SO with non-convex loss is still far from well-understood due to its intrinsic difficulties (see Section 2 for details).

ReLU regression, a fundamental non-convex model, is widely recognized for its effectiveness in deep learning applications and serves as a foundational step toward understanding multi-layer neural networks [Du et al., 2018]. Despite the extensive studies in the non-private setting that have been conducted, the theoretical exploration of ReLU regression in the DP model remains relatively limited. Particularly, in DP ReLU regression, we have an N -size dataset $D = \{(\mathbf{x}_i, y_i)\}_{i=0}^{N-1}$, where each data point consisting of a feature vector $\mathbf{x}_i \in \mathcal{X} \subseteq \mathbb{R}^d$ and a response variable $y_i \in \mathcal{Y}$ is i.i.d. sampled from a ReLU regression model. Specifically, each pair of (\mathbf{x}_i, y_i) is a realization of the ReLU regression model

$$y = \text{ReLU}(\mathbf{x}^\top \mathbf{w}_*) + z, \quad (1)$$

where $\text{ReLU}(\cdot) := \max\{\cdot, 0\}$; z is a zero mean randomized noise; $\mathbf{w}_* \in \mathbb{R}^d$ is the optimal model parameter.

^{*}Correspondence to: Di Wang <di.wang@kaust.edu.sa>, Jinhui Xu <jhxu@ustc.edu.cn>

The objective of the problem is to develop a DP model \mathbf{w}_{priv} that minimizes the excess population risk, defined as $\mathcal{L}(\mathbf{w}_{\text{priv}}) - \mathcal{L}(\mathbf{w}_*)$, where the risk function $\mathcal{L}(\mathbf{w})$ is given by:

$$\mathcal{L}(\mathbf{w}) = \frac{1}{2} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}}[(\text{ReLU}(\mathbf{x}^\top \mathbf{w}) - y)^2]. \quad (2)$$

Recently, Shen et al. [2023] explored DP ReLU regression in both well-specified and misspecified settings, yet the problem remains largely unexplored, with numerous challenges yet to be addressed. Specifically, their methods rely on stringent assumptions, including bounded norms for feature vectors and labels, with $\|\mathbf{x}\|_2 \leq O(1)$ and $\|y\| \leq O(1)$ —assumptions that do not hold even for typical Gaussian distributions. Even when $\|\mathbf{x}\|_2 \leq O(\sqrt{d})$ such as Bernoulli or uniform distributions, the bound in Shen et al. [2023] is only sub-optimal (see Remark 4 and Theorem 7 for details). Moreover, their proposed differentially private projected gradient descent (DP-PGD) requires at least $O(N^2)$ gradient computations, rendering it inefficient.

In this paper, we revisit the problem of DP ReLU regression and offer (nearly) optimal guarantees for excess population risk under more standard assumptions where the data can be i.i.d. sampled from $O(1)$ -sub-Gaussian distributions. Our contributions can be summarized as follows:

1) We provide the analysis on the Differentially Private Generalized Linear Model Perceptron algorithm (DP-GLMtron), which utilizes a one-pass training strategy where data points are permuted and sampled without replacement. To make the gradient norm bounded, instead of using a fixed clipping threshold, we incorporate adaptive clipping by estimating from additional public data points. This allows the noise to be set adaptively based on the excess error in each iteration. We demonstrate that our (ε, δ) -DP method can achieve an excess population risk upper bound of $\tilde{O}(\frac{d^2}{N^2 \varepsilon^2})$.

2) Key concerns with the analysis of DP-GLMtron include that its upper bound only holds with a small privacy budget $\varepsilon = O(\sqrt{\frac{\log(N/\delta)}{N}})$ and its reliance on additional public data for the adaptive clipping mechanism. To address these limitations, we modify DP-GLMtron to introduce a new method—DP-MBGLMtron (DP-Mini-Batch Generalized Linear Model Perceptron)—which divides the data into mini-batches and performs one pass of the mini-batch GLMtron. We show that DP-MBGLMtron can achieve the same excess population risk upper bound as DP-GLMtron, even with larger privacy budgets and without available public data.

3) To illustrate the tightness of our analysis, we derive a lower bound of the estimation error for any (ε, δ) -DP algorithms. Specifically, our analysis uses a tracing attack argument, illustrating that estimators with overly precise estimates would compromise privacy guarantees. According

to this property, we can establish that any such algorithm must incur an excess population risk of $\Omega(\frac{d^2}{N^2 \varepsilon^2})$, indicating that the upper bound is optimal up to logarithmic factors.

2 RELATED WORK

Private Convex Optimization. Differentially private convex optimization has been extensively studied over the past decade Chaudhuri et al. [2011], Jain et al. [2012], Kifer et al. [2012], Bassily et al. [2014], Jain and Thakurta [2014], Wang et al. [2017], Feldman et al. [2020]. Existing approaches in this field can broadly be categorized into three main categories: output perturbation, objective perturbation, and gradient perturbation. Output perturbation ensures differential privacy by adding calibrated noise to the final model parameters Dwork et al. [2006], Chaudhuri et al. [2011], Kifer et al. [2012], Zhang et al. [2017], Wu et al. [2017]; Objective perturbation modifies the optimization objective itself by injecting noise into the loss function before solving the problem, thereby inherently privatizing the optimization process Chaudhuri et al. [2011], Kifer et al. [2012], Talwar et al. [2014], Iyengar et al. [2019]; Gradient perturbation privatizes iterative optimization algorithms (e.g., stochastic gradient descent) by perturbing the gradient updates at each iteration Bassily et al. [2014], Wang et al. [2017], Jayaraman et al. [2018], Wang and Xu [2019a], Bassily et al. [2019]. All of these approaches have been demonstrated to achieve the asymptotically optimal bound $\tilde{O}(\frac{\sqrt{d}}{\varepsilon N})$ for smooth convex loss.

Private Nonconvex Optimization. In the domain of DP-SO and DP-ERM with convex loss functions, excess population risk has traditionally been the main metric for utility evaluation. However, in non-convex settings, utility assessment methods generally fall into three categories: first-order stationarity-based, second-order stationarity-based, and direct measurement of excess population risk. First-order stationarity-based methods [Wang and Xu, 2019a, Zhou et al., 2020, Song et al., 2021, Bassily et al., 2021a, Zhang et al., 2021, Xiao et al., 2023, Wang et al., 2023a, Tao et al., 2025] evaluate utility by analyzing the ℓ_2 -norm of the gradient of the population risk function. While widely adopted, these methods face notable challenges. For example, Agarwal et al. [2017] showed that as the sample size increases indefinitely, the gradient norm approaches zero. However, a vanishing gradient does not necessarily indicate that a differentially private estimator converges to, or is near, a local minimum. Second-order stationarity-based methods [Wang and Xu, 2019a, 2021] assess both the gradient norm and the minimal eigenvalue of the Hessian matrix of the population risk function. These approaches work well in specific settings where any second-order stationary point is a local minimum, and all local minima are global minima, such as in problems like matrix completion and dictionary learning. The third category directly uses excess population

risk to evaluate utility [Shen et al., 2023, Wang and Xu, 2019a], which aligns with the focus of our work.

One of the concurrent works, Ding et al., addresses the same private ReLU regression problem with similar assumptions, but our work differs significantly in several key aspects, including threshold estimation, privacy amplification techniques, theoretical bounds, and data assumptions. Specifically, Ding et al. uses a threshold estimation method based on Liu et al. [2023b] and a tree aggregation mechanism for privacy amplification, whereas we leverage statistical properties and minibatch sampling. Additionally, the theoretical results in Ding et al. include an upper bound with a term Γ dependent on unknown intermediate parameters \mathbf{w}_t , while our results depend only on the problem parameters d, n , and ε , making them more natural. Furthermore, the lower bound in Ding et al. is algorithm-specific and relies on intermediate models, whereas our lower bound is general, depending solely on d, n , and ε , and achieves nearly optimal rates. Finally, Ding et al. assumes the eigenvalue decomposition of the data covariance matrix is well-defined, which our approach does not require.

3 PRELIMINARIES

Notations: We use boldface lower letters such as \mathbf{x}, \mathbf{w} for vectors and boldface capital letters (e.g., \mathbf{A}, \mathbf{H}) for matrices. Let $\|\mathbf{A}\|_2$ denote the spectral norm of \mathbf{A} . For two matrices \mathbf{A} and \mathbf{B} of appropriate dimension, their inner product is defined as $\langle \mathbf{A}, \mathbf{B} \rangle := \text{tr}(\mathbf{A}^\top \mathbf{B})$. For a positive semi-definite (PSD) matrix \mathbf{A} and a vector \mathbf{v} of appropriate dimension, we write $\|\mathbf{v}\|_{\mathbf{A}}^2 := \mathbf{v}^\top \mathbf{A} \mathbf{v}$. The outer product is denoted by \otimes .

In this paper, we will employ the definition of classical DP Dwork et al. [2006] for privacy guarantees.

Definition 1 (Differential Privacy Dwork et al. [2006]). *A randomized algorithm \mathcal{A} is considered (ε, δ) -differentially private (abbreviated as (ε, δ) -DP) if, for any two datasets D and D' that differ by a single element, and for any event S in the output space of \mathcal{A} , the following condition holds: $\mathbb{P}[\mathcal{A}(D) \in S] \leq e^\varepsilon \cdot \mathbb{P}[\mathcal{A}(D') \in S] + \delta$*

In the following, we will introduce some definitions related to the model. We first consider the ReLU regression model to satisfy the following condition, which is commonly referred to in the literature as the "noisy teacher" setting Frei et al. [2020] or the well-structured noise model [Goel and Klivans, 2019], has been extensively studied in prior research [Zou et al., 2021, Varshney et al., 2022, Shen et al., 2023].

Definition 2 (Well-specified Condition). *Assume that there exists a parameter $\mathbf{w}_* \in \mathbb{R}^d$ such that $\mathbb{E}[y \mid \mathbf{x}] = \text{ReLU}(\mathbf{x}^\top \mathbf{w}_*)$, and the variance of the model noise can be denoted by $\sigma^2 := \mathbb{E}[(y - \text{ReLU}(\mathbf{x}^\top \mathbf{w}_*))^2]$.*

Moreover, we give some assumptions on the data to ensure the analysis of algorithms.

Assumption 1 (Data Covariance). *Define $\mathbf{H} := \mathbb{E}[\mathbf{x}\mathbf{x}^\top]$ as the expected data covariance matrix and assume that each entry and the trace of \mathbf{H} are finite.*

Assumption 2 (Fourth Moment Conditions). *Assume that the fourth moment of \mathbf{x} is finite and there exists a constant $\alpha > 0$ such that for any Positive Semi-Definite (PSD) matrix \mathbf{A} , the following holds:*

$$\mathbb{E}[\mathbf{x}\mathbf{x}^\top \mathbf{A} \mathbf{x}\mathbf{x}^\top] \preceq \alpha \cdot \text{tr}(\mathbf{H}\mathbf{A}) \cdot \mathbf{H}.$$

Remark 1. For normal Gaussian distribution, it can be verified that Assumption 2 holds with $\alpha = 3$ Zou et al. [2021]. Moreover, when the data follows a sub-Gaussian distribution—more precisely, when $\mathbf{x} = \mathbf{H}^{-\frac{1}{2}} \mathbf{z}$, where \mathbf{z} is a sub-Gaussian random vector with variance σ_z^2 —Assumption 2 remains valid with $\alpha = 16\sigma_z^2$ [Wang and Xu, 2019b, Varshney et al., 2022, Zhu et al., 2023, 2024, Liu et al., 2023b, Ding et al., 2024].

Definition 3 ((\mathbf{H}, C_2, a, b) -Tail). *A random vector \mathbf{x} satisfies $(\mathbf{H}, C_2, a, b_{\mathbf{x}})$ -Tail if the following holds:*

- $\exists a > 0$ s.t. with probability $\geq 1 - b_{\mathbf{x}}$,

$$\|\mathbf{x}\|_2^2 \leq \mathbb{E}[\|\mathbf{x}\|_2^2] \cdot \log^{2a}(1/b_{\mathbf{x}}), \quad (3)$$

- We have,

$$\max_{\mathbf{v}, \|\mathbf{v}\|=1} \mathbb{E}[\exp((\frac{|\langle \mathbf{x}, \mathbf{v} \rangle|^2}{C_2^2 \|\mathbf{H}\|_2})^{1/2a})] \leq 1,$$

That is, for any fixed \mathbf{v} , with probability $\geq 1 - b_{\mathbf{x}}$:

$$(\langle \mathbf{x}, \mathbf{v} \rangle)^2 \leq C_2^2 \|\mathbf{H}\|_2 \|\mathbf{v}\|^2 \log^{2a}(1/b_{\mathbf{x}}).$$

Definition 3 has been extensively employed in recent studies on differential privacy analysis for sub-Gaussian data, as seen in [Varshney et al., 2022, Liu et al., 2022, 2023a]. In this work, we assume that each sample \mathbf{x} satisfies the $(\mathbf{H}, C_2, a, b_{\mathbf{x}})$ -Tail condition, while the inherent noise z satisfies the $(\sigma^2, C_2, a, b_{\mathbf{x}})$ -Tail condition. Furthermore, based on Assumption 2, it directly follows that $\|\mathbf{x}\|_2^2 \leq \alpha \text{tr}(\mathbf{H}) \cdot \log^{2a}(1/b_{\mathbf{x}})$ with probability at least $1 - b_{\mathbf{x}}$, as shown in Equation (3).

Assumption 3 (Symmetricity conditions). *Assume that for every $\mathbf{u}, \mathbf{v} \in \mathbb{R}^d$, it holds that:*

$$\begin{aligned} & \mathbb{E}[\mathbf{x}\mathbf{x}^\top \cdot \mathbb{1}[\mathbf{x}^\top \mathbf{u} > 0, \mathbf{x}^\top \mathbf{v} > 0]] \\ &= \mathbb{E}[\mathbf{x}\mathbf{x}^\top \cdot \mathbb{1}[\mathbf{x}^\top \mathbf{u} < 0, \mathbf{x}^\top \mathbf{v} < 0]], \\ & \mathbb{E}[(\mathbf{x}^\top \mathbf{v})^2 \mathbf{x}\mathbf{x}^\top \cdot \mathbb{1}[\mathbf{x}^\top \mathbf{u} > 0, \mathbf{x}^\top \mathbf{v} > 0]] \\ &= \mathbb{E}[(\mathbf{x}^\top \mathbf{v})^2 \mathbf{x}\mathbf{x}^\top \cdot \mathbb{1}[\mathbf{x}^\top \mathbf{u} < 0, \mathbf{x}^\top \mathbf{v} < 0]]. \end{aligned}$$

Remark 2. Here, we impose the assumptions that both the second and fourth moments of \mathbf{x} exhibit symmetry. Assumption 3 is satisfied when \mathbf{x} and $-\mathbf{x}$ follow the same distribution. This condition naturally holds for symmetric sub-Gaussian distributions, including symmetric Bernoulli and Gaussian distributions.

4 DP-GLMTRON ALGORITHM

Before presenting our analysis on DP-GLMtron, we first recall the proposed DP-PGD algorithm in [Shen et al., 2023]. The central principle of DP-PGD in ensuring privacy protection involves adding noise to the gradient and executing a projection operation post-model update. This process ensures that the model parameter \mathbf{w} remains bounded, thereby keeping the gradient within manageable limits as well. However, this method leaves several unresolved issues. Primarily, their algorithm assumes that the data are bounded with $\|\mathbf{x}\|_2 \leq 1$, which enables the control of the gradient $\nabla \mathcal{L}(\mathbf{w}) = (\text{ReLU}(\mathbf{x}^\top \mathbf{w}) - y) \mathbf{x} \cdot \mathbb{1}[\mathbf{x}^\top \mathbf{w} > 0]$ via the model \mathbf{w} its subsequent projection. If the data exhibit $O(1)$ -sub-Gaussian properties, then we can see $\|\nabla \mathcal{L}(\mathbf{w})\| \leq O(d)$ (with high probability), which means the Gaussian noise added in each iteration has a scale of $\Omega(d^2)$, making a large estimation error (see Remark 4 for a detailed comparison). Additionally, it is noticed that at each iteration, DP-PGD requires computing a full gradient. This process is highly costly and inefficient, particularly in settings involving large datasets or high-dimensional data.

To address the above-mentioned challenges, we consider the DP-GLMtron method built upon the Generalized Linear Model Perceptron (GLMtron) algorithm of [Kakade et al., 2011] with a one-pass strategy. The fundamental distinction between SGD and GLMtron lies in their respective update rules. Specifically, it takes the following rules:

$$\text{SGD: } \mathbf{w}_t = \mathbf{w}_{t-1} - \eta \cdot \mathbf{l}_t$$

$$\text{where } \mathbf{l}_t = (\text{ReLU}(\mathbf{x}_t^\top \mathbf{w}_{t-1}) - y_t) \mathbf{x}_t \cdot \mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0]$$

$$\text{GLMtron: } \mathbf{w}_t = \mathbf{w}_{t-1} - \eta \cdot (\text{ReLU}(\mathbf{x}_t^\top \mathbf{w}_{t-1}) - y_t) \mathbf{x}_t.$$

The algorithm begins from an initial point \mathbf{w}_0 and iterates from $t = 0$ to $t = N - 1$ with a step size η . In contrast to the typical update rule of SGD, GLMtron diverges by modifying the derivative of the ReLU function in its update mechanism. The exclusion of this derivative in GLMtron’s framework not only simplifies the computational process but also enhances efficiency. Furthermore, [Kakade et al., 2011] demonstrates that this specific omission significantly contributes to GLMtron’s ability to efficiently identify a predictor that closely approximates the optimal solution. Building upon these foundations, we now present the detailed implementation of the proposed DP-GLMtron. The process starts with a random permutation of the dataset to amplify privacy via shuffling [Feldman et al., 2022]. In contrast to Shen et al. [2023], our method adopts the one-pass DP strategy without data replacement, ensuring that the time complexity is linear in N and each iterate of model \mathbf{w}_t is independent of data \mathbf{x}_t . See Algorithm 1 for details.

A critical step in our approach involves determining the clipping threshold prior to the iterative updates for \mathbf{w}_t . An excessively low clipping threshold can result in the loss of important gradient information, leading to high bias McMa-

Algorithm 1 DP-GLMtron

```

1: Input: Samples:  $\{(\mathbf{x}_i, \mathbf{y}_i)\}_{i=0}^{N-1}$ , Clipping Norm:  $\xi$ , DP
   Noise Multiplier:  $f$ , Learning Rate:  $\eta$ , Public Data
    $D' = \{(\mathbf{x}'_i, \mathbf{y}'_i)\}_{i=1}^m$ , Parameters  $\Upsilon, \Delta$ 
2: Randomly permute  $\{(\mathbf{x}_i, \mathbf{y}_i)\}_{i=0}^{N-1}$ 
3: Initialize  $\mathbf{w}_0 \leftarrow 0$ 
4: for  $t = 0, \dots, N - 1$  do
5:    $s_t \leftarrow \text{DP-Threshold}(\{(\mathbf{x}'_i, \mathbf{y}'_i)\}_{i=1}^m, \mathbf{w}_t, \Upsilon, \Delta)$ 
6:   Sample  $\mathbf{g}_t \sim \mathcal{N}(0, \mathbf{I}_{d \times d})$ 
7:    $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - \eta(\text{clip}_{s_t}(\mathbf{x}_t^\top (\text{ReLU}(\mathbf{x}_t^\top \mathbf{w}_t) - \mathbf{y}_t)) +$ 
      $2f s_t \mathbf{g}_t)$ 
8: end for
9: return  $\mathbf{w} \leftarrow \frac{1}{N} \sum_{t=0}^{N-1} \mathbf{w}_t$ 

```

han et al. [2017], Amin et al. [2019]. Therefore, we employ an adaptive clipping by estimating additional public data points Andrew et al. [2021], Varshney et al. [2022]. Specifically, Algorithm 2 sets the initial threshold s_0 , which seems to be a threshold that will be iteratively refined to find the approximate maximum. The loop runs for $\lceil \log_2(\Upsilon/\Delta) \rceil$ iterations, covering a range of possible maximum values scaled by the parameter Υ and the discretization width Δ . In each iteration, the Algorithm 2 counts the number of samples for which the value $|\text{ReLU}(\mathbf{x}_j^\top \mathbf{w}) - \mathbf{y}_j|$ is less than or equal to the current threshold s_t . If the private count is less than the sample size of public data m , the threshold is updated for the next iteration to double of its current value. If the count meets m , the Algorithm 2 exits the loop. When determining the clipping threshold, the model updates via the classical Gaussian mechanism. Finally, Algorithm 1 returns to the average of the iterates.

Algorithm 2 DP-Threshold

```

1: Input: Estimating Samples:  $\{(\mathbf{x}'_i, \mathbf{y}'_i)\}_{i=1}^m$ , Current
   Model:  $w$ , DP Noise Multiplier:  $f$ , Domain Size:  $\Upsilon$ ,
   Discretization Width:  $\Delta$ 
2:  $s_0 \leftarrow \Delta$ 
3: for  $i \in \{0, \dots, \lceil \log_2(\Upsilon/\Delta) \rceil\}$  do
4:    $u \leftarrow |\{|\text{ReLU}(\mathbf{x}_j^\top \mathbf{w}) - \mathbf{y}_j| \leq s_i : j \in \{0, \dots, m\}\}|$ 
5:   if Estimating samples are public then
6:      $u_{\text{priv}} \leftarrow u$ 
7:   else
8:      $u_{\text{priv}} \leftarrow u + \mathcal{N}(0, \lceil \log_2(\Upsilon/\Delta) \rceil f^2)$ 
9:   end if
10:  if  $u_{\text{priv}} < m$  then
11:     $s_{i+1} \leftarrow 2 * s_i$ 
12:  else
13:    break
14:  end if
15: end for
16: return  $s_{\text{priv}} \leftarrow s_i$ 

```

Theorem 1 (Privacy Guarantee). *DP-GLMtron satisfies (ε, δ) -DP with a noise multiplier set to $f = \Omega(\frac{\log(N/\delta)}{\varepsilon\sqrt{N}})$ if $\varepsilon = O(\sqrt{\frac{\log(N/\delta)}{N}})$ and $0 < \delta < 1$.*

Remark 3. Note that the privacy budget is limited to $\varepsilon = O(\sqrt{\frac{\log(N/\delta)}{N}})$ because of privacy amplification via shuffling in Feldman et al. [2022]. If there is no shuffling, plainly using the Gaussian mechanism will make $f = \Omega(\frac{\log(N/\delta)}{\varepsilon})$. Thus, privacy amplification can improve a factor of $\tilde{O}(\sqrt{N})$. However, this highlights a key limitation in DP-GLMtron: as the dataset size N increases, the algorithm is constrained by a smaller privacy budget ε .

Theorem 2 (Utility Guarantee). *Let $D = \{(\mathbf{x}_i, y_i)\}_{i=0}^{N-1}$ be sampled i.i.d. with $\mathbf{x}_i \sim \mathcal{D}$ satisfying $(\mathbf{H}, C_2, a, b_{\mathbf{x}})$ -Tail, and the distribution of the inherent noise z satisfies $(\sigma^2, C_2, a, b_{\mathbf{x}})$ -Tail with $b_{\mathbf{x}} = \frac{1}{\text{Poly}(N)}$. Let κ be the condition number of the covariance matrix \mathbf{H} and denote $R_x^2 = \alpha \text{tr}(\mathbf{H}) \cdot \log^{2a}(1/b_{\mathbf{x}})$.*

Initialize parameters in DP-GLMtron as follows: stepsize $\eta = \min\{\frac{1}{2R_x^2}, \frac{c_1}{\log^{4a} N} \cdot \frac{1}{C_2^2 R_x^2 \kappa^2} \cdot \frac{1}{d^2}\}$, where $c_1, c_2 > 0$ are global constants, noise multiplier $f = \Omega(\frac{\log(N/\delta)}{\varepsilon\sqrt{N}})$, domain size $\Upsilon = C_2 R_x (\|\mathbf{w}^\|_{\mathbf{H}} + \sigma) \log^{2a} N$, granularity $\Delta = \frac{\|\mathbf{w}^*\|_{\mathbf{H}} + \sigma}{\text{Poly}(N)}$, public datasize $m = \Omega(\frac{\log(N/\delta)\sqrt{\log(N \log N)}}{\varepsilon\sqrt{N}})$. Then, the output $\bar{\mathbf{w}}$ of DP-GLMtron achieves the following excess risk w.p. $\geq 1 - 1/\text{Poly}(N)$ over randomness in data and algorithm:*

$$\begin{aligned} \mathcal{L}(\bar{\mathbf{w}}) - \mathcal{L}(\mathbf{w}^*) &\lesssim \frac{\|\mathbf{w}_*\|_{\mathbf{H}}^2}{\text{Poly}(N)} + \frac{\sigma^2 d}{N} \\ &+ \frac{d^2 \log N \log(1/\delta)}{N^2 \varepsilon^2} \cdot C_2^2 \kappa^2 (\sigma^2 + \|\mathbf{w}_*\|_{\mathbf{H}}^2). \end{aligned}$$

Remark 4. Theorem 2 provides a utility guarantee for the DP-GLMtron algorithm, balancing privacy and performance. The excess risk is composed of three key components: The first component, dependent on $\|\mathbf{w}_*\|_{\mathbf{H}}^2$, diminishes polynomially in N . The second component corresponds to the inherent model noise, achieving the optimal rate (up to a constant factor) for non-private ReLU regression as established by Wu et al. [2023]. The third component is of the order $\tilde{O}(\frac{d^2}{N^2 \varepsilon^2})$. For $N = \Omega(d)$, the bound implies nearly optimal sample complexity, further supported by the lower bound derived in Section 6, disregarding constant factors.

Compared to Shen et al. [2023], our analysis here relax the data assumption $\|\mathbf{x}\|_2 \leq 1$. If we assume that $\|\mathbf{x}\|_2 \leq O(\sqrt{d})$, via the same analysis as in Shen et al. [2023], we can show the utility bound will be $O(\min\{\frac{d\sqrt{d}}{N\varepsilon}, (\frac{d}{N\varepsilon})^{\frac{2}{3}}\})$, which is worse than the one in Theorem 2.

5 ADVANCED DP-MINI-BATCH-GLMTRON

A key concern with the DP-GLMtron algorithm is its limited practicality where the privacy budget ε is small, potentially restricting its utility in real-world applications (see Theorem 1 for more details). Furthermore, Algorithm 1 may require additional public data to estimate the threshold. To overcome these challenges, we introduce DP-Mini-batch-GLMtron (Algorithm 3) in this section.

Specifically, the algorithm first operates by randomly partitioning the training samples $\{(\mathbf{x}_i, y_i)\}_{i=0}^{N-1}$, and setting the number of iterations $T = N/(b + m)$, where b and m are batch sizes and estimating sample size for determining the threshold. It is worth noting that, in this approach, a separate public dataset is not required to estimate the clipping threshold. Instead, we divide each batch of data and use a portion of it as the estimation data for the threshold. Therefore, in each iteration, the algorithm processes a mini-batch of data with size m and computes the DP-Threshold γ_t using estimating samples $\{(\mathbf{x}'_i, y'_i)\}_{i=1}^m$, and updates the clipping parameter s_t . In contrast to DP-GLMtron, we need to protect the counting numbers during the estimation process as we are using private data. Noise g_t is sampled from a Gaussian distribution and added to the gradient step for privacy preservation. The model weights are updated using step 9, where \mathbf{l}_{t+1} is the averaged clipped gradient. After iterating T times, the final weight estimate $\bar{\mathbf{w}}$ is returned as the average of all weight updates.

Theorem 3. *Algorithm DP-mini-batch-GLMtron with noise multiplier $f \geq \frac{2\sqrt{\log(1/\delta)+\varepsilon}}{\varepsilon}$ satisfies (ε, δ) -DP. Furthermore, if $\varepsilon \leq \log(1/\delta)$, then $f \geq \frac{\sqrt{8\log(1/\delta)}}{\varepsilon}$ suffices to ensure (ε, δ) -DP.*

Theorem 3 addresses the limitations of the DP-GLMtron algorithm, particularly requiring a small privacy budget ε , which can severely limit its utility in practical scenarios. By processing a subset of data in each iteration, the algorithm effectively reduces the sensitivity of the overall computation. This reduction allows for less noise to be added while maintaining the same level of privacy, thus improving the utility of the model.

Theorem 4. *Let $D = \{(\mathbf{x}_i, y_i)\}_{i=0}^{N-1}$ be sampled i.i.d. with $\mathbf{x}_i \sim \mathcal{D}$ satisfying $(\mathbf{H}, C_2, a, b_{\mathbf{x}})$ -Tail, and the distribution of the inherent noise z satisfies $(\sigma^2, C_2, a, b_{\mathbf{x}})$ -Tail with $b_{\mathbf{x}} = \frac{1}{\text{Poly}(N)}$.*

Initialize parameters in DP-MBGLMtron as follows: batch size $b = \frac{N}{T} - m$, estimating sample size $m = \frac{b}{10}$, appropriate stepsize $\eta = O(\frac{1}{R_x^2})$, number of iterations $T = O(\kappa \log(N))$, domain size $\Upsilon = C_2 R_x (\|\mathbf{w}^\|_{\mathbf{H}} + \sigma) \log^{2a} N$, granularity $\Delta = \frac{\|\mathbf{w}^*\|_{\mathbf{H}} + \sigma}{\text{Poly}(N)}$ and noise multiplier $f = \frac{\sqrt{8\log(1/\delta)}}{\varepsilon}$. Then, the output $\bar{\mathbf{w}}$ achieves the*

Algorithm 3 DP-MBGLMtron

```

1: Input: Training Samples:  $\{(x_i, y_i)\}_{i=0}^{N-1}$ , Learning
   Rate:  $\eta$ , DP Noise Multiplier:  $f$ , Expected  $x$  Norm:
    $\sqrt{\alpha \text{tr}(\mathbf{H})}$ , Parameters  $\Upsilon, \Delta$ 
2: Initialize  $\mathbf{w}_0 \leftarrow \mathbf{0}$  and  $s_0 \leftarrow \Delta$ 
3: Set  $T \leftarrow N/(b+m)$ 
4: for  $t = 0 \dots T-1$  do
5:   Set  $\tau(t) \leftarrow (b+m)t$ 
6:    $\gamma_t \leftarrow \text{DP-Threshold}(D_t, \mathbf{w}_t, f, \Upsilon, \Delta)$ , where
      $D_t = \{(\mathbf{x}_{\tau(t)+j}, y_{\tau(t)+j})\}_{j=0}^{m-1}$ 
7:    $s_t = \sqrt{2\alpha \text{tr}(\mathbf{H})} C_2 \log^{2a} N \cdot \gamma_t$  and add  $s_t$  to the
     list  $\mathbf{s}$ 
8:   Sample  $\mathbf{g}_t \sim \mathcal{N}(\mathbf{0}, \mathbf{I}_{d \times d})$ 
9:    $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - \eta \mathbf{l}_{t+1} - \frac{2f s_t \eta}{b} \mathbf{g}_t$ , where  $\mathbf{l}_{t+1} :=$ 
      $\frac{1}{b} \sum_{i=0}^{b-1} \text{clip}_{s_t}(\mathbf{x}_{\tau(t)+m+i}(\text{ReLU}(\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_t) -$ 
      $y_{\tau(t)+m+i}))$ 
10: end for
11: return  $\mathbf{w} := \frac{1}{T} \sum_{t=0}^{T-1} \mathbf{w}_t$ 

```

following excess risk with probability $\geq 1 - 1/\text{Poly}(N)$ over the randomness in data and algorithm:

$$\begin{aligned} \mathcal{L}(\bar{\mathbf{w}}) - \mathcal{L}(\mathbf{w}^*) &\lesssim \frac{\|\mathbf{w}_*\|_{\mathbf{H}}^2}{\text{Poly}(N)} + \frac{\sigma^2 d}{N} \\ &\quad + \frac{d^2 \log N \log(1/\delta)}{N^2 \varepsilon^2} \cdot C_2^2 \kappa^2 (\sigma^2 + \|\mathbf{w}_*\|_{\mathbf{H}}^2). \end{aligned}$$

To prove the utility, we have the following utility for the DP-Threshold algorithm.

Theorem 5 (DP-Threshold). *Suppose that DP-Threshold is applied to m estimated data with certain parameters $\{\mathbf{w}_t, \Upsilon, f, \Delta\}$, Algorithm 2 satisfies $(\varepsilon/2, \delta/2)$ -DP with $f \geq \frac{2\sqrt{\log(1/\delta)+\varepsilon}}{\varepsilon}$. Given $\Lambda = f\sqrt{2\log(\Upsilon/\Delta)\log(\log(\Upsilon/\Delta)/b_{\mathbf{x}})}$, then with probability at least $1 - b_{\mathbf{x}}$, Algorithm 2 outputs a private threshold s_{priv} such that*

- $|\{i \mid \text{ReLU}(\mathbf{x}_i^\top \mathbf{w}) - y_i \leq s_{\text{priv}} : i \in \{0, \dots, m\}\}| \geq m - \Lambda,$
- $|\{i \mid \text{ReLU}(\mathbf{x}_i^\top \mathbf{w}) - y_i \leq \max\{\frac{s_{\text{priv}}}{2}, \Delta\} : i \in \{0, \dots, m\}\}| < m - \Lambda.$

Remark 5. We provide further details regarding the threshold here. Suppose $\Lambda = \Omega(f \log N)$. With probability at least $1 - 1/\text{Poly}(N)$, at least $m - \Lambda$ data points satisfy the condition $|\text{ReLU}(\mathbf{x}_i^\top \mathbf{w}) - y_i| \leq s_{\text{priv}}$. According to Definition 3, and considering that \mathbf{w}_t is independent of $\mathbf{x}_{\tau(t)+j}$, we have the following with probability $\geq 1 - 1/\text{Poly}(N)$:

$$\|\mathbf{x}_{\tau(t)+j}(\text{ReLU}((\mathbf{x}_{\tau(t)+j})^\top \mathbf{w}_t) - y_{\tau(t)+j})\| \leq s_t,$$

by setting $s_t = O(R_x \gamma_t \log^a N)$ for all iterations. Moreover, recalling that Δ is the granularity of the search for the

approximate maximum threshold and Algorithm 2 will terminate once most of the samples fit under the current guess for γ_t , meaning γ_t will not exceed the current value plus the granularity Δ . That is, $\gamma_t \leq C_2 \log^a N (\sqrt{\kappa} \|\mathbf{w}_t - \mathbf{w}^*\|_{\mathbf{H}} + \sigma + \Delta)$ and $\Delta = \frac{\|\mathbf{w}^*\|_{\mathbf{H}} + \sigma}{\text{Poly}(N)}$. The choice of Δ reflects the granularity needed as the current weight approaches the optimal one. Therefore, with probability $\geq 1 - 1/\text{Poly}(N)$, both events hold: 1) the threshold is not required for any data point in its batch, and 2) the above condition on γ_t is satisfied in each iteration.

6 LOWER BOUND

In this section, we demonstrate that the minimax rate of the excess population risk for (ε, δ) -DP algorithms is lower bounded by $\Omega(\frac{d^2}{N^2 \varepsilon^2})$, indicating that the bound mentioned above is optimal up to logarithmic factors. To show this, we consider the following class of distributions for (\mathbf{x}, y) :

$$\begin{aligned} \mathcal{P}(\sigma, d, \mathcal{W}) &= \{P(\mathbf{x}, y) \mid \mathbf{w} \in \mathcal{W}, \mathbf{x} \sim \text{Uni}([-1, 1]^d), \\ f_{\mathbf{w}}(y|\mathbf{x}) &= \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y - \text{ReLU}(\mathbf{w}^\top \mathbf{x}))^2}{2\sigma^2}\right)\}, \end{aligned} \quad (4)$$

where $\mathcal{W} = \{\mathbf{w} \in \mathbb{R}^d \mid \|\mathbf{w}\|_2 \leq 1\}$, and $f_{\mathbf{w}}(y|\mathbf{x})$ is the density function of y given \mathbf{w} and \mathbf{x} . Thus, for any $(\mathbf{x}, y) \sim P \in \mathcal{P}(\sigma, d, \mathcal{W})$ we have $y = \text{ReLU}(\mathbf{w}^\top \mathbf{x}) + z$, where $z \sim \mathcal{N}(0, \sigma^2)$, and the covariate \mathbf{x} satisfies Assumption 2 with $\alpha, \beta = O(1)$ and Assumption 3. It also satisfies $(\mathbf{I}_d, C_2, a, b)$ -Tail with some $a, b = O(1)$.

Our lower bounds will be in the form of private minimax risk. Let \mathcal{P} be a class of distributions over a data universe \mathcal{X} . For each distribution $p \in \mathcal{P}$, there is a deterministic function $\mathbf{w}(p) \in \mathcal{W}$, where \mathcal{W} is the parameter space. Let $\rho : \mathcal{W} \times \mathcal{W} \rightarrow \mathbb{R}_+$ be a semi-metric function on the space \mathcal{W} and $\Phi : \mathbb{R}_+ \rightarrow \mathbb{R}_+$ be a non-decreasing function with $\Phi(0) = 0$.¹ We further assume that $D = \{X_i\}_{i=1}^n$ are n i.i.d observations drawn according to some distribution $p \in \mathcal{P}$, and $\hat{\mathbf{w}} : \mathcal{X}^N \rightarrow \mathcal{W}$ be some estimator. In the (ε, δ) -DP model, the estimator $\hat{\mathbf{w}}$ is obtained via some (ε, δ) -DP mechanism Q . The (ε, δ) -private minimax risk is defined as:

$$\mathcal{M}_n(\theta(\mathcal{P}), \Phi \circ \rho) := \inf_{Q \in \mathcal{Q}_{\varepsilon, \delta}} \sup_{p \in \mathcal{P}} \mathbb{E}_{p, Q}[\Phi(\rho(Q(D), \mathbf{w}(p)))],$$

where $\mathcal{Q}_{\varepsilon, \delta}$ is the set of all the (ε, δ) -DP mechanisms.

To prove the lower bound, we aim to use the tracing attack argument in [Cai et al., 2021]. Specifically, a tracing attacker attempts to construct an attack to detect the absence/presence of a sample \mathbf{x} in a target dataset D by looking at the (private) estimator $M(D)$ for the dataset. If one

¹In this paper, we assume that $\rho(\mathbf{w}, \mathbf{w}') = \|\mathbf{w} - \mathbf{w}'\|_{\Sigma_{\mathbf{x}}}$ and $\Phi(\mathbf{x}) = \mathbf{x}^2$ unless specified otherwise, where $\Sigma_{\mathbf{x}} = \mathbf{I}_d$ is the covariance matrix of x . Here, we do not omit $\Sigma_{\mathbf{x}}$ to make our results consistent with previous results.

can construct a tracing attack that is powerful, given an accurate estimator, an argument by contradiction leads to a lower bound: suppose a differentially private estimator computed from the target data set is sufficiently accurate, the tracing adversary will be able to determine whether a given sample belongs to the dataset or not, thereby contradicting with the differential privacy guarantee. The privacy guarantee and the tracing adversary together ensure that a differentially private estimator cannot be "too accurate". In detail, for a dataset D and a target sample (\mathbf{x}, y) , we consider the following tracing attack:

$$\mathcal{T}_{\mathbf{w}}((\mathbf{x}, y), M(D)) = \langle M(D) - \mathbf{w}, (y - \text{ReLU}(\mathbf{w}^\top \mathbf{x})) \mathbf{x} \cdot \mathbb{1}(\mathbf{w}^\top \mathbf{x} > 0) \rangle. \quad (5)$$

We will first show that if $(\mathbf{x}, y) \in D$, then the attack value is small; otherwise, it will be large.

Lemma 6. *Consider $D = (Y, X) = \{(\mathbf{x}_i, y_i)\}_{i=1}^N$ be i.i.d. sampled from $P \in \mathcal{P}(\sigma, d, \mathcal{W})$ with the underlying \mathbf{w} . For every (ε, δ) -DP algorithm M satisfying $\mathbb{E}_{Y, X | \mathbf{w}} \|M(D) - \mathbf{w}\|_2^2 = o(1)$ for all $\mathbf{w} \in \mathcal{W}$, then we have the following:*

1. *For each $i \in [n]$, denote D'_i as the dataset obtained by replacing (\mathbf{x}_i, y_i) in D with an independent copy, then we have*

$$\begin{cases} \mathbb{E} \mathcal{T}_{\mathbf{w}}((\mathbf{x}_i, y_i), M(D'_i)) = 0, \\ \mathbb{E} |\mathcal{T}_{\mathbf{w}}((\mathbf{x}_i, y_i), M(D'_i))| \leq \sigma \sqrt{\mathbb{E} \|M(D) - \mathbf{w}\|_{\Sigma_{\mathbf{x}}}^2}, \end{cases}$$

2. *There exists a prior distribution of π for \mathbf{w} supported on \mathcal{W} such that*

$$\sum_{i \in [n]} \mathbb{E}_{\pi} \mathbb{E}_{Y, X | \mathbf{w}} [\mathcal{T}_{\mathbf{w}}((\mathbf{x}_i, y_i), M(D))] \geq \Omega(\sigma^2 d).$$

Remark 6. Lemma 6 establishes a connection between the accuracy of a DP estimator and the potential for privacy breaches via tracing attacks. Specifically, when (x_i, y_i) is independent on D'_i , we can control the variance of $\mathcal{T}_{\mathbf{w}}((\mathbf{x}_i, y_i), M(D'_i))$, which is upper bounded by $\sigma \sqrt{\mathbb{E} \|M(D) - \mathbf{w}\|_{\Sigma_{\mathbf{x}}}^2}$. Moreover, if they are dependent, then from part 2 we can see there exists w such that $\mathcal{T}_{\mathbf{w}}((\mathbf{x}_i, y_i), M(D)) \geq \Omega(\frac{\sigma^2 d}{n})$. These results show that when $\|M(D) - \mathbf{w}\|_{\Sigma_{\mathbf{x}}}^2$ is small enough, then the attacker can distinguish D'_i and D , making DP failed. Specifically, we have the following result:

Theorem 7. *For $0 < \varepsilon < 1$ and $\delta \leq N^{-(1+u)}$ for some $u > 0$, we have*

$$\inf_{M \in \mathcal{Q}_{\varepsilon, \delta}} \sup_{p \in \mathcal{P}} \mathbb{E}_{D \sim p^N, M} [\mathcal{L}(M(D)) - \mathcal{L}(\mathbf{w})] \geq \frac{1}{4} \inf_{M \in \mathcal{Q}_{\varepsilon, \delta}} \sup_{p \in \mathcal{P}} \mathbb{E}_{D \sim p^N, M} [\|M(D) - \mathbf{w}\|_{\Sigma_{\mathbf{x}}}^2] \geq O(\frac{\sigma^2 d^2}{N^2 \varepsilon^2}).$$

Remark 7. Theorem 7 shows that under differential privacy, if the estimator M is too accurate, it may inadvertently leak information about the presence of specific data points. Therefore, the mechanism must maintain the error rate of $O(\frac{d^2}{N^2 \varepsilon^2})$, which aligns with our previous upper bound, thus confirming the tightness of our results.

7 EXPERIMENTS

In this section, we present experimental results to validate our theoretical findings. Due to space constraints, the detailed experimental setup and implementation details are provided in Appendix A.

Datasets and Models. We conducted experiments using three regression datasets: California Housing [Pace and Barry, 1997], Gas Turbine CO and NOx Emission DataSet [gas, 2019], and Wine Quality [Cortez et al., 2009]. The information of three datasets used in our experiments is summarized in Section 7. For each dataset, the data was randomly split into an 80% training set and a 20% test set. All numeric attributes were standardized to have a mean of zero and a standard deviation of one. The target variables were normalized by dividing them by the maximum absolute value of the target variable across the entire dataset. The model used for the experiments was ReLU regression, and evaluations were performed under three different privacy budgets with $\delta = \frac{1}{N^{1.1}}$, $\varepsilon = \{0.05, 0.2, 0.5\}$. See Appendix A for more details.

Table 1: Summary of Dataset Statistics.

Dataset	Samples	Attributes
California Housing	20640	8
Gas Turbine CO and NOx Emission	36733	9
Wine Quality	4898	11

Implementation Details. We implemented DP-SGD, DP-GLMtron, and DP-MBGLMtron for regression tasks, tuning hyperparameters to ensure fair comparisons. Specifically, we set the learning rate to 0.01 for DP-SGD and DP-MBGLMtron, while DP-GLMtron used a higher learning rate of 0.05 to account for its single-pass training strategy. Each model was trained for 500 epochs to allow sufficient training progress, with DP-MBGLMtron utilizing a mini-batch size of 32. To ensure the robustness of our findings, every experiment was repeated five times, and the average performance was reported along with standard deviations where applicable. The experiments were conducted on an NVIDIA A6000 GPU. Throughout the training, we monitored the training loss, validation loss, and gradient norms to track convergence and model stability under varying privacy constraints.

Experiment Results. We evaluated the implemented algorithms using two criteria: training loss and test loss, both measured against the number of training epochs. We report

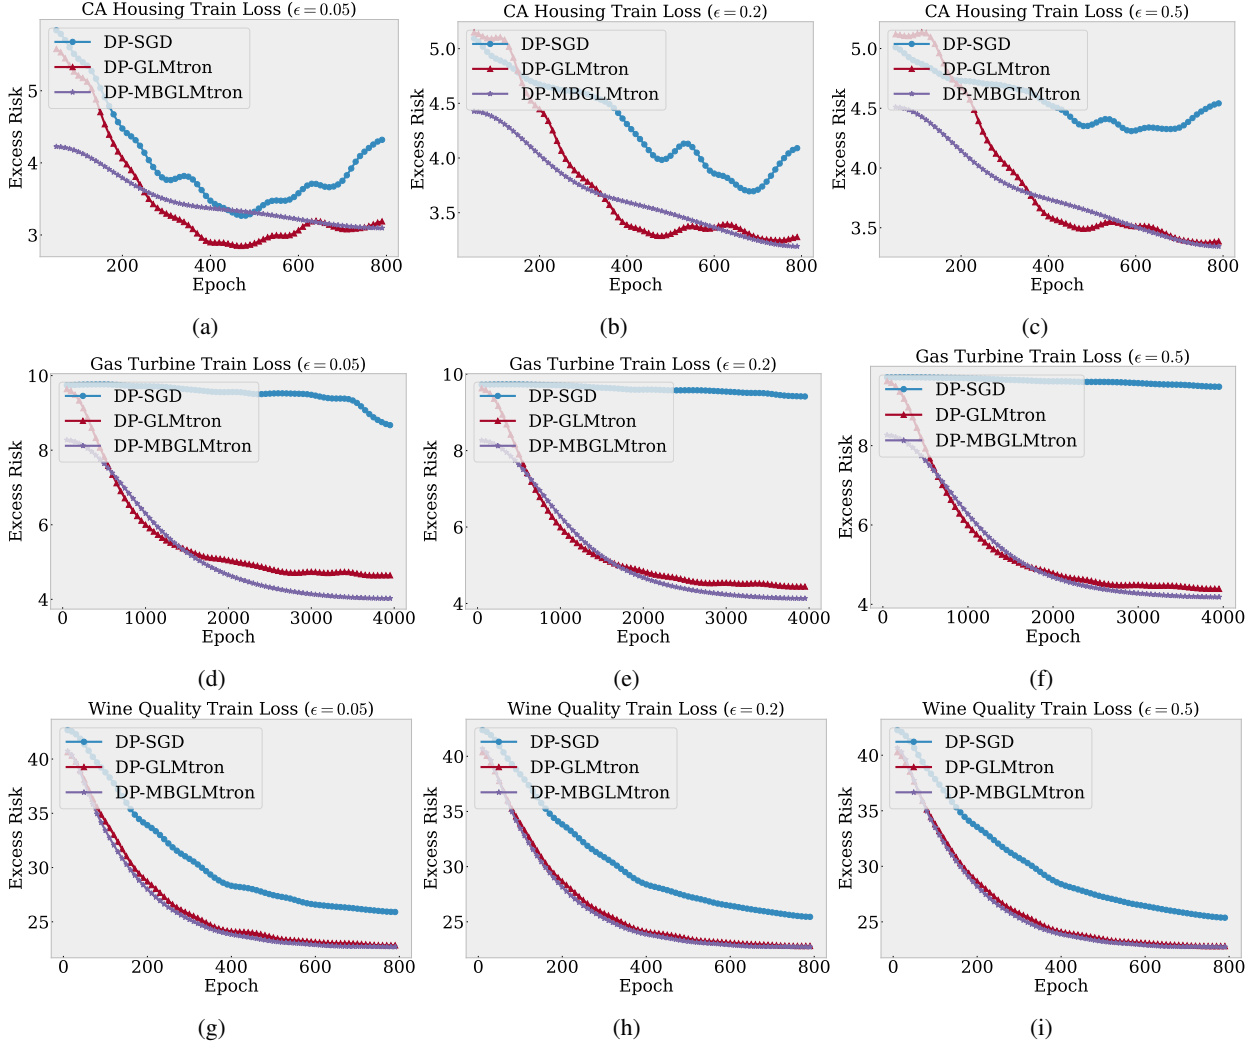


Figure 1: Training loss over epochs for DP-SGD, DP-GLMtrn, and DP-MBGLMtrn on three regression datasets: California Housing, Gas Turbine, and Wine Quality, under varying privacy budgets ($\epsilon = 0.05, 0.2, 0.5$)

the training loss here, with additional experimental results provided in Appendix A. From Fig. 1, we can see across all datasets and privacy budgets, DP-MBGLMtrn and DP-GLMtrn consistently outperform DP-SGD, achieving lower excess risk and faster convergence. These results indicate that the minibatch approach in DP-MBGLMtrn is particularly effective under strict privacy constraints, improving both stability and performance in differential privacy settings. The minibatch strategy in DP-MBGLMtrn allows for more frequent gradient updates, which helps mitigate the negative effects of privacy-induced noise and stabilize training. In contrast, DP-SGD struggles with convergence, particularly at smaller privacy budgets (e.g., $\epsilon = 0.05$). In Figures 1a-1f, the excess risks remain high or do not decrease as effectively as with DP-MBGLMtrn and DP-GLMtrn, highlighting DP-SGD’s limitations in maintaining performance in the ReLU regression model.

8 CONCLUSION

In this work, we revisited differentially private learning in the ReLU regression model under standard assumptions of i.i.d. data from $O(1)$ -sub-Gaussian distributions, presenting nearly optimal guarantees for excess population risk. We introduced and analyzed two algorithms: DP-GLMtrn and DP-MBGLMtrn. DP-GLMtrn leverages adaptive gradient clipping derived from additional public data, achieving an excess risk upper bound of $\tilde{O}(\frac{d^2}{N^2\epsilon^2})$. To mitigate its limitations regarding privacy budgets and public data dependence, we proposed DP-MBGLMtrn, which uses mini-batching to eliminate the need for public data and supports larger privacy budgets without compromising performance. Additionally, we established a matching lower bound through a tracing attack, confirming the tightness of our theoretical results. Empirical evaluations on regression tasks further validated our theoretical insights.

Acknowledgements

Di Wang is supported in part by the funding BAS/1/1689-01-01, URF/1/4663-01-01, REI/1/5232-01-01, REI/1/5332-01-01, and URF/1/5508-01-01 from KAUST, and funding from KAUST - Center of Excellence for Generative AI, under award number 5940.

References

- Gas Turbine CO and NO_x Emission Data Set. UCI Machine Learning Repository, 2019. DOI: <https://doi.org/10.24432/C5WC95>.
- Naman Agarwal, Zeyuan Allen-Zhu, Brian Bullins, Elad Hazan, and Tengyu Ma. Finding approximate local minima faster than gradient descent. In *Proceedings of the 49th Annual ACM SIGACT Symposium on Theory of Computing*, pages 1195–1199, 2017.
- Kareem Amin, Alex Kulesza, Andres Munoz, and Sergei Vassilvitskii. Bounding user contributions: A bias-variance trade-off in differential privacy. In *International Conference on Machine Learning*, pages 263–271. PMLR, 2019.
- Galen Andrew, Om Thakkar, Brendan McMahan, and Swaroop Ramaswamy. Differentially private learning with adaptive clipping. *Advances in Neural Information Processing Systems*, 34:17455–17466, 2021.
- Hilal Asi, John Duchi, Alireza Fallah, Omid Javidi, and Kunal Talwar. Private adaptive gradient methods for convex optimization. In *International Conference on Machine Learning*, pages 383–392. PMLR, 2021.
- Raef Bassily, Adam Smith, and Abhradeep Thakurta. Private empirical risk minimization: Efficient algorithms and tight error bounds. In *2014 IEEE 55th annual symposium on foundations of computer science*, pages 464–473. IEEE, 2014.
- Raef Bassily, Vitaly Feldman, Kunal Talwar, and Abhradeep Guha Thakurta. Private stochastic convex optimization with optimal rates. *Advances in neural information processing systems*, 32, 2019.
- Raef Bassily, Cristóbal Guzmán, and Michael Menart. Differentially private stochastic optimization: New results in convex and non-convex settings. *Advances in Neural Information Processing Systems*, 34:9317–9329, 2021a.
- Raef Bassily, Cristóbal Guzmán, and Anupama Nandi. Non-euclidean differentially private stochastic convex optimization. In *Conference on Learning Theory*, pages 474–499. PMLR, 2021b.
- Mark Bun and Thomas Steinke. Concentrated differential privacy: Simplifications, extensions, and lower bounds. In *Theory of Cryptography Conference*, pages 635–658. Springer, 2016.
- T Tony Cai, Yichen Wang, and Linjun Zhang. The cost of privacy: Optimal rates of convergence for parameter estimation with differential privacy. *The Annals of Statistics*, 49(5):2825–2850, 2021.
- Kamalika Chaudhuri, Claire Monteleoni, and Anand D Sarwate. Differentially private empirical risk minimization. *Journal of Machine Learning Research*, 12(3), 2011.
- Paulo Cortez, A. Cerdeira, F. Almeida, T. Matos, and J. Reis. Wine Quality. UCI Machine Learning Repository, 2009. DOI: <https://doi.org/10.24432/C56S3T>.
- Meng Ding, Mingxi Lei, Liyang Zhu, Shaowei Wang, Di Wang, and Jinhui Xu. Revisiting differentially private relu regression. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*.
- Meng Ding, Kaiyi Ji, Di Wang, and Jinhui Xu. Understanding forgetting in continual learning with linear regression. *arXiv preprint arXiv:2405.17583*, 2024.
- Simon S Du, Xiyu Zhai, Barnabas Poczos, and Aarti Singh. Gradient descent provably optimizes over-parameterized neural networks. *arXiv preprint arXiv:1810.02054*, 2018.
- Cynthia Dwork, Frank McSherry, Kobbi Nissim, and Adam Smith. Calibrating noise to sensitivity in private data analysis. In *Theory of Cryptography: Third Theory of Cryptography Conference, TCC 2006, New York, NY, USA, March 4-7, 2006. Proceedings 3*, pages 265–284. Springer, 2006.
- Vitaly Feldman, Tomer Koren, and Kunal Talwar. Private stochastic convex optimization: optimal rates in linear time. In *Proceedings of the 52nd Annual ACM SIGACT Symposium on Theory of Computing*, pages 439–449, 2020.
- Vitaly Feldman, Audra McMillan, and Kunal Talwar. Hiding among the clones: A simple and nearly optimal analysis of privacy amplification by shuffling. In *2021 IEEE 62nd Annual Symposium on Foundations of Computer Science (FOCS)*, pages 954–964. IEEE, 2022.
- Spencer Frei, Yuan Cao, and Quanquan Gu. Agnostic learning of a single neuron with gradient descent. *Advances in Neural Information Processing Systems*, 33:5417–5428, 2020.
- Surbhi Goel and Adam R Klivans. Learning neural networks with two nonlinear layers in polynomial time. In *Conference on Learning Theory*, pages 1470–1499. PMLR, 2019.

- Lijie Hu, Shuo Ni, Hanshen Xiao, and Di Wang. High dimensional differentially private stochastic optimization with heavy-tailed data. In *Proceedings of the 41st ACM SIGMOD-SIGACT-SIGAI Symposium on Principles of Database Systems*, pages 227–236, 2022.
- Roger Iyengar, Joseph P Near, Dawn Song, Om Thakkar, Abhradeep Thakurta, and Lun Wang. Towards practical differentially private convex optimization. In *2019 IEEE symposium on security and privacy (SP)*, pages 299–316. IEEE, 2019.
- Prateek Jain and Abhradeep Guha Thakurta. (near) dimension independent risk bounds for differentially private learning. In *International Conference on Machine Learning*, pages 476–484. PMLR, 2014.
- Prateek Jain, Pravesh Kothari, and Abhradeep Thakurta. Differentially private online learning. In *Conference on Learning Theory*, pages 24–1. JMLR Workshop and Conference Proceedings, 2012.
- Prateek Jain, Sham Kakade, Rahul Kidambi, Praneeth Netrapalli, and Aaron Sidford. Parallelizing stochastic gradient descent for least squares regression: mini-batching, averaging, and model misspecification. *Journal of machine learning research*, 18, 2018.
- Bargav Jayaraman, Lingxiao Wang, David Evans, and Quanquan Gu. Distributed learning without distrust: Privacy-preserving empirical risk minimization. *Advances in Neural Information Processing Systems*, 31, 2018.
- Sham M Kakade, Varun Kanade, Ohad Shamir, and Adam Kalai. Efficient learning of generalized linear and single index models with isotonic regression. *Advances in Neural Information Processing Systems*, 24, 2011.
- Vishesh Karwa and Salil Vadhan. Finite sample differentially private confidence intervals. *arXiv preprint arXiv:1711.03908*, 2017.
- Daniel Kifer, Adam Smith, and Abhradeep Thakurta. Private convex empirical risk minimization and high-dimensional regression. In *Conference on Learning Theory*, pages 25–1. JMLR Workshop and Conference Proceedings, 2012.
- Janardhan Kulkarni, Yin Tat Lee, and Daogao Liu. Private non-smooth empirical risk minimization and stochastic convex optimization in subquadratic steps. *arXiv preprint arXiv:2103.15352*, 2021.
- Xiyang Liu, Weihao Kong, Prateek Jain, and Sewoong Oh. Dp-pca: Statistically optimal and differentially private pca. *Advances in Neural Information Processing Systems*, 35:29929–29943, 2022.
- Xiyang Liu, Prateek Jain, Weihao Kong, Sewoong Oh, and Arun Suggala. Label robust and differentially private linear regression: Computational and statistical efficiency. In *Thirty-seventh Conference on Neural Information Processing Systems*, 2023a.
- Xiyang Liu, Prateek Jain, Weihao Kong, Sewoong Oh, and Arun Sai Suggala. Near optimal private and robust linear regression. *arXiv preprint arXiv:2301.13273*, 2023b.
- H Brendan McMahan, Daniel Ramage, Kunal Talwar, and Li Zhang. Learning differentially private recurrent language models. *arXiv preprint arXiv:1710.06963*, 2017.
- R. Kelley Pace and Ronald Barry. California housing data. StatLib Repository, 1997. URL <http://www.spatial-statistics.com>. Data obtained from the 1990 U.S. Census. The manuscript describing the data can be found at <http://www.spatial-statistics.com>. DOI: [https://doi.org/10.1016/S0167-7152\(97\)00107-X](https://doi.org/10.1016/S0167-7152(97)00107-X).
- Mark Rudelson and Roman Vershynin. Hanson-wright inequality and sub-gaussian concentration. 2013.
- Hanpu Shen, Cheng-Long Wang, Zihang Xiang, Yiming Ying, and Di Wang. Differentially private non-convex learning for multi-layer neural networks. *arXiv preprint arXiv:2310.08425*, 2023.
- Shuang Song, Om Thakkar, and Abhradeep Thakurta. Characterizing private clipped gradient descent on convex generalized linear problems. *arXiv preprint arXiv:2006.06783*, 2020.
- Shuang Song, Thomas Steinke, Om Thakkar, and Abhradeep Thakurta. Evading the curse of dimensionality in unconstrained private glms. In *International Conference on Artificial Intelligence and Statistics*, pages 2638–2646. PMLR, 2021.
- Jinyan Su and Di Wang. Faster rates of differentially private stochastic convex optimization. *arXiv preprint arXiv*, 2108, 2021.
- Jinyan Su, Changhong Zhao, and Di Wang. Differentially private stochastic convex optimization in (non)-euclidean space revisited. In *Uncertainty in Artificial Intelligence*, pages 2026–2035. PMLR, 2023.
- Jinyan Su, Lijie Hu, and Di Wang. Faster rates of differentially private stochastic convex optimization. *Journal of Machine Learning Research*, 25(114):1–41, 2024.
- Kunal Talwar, Abhradeep Thakurta, and Li Zhang. Private empirical risk minimization beyond the worst case: The effect of the constraint set geometry. *arXiv preprint arXiv:1411.5417*, 2014.
- Yuming Tao, Zuyuan Zhang, Dongxiao Yu, Xiuzhen Cheng, Falko Dressler, and Di Wang. Second-order convergence in private stochastic non-convex optimization. *arXiv preprint arXiv:2505.15647*, 2025.

- Prateek Varshney, Abhradeep Thakurta, and Prateek Jain. (nearly) optimal private linear regression via adaptive clipping. *arXiv preprint arXiv:2207.04686*, 2022.
- Di Wang and Jinhui Xu. Differentially private empirical risk minimization with smooth non-convex loss functions: A non-stationary view. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 33, pages 1182–1189, 2019a.
- Di Wang and Jinhui Xu. On sparse linear regression in the local differential privacy model. In *International Conference on Machine Learning*, pages 6628–6637. PMLR, 2019b.
- Di Wang and Jinhui Xu. Escaping saddle points of empirical risk privately and scalably via dp-trust region method. In *Machine Learning and Knowledge Discovery in Databases: European Conference, ECML PKDD 2020, Ghent, Belgium, September 14–18, 2020, Proceedings, Part III*, pages 90–106. Springer, 2021.
- Di Wang and Jinhui Xu. Gradient complexity and non-stationary views of differentially private empirical risk minimization. *Theoretical Computer Science*, 982: 114259, 2024.
- Di Wang, Minwei Ye, and Jinhui Xu. Differentially private empirical risk minimization revisited: Faster and more general. *Advances in Neural Information Processing Systems*, 30, 2017.
- Di Wang, Changyou Chen, and Jinhui Xu. Differentially private empirical risk minimization with non-convex loss functions. In *International Conference on Machine Learning*, pages 6526–6535. PMLR, 2019.
- Di Wang, Jiahao Ding, Lijie Hu, Zejun Xie, Miao Pan, and Jinhui Xu. Finite sample guarantees of differentially private expectation maximization algorithm. In *ECAI 2023*, pages 2435–2442. IOS Press, 2023a.
- Di Wang, Lijie Hu, Huanyu Zhang, Marco Gaboardi, and Jinhui Xu. Generalized linear models in non-interactive local differential privacy with public data. *Journal of Machine Learning Research*, 24(132):1–57, 2023b.
- Lingxiao Wang, Bargav Jayaraman, David Evans, and Quanquan Gu. Efficient privacy-preserving stochastic nonconvex optimization. In *Uncertainty in Artificial Intelligence*, pages 2203–2213. PMLR, 2023c.
- Jingfeng Wu, Difan Zou, Zixiang Chen, Vladimir Braverman, Quanquan Gu, and Sham M Kakade. Finite-sample analysis of learning high-dimensional single relu neuron. 2023.
- Xi Wu, Fengan Li, Arun Kumar, Kamalika Chaudhuri, Somesh Jha, and Jeffrey Naughton. Bolt-on differential privacy for scalable stochastic gradient descent-based analytics. In *Proceedings of the 2017 ACM International Conference on Management of Data*, pages 1307–1322, 2017.
- Hanshen Xiao, Zihang Xiang, Di Wang, and Srinivas Devadas. A theory to instruct differentially-private learning via clipping bias reduction. In *2023 IEEE Symposium on Security and Privacy (SP)*, pages 2170–2189. IEEE Computer Society, 2023.
- Jiaqi Zhang, Kai Zheng, Wenlong Mou, and Liwei Wang. Efficient private erm for smooth objectives. *arXiv preprint arXiv:1703.09947*, 2017.
- Qiuchen Zhang, Jing Ma, Jian Lou, and Li Xiong. Private stochastic non-convex optimization with improved utility rates. In *Proceedings of the Thirtieth International Joint Conference on Artificial Intelligence*, 2021.
- Ruijia Zhang, Mingxi Lei, Meng Ding, Zihang Xiang, Jinhui Xu, and Di Wang. Improved rates of differentially private nonconvex-strongly-concave minimax optimization. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 22524–22532, 2025.
- Yingxue Zhou, Zhiwei Steven Wu, and Arindam Banerjee. Bypassing the ambient dimension: Private sgd with gradient subspace identification. *arXiv preprint arXiv:2007.03813*, 2020.
- Liyang Zhu, Meng Ding, Vaneet Aggarwal, Jinhui Xu, and Di Wang. Improved analysis of sparse linear regression in local differential privacy model. *arXiv preprint arXiv:2310.07367*, 2023.
- Liyang Zhu, Amina Manseur, Meng Ding, Jinyan Liu, Jinhui Xu, and Di Wang. Truthful high dimensional sparse linear regression. *arXiv preprint arXiv:2410.13046*, 2024.
- Difan Zou, Jingfeng Wu, Vladimir Braverman, Quanquan Gu, and Sham Kakade. Benign overfitting of constant-stepsizes sgd for linear regression. In *Conference on Learning Theory*, pages 4633–4635. PMLR, 2021.

Nearly Optimal Differentially Private ReLU Regression (Supplementary Material)

Meng Ding² Mingxi Lei² Shaowei Wang³ Tianhang Zheng⁴ Di Wang^{†5} Jinhui Xu^{*1}

¹School of Information Science and Technology, University of Science and Technology of China

²Department of Computer Science and Engineering, State University of New York at Buffalo

³Institute of Artificial Intelligence and Blockchain, Guangzhou University

⁴The State Key Laboratory of Blockchain and Data Security, Zhejiang University

⁵Division of CEMSE, King Abdullah University of Science and Technology

A ADDITIONAL EXPERIMENT

Datasets Information. The information of three datasets used in our experiments is summarized in Appendix A.

Table 2: Summary of Dataset Statistics.

Dataset	Samples	Attributes
California Housing	20640	8
Gas Turbine CO and NOx Emission	36733	9
Wine Quality	4898	11

Experimental Results. Across all datasets and privacy budgets, DP-GLMtron and DP-MBGLMtron consistently outperform DP-SGD. The two methods converge faster and stabilize at test loss, suggesting that they are more effective at maintaining performance while adhering to privacy constraints. The trend holds consistent across varying privacy budgets ($\epsilon = 0.05, 0.2, 0.5$), further highlighting the robustness of our approaches.

Computing Infrastructures. The information of training configuration used in our experiments is summarized in Table 3.

Table 3: Hardware and Software Configuration.

Components	Details
Operating System	Ubuntu 16.04.6
CPU	AMD EPYC 7552, 48-Core Processor
CPU Memory	1.0 TB
GPU	NVIDIA RTX A6000
Programming Language	Python 3.9.12
Deep Learning Framework	Pytorch 1.12.1

To better illustrate our theoretical findings, we have added an additional experiment using synthetic data that satisfies our assumptions. In this setting, we could observe that the excess risk is closely related to the privacy budget in the table 4.

B ADDITIONAL DEFINITIONS

Definition 4 (zCDP Bun and Steinke [2016]). *A randomized algorithm \mathcal{A} is ρ -zCDP if for any pair of data sets D and D' that differ in one record, we have $D_p(\mathcal{A}(D) \parallel \mathcal{A}(D')) \leq \rho p$ for all $p > 1$, where D_p is the Rényi divergence of order p .*

*Correspondence to: Di Wang <di.wang@kaust.edu.sa>, Jinhui Xu <jh xu@ustc.edu.cn>

†Correspondence to: Di Wang <di.wang@kaust.edu.sa>, Jinhui Xu <jh xu@ustc.edu.cn>

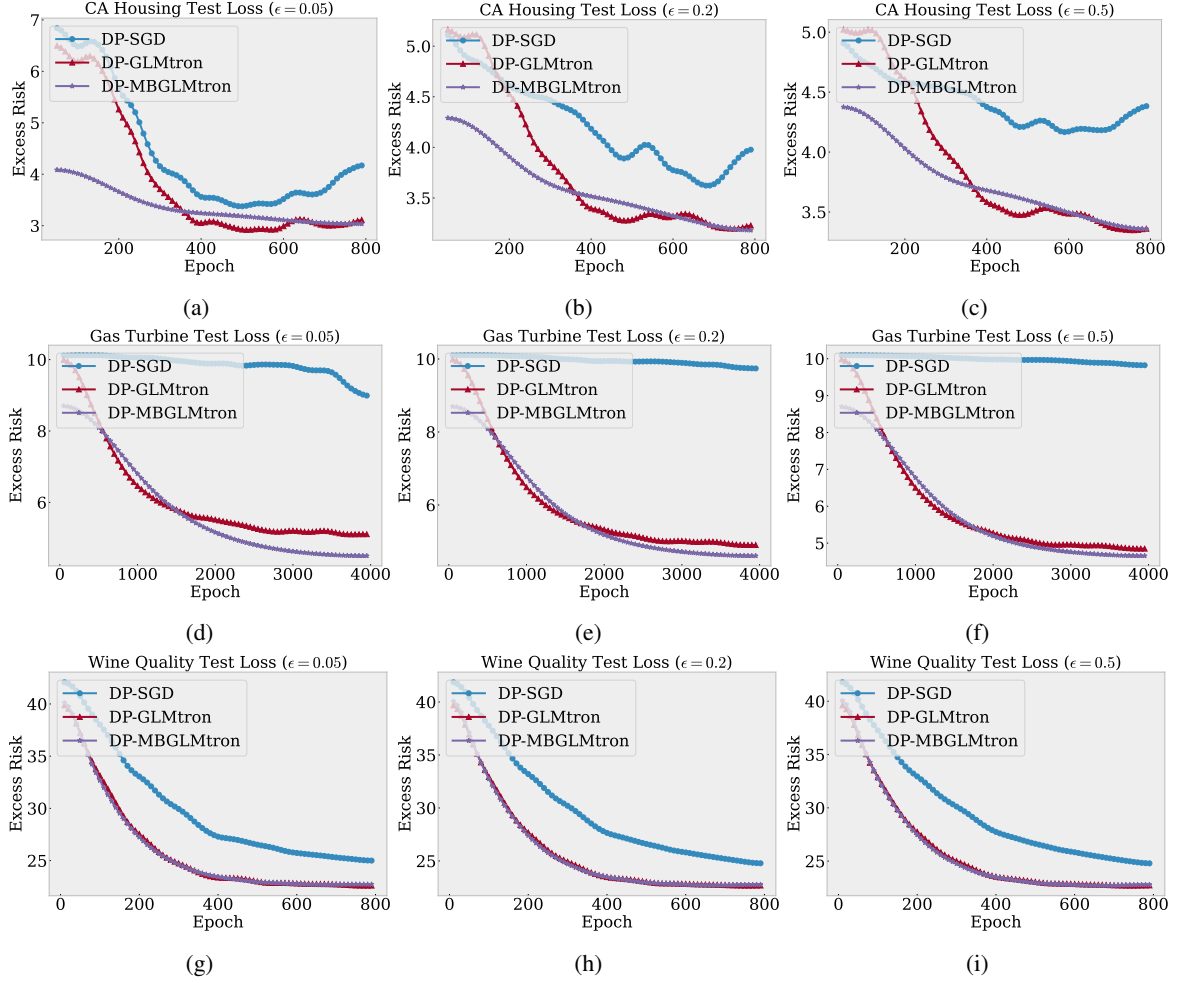


Figure 2: Test loss over epochs for DP-SGD, DP-GLMtron, and DP-MiniBatch GLMtron on three regression datasets: California Housing, Gas Turbine, and Blog Feedback, under varying privacy budgets ($\epsilon = 0.05, 0.2, 0.5$)

Definition 5 (Sub-Gaussian random variable). A zero-mean random variable $X \in \mathbb{R}$ is said to be sub-Gaussian with variance σ^2 ($X \sim \text{subG}(\sigma^2)$) if its moment generating function satisfies $\mathbb{E}[\exp(tX)] \leq \exp\left(\frac{\sigma^2 t^2}{2}\right)$ for all $t > 0$. For a sub-Gaussian random variable X , its sub-Gaussian norm $\|X\|_{\psi_2}$ is defined as $\|X\|_{\psi_2} = \inf\{c > 0 : \mathbb{E}[\exp\left(\frac{X^2}{c^2}\right)] \leq 2\}$. Specifically, if $X \sim \text{subG}(\sigma^2)$ we have $\|X\|_{\psi_2} \leq O(\sigma)$.

Definition 6 (Sub-Gaussian random vector). A zero mean random vector $\mathbf{X} \in \mathbb{R}^d$ is said to be sub-Gaussian with variance σ^2 (for simplicity, we call it σ^2 -sub-Gaussian), which is denoted as $(\mathbf{X} \sim \text{subG}_d(\sigma^2))$, if $\langle \mathbf{X}, \mathbf{u} \rangle$ is sub-Gaussian with variance σ^2 for any unit vector $\mathbf{u} \in \mathbb{R}^d$.

C DP-GLMTRON

C.1 PRIVACY GUARANTEE

To guarantee the privacy of DP-GLMtron, we should first ensure that each step of DP-GLMtron is private.

Lemma 8. Each update step of DP-GLMtron (Algorithm 1) ensures (ϵ_0, δ_0) -differential privacy, provided that $f = \frac{c}{\epsilon_0}$, where $c \geq \sqrt{2 \log(1.25/\delta_0)}$, and s_t denotes the clipping norm.

Proof. We first consider $\{\mathbf{w}_t\}_{t=0}^{N-1}$.

	ε	DP-SGD	DP-GLMtron	DP-MBGLMtron	DP-TAGLMtron
Train	0.05	9.14	8.26	5.19	7.77
	0.2	6.40	5.47	4.40	4.90
	0.5	5.77	4.83	4.38	4.55
Test	0.05	9.74	8.74	5.63	8.06
	0.2	7.38	6.25	4.82	5.14
	0.5	6.82	5.63	4.81	5.25

Table 4: Excess Risk (lower is better) under different privacy budgets.

Each update step (excluding the DP-noise addition) is of the form:

$$\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - \eta \text{clip}_{s_t}(\mathbf{x}_t(\text{ReLU}(\mathbf{x}_t^\top \mathbf{w}_t) - y_t)),$$

where $\text{clip}_{s_t}(\mathbf{v}) = \mathbf{v} \cdot \max\{1, \frac{s_t}{\|\mathbf{v}\|_2}\}$. Consequently, the local L_2 sensitivity of \mathbf{w}_{t+1} is determined by analyzing the variation in the t^{th} iteration data sample, as follows:

$$\begin{aligned} \Delta_2 &= \|\mathbf{w}'_{t+1} - \mathbf{w}_{t+1}\| \\ &= \|\eta \text{clip}_{s_t}(\mathbf{x}'_t(\text{ReLU}(\mathbf{x}'_t^\top \mathbf{w}_t) - y'_t)) - \eta \text{clip}_{s_t}(\mathbf{x}_t(\text{ReLU}(\mathbf{x}_t^\top \mathbf{w}_t) - y_t))\| \\ &\leq 2\eta \|\text{clip}_{s_t}(\mathbf{x}_t(\text{ReLU}(\mathbf{x}_t^\top \mathbf{w}_t) - y_t))\| \\ &= 2\eta s_t. \end{aligned}$$

Moreover, denoting $\varepsilon' = \varepsilon_0 / \sqrt{8N \log(2/\delta_0)}$ and $\delta' = \delta_0 / (2N)$, according to Lemma 2.3 in Karwa and Vadhan [2017], we could know that $\{s_t\}_{t=0}^N$ is (ε', δ') -DP.

□

Lemma 9 (Feldman et al. [2022]). *For a domain \mathcal{D} , let $\mathcal{R}^{(i)} : f \times \mathcal{D} \rightarrow \mathcal{S}^{(i)}$ for $i \in [n]$ (where $\mathcal{S}^{(i)}$ is the range space of $\mathcal{R}^{(i)}$) be a sequence of algorithms such that $\mathcal{R}^{(i)}(z_{1:i-1}, \cdot)$ is a $(\varepsilon_0, \delta_0)$ -DP local randomizer for all values of auxiliary inputs $z_{1:i-1} \in \mathcal{S}^{(1)} \times \dots \times \mathcal{S}^{(i-1)}$. Let $\mathcal{A}_s : \mathcal{D}^N \rightarrow \mathcal{S}^{(1)} \times \dots \times \mathcal{S}^{(n)}$ be the algorithm that given a dataset $x_{1:N} \in \mathcal{D}^N$, samples a uniformly random permutation π , then sequentially computes $z_i = \mathcal{R}^{(i)}(z_{1:i-1}, x_{\pi(i)})$ for $i \in [n]$, and outputs $z_{1:N}$. Then for any $\delta \in [0, 1]$ such that $\varepsilon_0 \leq \log(\frac{n}{16 \log(2/\delta)})$, \mathcal{A}_s is $(\varepsilon, \delta + O(e^\varepsilon \delta_0 n))$ -DP where ε is:*

$$\varepsilon = O((1 - e^{-\varepsilon_0}) \left(\frac{\sqrt{e^{\varepsilon_0} \log(1/\delta)}}{\sqrt{n}} + \frac{e^{\varepsilon_0}}{n} \right)).$$

We are now prepared to prove the privacy of DP-GLMtron, utilizing the lemmas discussed above.

Firstly, we reformulate the update rule into a sequence of one-step algorithms as follows:

$$\mathcal{R}^{(t+1)}(u_{0:t}, (\mathbf{x}, y)) := \mathbf{w}_{t+1} \leftarrow \mathbf{w}_t(u_{0:t}) - \eta \text{clip}_{s_t}(\mathbf{x}_{\pi(t)}(\text{ReLU}(\mathbf{x}_{\pi(t)}^\top \mathbf{w}_t(u_{0:t})) - y_{\pi(t)})) - 2\eta s_t f g_t,$$

where u denotes auxiliary inputs, and $\pi(t)$ represents the sample at the t -th iteration after randomly permuting the input data.

From Lemma 8, each $\mathcal{R}^{(t+1)}(u_{0:t}, \cdot)$ is a $(\varepsilon_0, \delta_0)$ -DP local randomizer algorithm, where $\varepsilon_0 \leq \log(\frac{N}{16 \log(2/\delta)})$. The output of DP-GLMtron is derived through post-processing of the shuffled outputs $u_{t+1} = \mathcal{R}^{(t+1)}(u_{0:t}, (\mathbf{x}, y))$ for $t \in 0, \dots, N-1$. Therefore, by Lemma 9, Algorithm DP-GLMtron adheres to $(\hat{\varepsilon}, \hat{\delta} + O(e^{\hat{\varepsilon}} \delta_0 N))$ -DP, where:

$$\hat{\varepsilon} = O((1 - e^{-\varepsilon_0}) \left(\frac{\sqrt{e^{\varepsilon_0} \log(1/\hat{\delta})}}{\sqrt{N}} + \frac{e^{\varepsilon_0}}{N} \right)).$$

Assuming $\varepsilon_0 \leq \frac{1}{2}$, we can infer the existence of some constant $c_1 > 0$ such that:

$$\hat{\varepsilon} \leq c_1 \cdot (1 - e^{-\varepsilon_0}) \left(\frac{\sqrt{e^{\varepsilon_0} \log(1/\hat{\delta})}}{\sqrt{N}} + \frac{e^{\varepsilon_0}}{N} \right)$$

$$\begin{aligned}
&\leq c_1 \cdot ((e^{\varepsilon_0/2} - e^{-\varepsilon_0/2})\sqrt{\frac{\log(1/\hat{\delta})}{N}} + (e^{\varepsilon_0} - 1)\frac{1}{N}) \\
&\leq c_1 \cdot (((1 + \varepsilon_0) - (1 - \varepsilon_0/2))\sqrt{\frac{\log(1/\hat{\delta})}{N}} + ((1 + 2\varepsilon_0) - 1)\frac{1}{N}) \\
&= c_1 \cdot \varepsilon_0 \left(\frac{1}{2}\sqrt{\frac{\log(1/\hat{\delta})}{N}} + \frac{2}{N} \right).
\end{aligned} \tag{6}$$

By setting $f = \frac{\sqrt{2 \log(1.25/\delta_0)}}{\varepsilon_0}$ in Lemma 8, we ensure that each update step of DP-GLMtron independently satisfies $(\varepsilon_0, \delta_0)$ -DP, based on standard Gaussian mechanism. Replacing $\varepsilon_0 = \frac{\sqrt{2 \log(1.25/\delta_0)}}{f}$, we obtain:

$$\hat{\varepsilon} \leq c_1 \cdot \frac{\sqrt{2 \log(1.25/\delta_0) \log(1/\hat{\delta})}}{f\sqrt{N}} \tag{7}$$

To satisfy overall (ε, δ) -DP, set $\hat{\delta} = \frac{\delta}{2}$, and $\delta_0 = c_2 \cdot \frac{\delta}{e^{\varepsilon N}}$ for some constant $c_2 > 0$. From this, we have:

$$\hat{\varepsilon} \leq c_1 \cdot \frac{\sqrt{2 \log(c_2 \cdot 1.25 \cdot e^{\varepsilon N}/\delta) \cdot \log(2/\delta)}}{f\sqrt{N}} \tag{8}$$

For any $\varepsilon \leq 1$, setting $f = c_3 \cdot \frac{\log(N/\delta)}{\varepsilon\sqrt{N}} \geq c_3 \frac{\sqrt{\log(N/\delta) \log(1/\delta)}}{\varepsilon\sqrt{N}}$ for a sufficiently large $c_3 > 0$ ensures that $\hat{\varepsilon} \leq \varepsilon$. Additionally, to fulfill Lemma 9's assumption, $\varepsilon_0 < \frac{1}{2}$ must be satisfied, which is attainable by setting $\varepsilon = O(\sqrt{\frac{\log(N/\delta)}{N}})$.

This implies that for $f = \Omega(\frac{\log(N/\delta)}{\varepsilon\sqrt{N}})$, DP-GLMtron achieves (ε, δ) -DP as long as $\varepsilon = O(\sqrt{\frac{\log(N/\delta)}{N}})$, thereby completing the proof.

C.2 UTILITY GUARANTEE

Here we first provide several auxiliary results that will be used in our DP-GLMtron utility bound.

Assumption 4 (Moment symmetricity conditions Wu et al. [2023]). *Assume that*

(A). *For every $\mathbf{u} \in \mathbb{H}$, it holds that*

$$\mathbb{E}[\mathbf{xx}^\top \cdot \mathbb{1}[\mathbf{x}^\top \mathbf{u} > 0]] = \mathbb{E}[\mathbf{xx}^\top \cdot \mathbb{1}[\mathbf{x}^\top \mathbf{u} < 0]].$$

(B). *For every $\mathbf{u} \in \mathbb{H}$ and $\mathbf{v} \in \mathbb{H}$, it holds that*

$$\mathbb{E}[\mathbf{xx}^\top \cdot \mathbb{1}[\mathbf{x}^\top \mathbf{u} > 0, \mathbf{x}^\top \mathbf{v} > 0]] = \mathbb{E}[\mathbf{xx}^\top \cdot \mathbb{1}[\mathbf{x}^\top \mathbf{u} < 0, \mathbf{x}^\top \mathbf{v} < 0]].$$

(C). *For every $\mathbf{u} \in \mathbb{H}$, it holds that*

$$\mathbb{E}[\mathbf{x}^{\otimes 4} \cdot \mathbb{1}[\mathbf{x}^\top \mathbf{u} > 0]] = \mathbb{E}[\mathbf{x}^{\otimes 4} \cdot \mathbb{1}[\mathbf{x}^\top \mathbf{u} < 0]].$$

(D). *For every $\mathbf{u} \in \mathbb{H}$ and $\mathbf{v} \in \mathbb{H}$, it holds that*

$$\mathbb{E}[(\mathbf{x}^\top \mathbf{v})^2 \mathbf{xx}^\top \cdot \mathbb{1}[\mathbf{x}^\top \mathbf{u} > 0, \mathbf{x}^\top \mathbf{v} > 0]] = \mathbb{E}[(\mathbf{x}^\top \mathbf{v})^2 \mathbf{xx}^\top \cdot \mathbb{1}[\mathbf{x}^\top \mathbf{u} < 0, \mathbf{x}^\top \mathbf{v} < 0]].$$

The following results are direct consequences of Assumption 4.

Lemma 10 (Wu et al. [2023]). *The following results are direct consequences of Assumption 4.*

1. *Under Assumption 4 (A), it holds that: for every vector $\mathbf{u} \in \mathbb{H}$,*

$$\mathbb{E}[\mathbf{xx}^\top \cdot \mathbb{1}[\mathbf{x}^\top \mathbf{u} > 0]] = \frac{1}{2} \cdot \mathbb{E}[\mathbf{xx}^\top] =: \frac{1}{2} \cdot \mathbf{H}.$$

2. Under Assumption 4 (C), it holds that: for every vector $\mathbf{u} \in \mathbb{H}$,

$$\mathbb{E}[\mathbf{x}^{\otimes 4} \cdot \mathbb{1}[\mathbf{x}^\top \mathbf{u} > 0]] = \frac{1}{2} \cdot \mathbb{E}[\mathbf{x}^{\otimes 4}] =: \frac{1}{2} \cdot \mathcal{M}$$

Lemma 11 (Rudelson and Vershynin [2013]). *Hanson-Wright Inequality: For any $\mathbf{X} \sim \mathcal{N}(\mathbf{0}, \Sigma)$, the following holds for $t \geq 0$*

$$\mathbb{P}(\|\mathbf{X}\|^2 \geq \text{Tr}(\Sigma) + 2\sqrt{t}\|\Sigma\|_F + 2t\|\Sigma\|_{op}) \leq e^{-t}.$$

Lemma 12. *Let $\mathbf{P}_t := \mathbf{I} - \eta \mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0] \mathbf{x}_t \mathbf{x}_t^\top$, where each $\mathbf{x}_t \in \mathbb{R}^d$, $\forall t \in \{j, \dots, T-1\}$ has been sampled i.i.d. from \mathcal{D} and $\mathbf{w}_{t-1} \in \mathbb{R}^d$ is the weight parameter for iteration $t-1$. Let $\mathbf{z} \in \mathbb{R}^d$ be a vector independent of all \mathbf{P}_t 's. Then for $b > 0$ and $\eta < \frac{1}{R_x^2}$, we have with probability $\geq 1 - b$:*

$$\|\mathbf{P}_{T-1} \mathbf{P}_{T-2} \dots \mathbf{P}_j \mathbf{z}\|^2 \leq \frac{1}{b} e^{-\eta\mu(T-j)} \|\mathbf{z}\|^2.$$

Proof. Note that

$$\begin{aligned} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{P}_{T-1} \mathbf{P}_{T-2} \dots \mathbf{P}_j \mathbf{z}\|^2] &= \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [(\mathbf{P}_{T-1} \mathbf{P}_{T-2} \dots \mathbf{P}_j \mathbf{z})^\top \mathbf{P}_{T-1} \mathbf{P}_{T-2} \dots \mathbf{P}_j \mathbf{z}] \\ &= \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\mathbf{z}^\top \mathbf{P}_j^\top \dots \mathbf{P}_{T-2}^\top \mathbf{P}_{T-1}^\top \mathbf{P}_{T-1} \mathbf{P}_{T-2} \dots \mathbf{P}_j \mathbf{z}] \\ &= \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\mathbf{z}^\top \mathbf{P}_j^\top \dots \mathbf{P}_{T-2}^\top \mathbb{E}_{(\mathbf{x}_{T-1}, y_{T-1}) \sim \mathcal{D}} [\mathbf{P}_{T-1}^\top \mathbf{P}_{T-1}] \mathbf{P}_{T-2} \dots \mathbf{P}_j \mathbf{z}]. \end{aligned}$$

We undertake a focused examination of the expectation $\mathbb{E}_{(\mathbf{x}_{T-1}, y_{T-1}) \sim \mathcal{D}} [\mathbf{P}_{T-1}^\top \mathbf{P}_{T-1}]$, considered independently.

$$\begin{aligned} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\mathbf{P}_{T-1}^\top \mathbf{P}_{T-1}] &= \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [(\mathbf{I} - \eta \mathbb{1}[\mathbf{x}_{T-1}^\top \mathbf{w}_{T-2} > 0] \mathbf{x}_{T-1} \mathbf{x}_{T-1}^\top) (\mathbf{I} - \eta \mathbb{1}[\mathbf{x}_{T-1}^\top \mathbf{w}_{T-2} > 0] \mathbf{x}_{T-1} \mathbf{x}_{T-1}^\top)^\top] \\ &= \mathbf{I} - \frac{\eta}{2} \mathbf{H} - \frac{\eta}{2} \mathbf{H} + \eta^2 \mathcal{M} \\ &\leq \mathbf{I} - \eta \mathbf{H} + \eta^2 R_x^2 \mathbf{H}, \end{aligned}$$

where the first equality is a direct result of Lemma 10 and the last inequality drives from the ?? and Definition 3. As a result, we have

$$\begin{aligned} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{P}_{T-1} \mathbf{P}_{T-2} \dots \mathbf{P}_j \mathbf{z}\|^2] &\leq \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\lambda_{\max}(\mathbf{I} - 2\eta \mathbf{H} + \eta^2 R_x^2 \mathbf{H}) \|\mathbf{P}_{T-2} \dots \mathbf{P}_j \mathbf{z}\|^2] \\ &= \lambda_{\max}(\mathbf{I} - 2\eta \mathbf{H} + \eta^2 R_x^2 \mathbf{H}) \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{P}_{T-2} \dots \mathbf{P}_j \mathbf{z}\|^2] \\ &\leq (1 - \eta(2 - \eta R_x^2) \mu) \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{P}_{T-2} \dots \mathbf{P}_j \mathbf{z}\|^2] \\ &\leq (1 - \eta \mu) \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{P}_{T-2} \dots \mathbf{P}_j \mathbf{z}\|^2] \\ &\leq \text{repeat the same procedure} \\ &\leq (1 - \eta \mu)^{(T-j)} \|\mathbf{z}\|^2 \\ &\leq e^{-\eta\mu(T-j)} \|\mathbf{z}\|^2, \end{aligned}$$

With Markov Inequality, for $b > 0$ indicating $\Pr\{\mathbf{Z} \geq \frac{\mathbb{E}[\mathbf{Z}]}{b}\} \leq b$, we have

$$\Pr\{\|\mathbf{P}_{T-1} \dots \mathbf{P}_j \mathbf{z}\|^2 \leq \frac{\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{P}_{T-1} \dots \mathbf{P}_j \mathbf{z}\|^2]}{b}\} \geq 1 - b.$$

Therefore with probability at least $1 - b$:

$$\|\mathbf{P}_{T-1} \dots \mathbf{P}_0 \mathbf{z}\|^2 \leq \frac{1}{b} e^{-\eta\mu(T-j)} \|\mathbf{z}\|^2.$$

□

Lemma 13. Let η be stepsize such that $\eta \leq \min\{\frac{\lambda_d}{\lambda_1 R_x^2 \log^{2a} N}, \frac{1}{3f\sqrt{d}}\}$, where $c_1, c_2 > 0$ are global constants and $\Gamma = 4C_2 R_x \cdot \log^{2a} N \cdot (\sqrt{\|\mathbf{H}\|_2} \|\mathbf{w}_*\| + \sqrt{\kappa}\sigma)$. Furthermore, let $f = f_{\varepsilon, \delta, N}$ be a function of ε, δ, N . Then, with probability $\geq 1 - \frac{1}{N^{100}}$, $\|\mathbf{x}_t(\langle \mathbf{x}_t, \mathbf{w}_t \rangle - y_t)\| \leq \Gamma$ for all $0 \leq t \leq N-1$; \mathbf{w}_t is the t^{th} iterate of Algorithm DP-GLMtron.

Proof. We begin by examining the base case when $t = 0$ and the norm of the "gradient" can be expressed as:

$$\begin{aligned} \|\mathbf{x}_0(\text{ReLU}(\mathbf{x}_0^\top \mathbf{w}_0) - y_0)\| &= \|\mathbf{x}_0(\text{ReLU}(\mathbf{x}_0^\top \cdot \mathbf{0}) - y_0)\| \\ &= \|\mathbf{x}_0 y_0\| \\ &\leq \|\mathbf{x}_0\| |\mathbf{x}_0^\top \mathbf{w}_* + z_0|. \end{aligned}$$

By the distribution of \mathbf{x} and Definition 3, w.p. at least $1 - b_x$, we have:

$$\|\mathbf{x}_0\| \leq R_x \log^a(1/b_x),$$

and by the triangle inequality $|\mathbf{x}_0^\top \mathbf{w}_* + z_0| \leq \|\mathbf{x}_0^\top \mathbf{w}_*\| + \|z_0\|$, w.p. at least $1 - b_{\mathbf{w}_*} - b_\sigma$, we have:

$$|\mathbf{x}_0^\top \mathbf{w}_*| + |z_0| \leq C_2 \sqrt{\|\mathbf{H}\|_2} \|\mathbf{w}_*\| \log^a(1/b_{\mathbf{w}_*}) + \sigma C_2 \log^a(1/b_\sigma).$$

Since each b is $1/\text{poly}(N)$, the lemma holds.

Now let us assume that the Lemma is valid for the $(t-1)$ -th iteration. Proceeding with this assumption, we turn our attention to the t -th iteration:

$$\|\mathbf{x}_t(\text{ReLU}(\mathbf{x}_t^\top \mathbf{w}_t) - y_t)\| = \|\mathbf{x}_t(\max(0, \mathbf{x}_t^\top \mathbf{w}_t) - y_t)\|.$$

It is obvious that when $\mathbf{x}_t^\top \mathbf{w}_t \leq 0$, the norm of gradient simplifies to $\|\mathbf{x}_t y_t\|$, which aligns closely with base case.

If $\mathbf{x}_t^\top \mathbf{w}_t, \mathbf{x}_t^\top \mathbf{w}_* \geq 0$, then we will have

$$\begin{aligned} \|\mathbf{x}_t(\max(0, \mathbf{x}_t^\top \mathbf{w}_t) - y_t)\| &= \|\mathbf{x}_t \mathbf{x}_t^\top (\mathbf{w}_t - \mathbf{w}_*) + \mathbf{x}_t z_t\| \\ &\leq \|\mathbf{x}\| (\|\mathbf{x}_t^\top (\mathbf{w}_t - \mathbf{w}_*)\| + \|z_t\|) \\ &\leq R_x \log^a(1/b_x) (C_2 \sqrt{\|\mathbf{H}\|_2} \|\mathbf{w}_t - \mathbf{w}_*\| \log^a(1/b_{\mathbf{w}_t}) + \sigma C_2 \log^a(1/b_\sigma)) \\ &= C_2 R_x \log^{2a} N (\sqrt{\|\mathbf{H}\|_2} \|\mathbf{w}_t - \mathbf{w}_*\| + \sigma), \end{aligned}$$

where $b_x, b_{\mathbf{w}_t}, b_\sigma$ is $1/\text{poly}(N)$.

If $\mathbf{x}_t^\top \mathbf{w}_t \geq 0$ and $\mathbf{x}_t^\top \mathbf{w}_* \leq 0$, then we will have:

$$\begin{aligned} \|\mathbf{x}_t(\max(0, \mathbf{x}_t^\top \mathbf{w}_t) - y_t)\| &= \|\mathbf{x}_t \mathbf{x}_t^\top \mathbf{w}_t + \mathbf{x}_t z_t\| \\ &\leq \|\mathbf{x}\| (\|\mathbf{x}_t^\top (\mathbf{w}_t - \mathbf{w}_*)\| + \|\mathbf{w}_*\| + \|z_t\|) \\ &\leq R_x \log^a(1/b_x) (C_2 \sqrt{\|\mathbf{H}\|_2} \|\mathbf{w}_t - \mathbf{w}_*\| \log^a(1/b_{\mathbf{w}_t}) + \|\mathbf{w}_*\| + \sigma C_2 \log^a(1/b_\sigma)) \\ &= C_2 R_x \log^{2a} N (\sqrt{\|\mathbf{H}\|_2} \|\mathbf{w}_t - \mathbf{w}_*\| + \|\mathbf{w}_*\| + \sigma), \end{aligned} \tag{9}$$

Given that the threshold s_{t-1} has not been exceeded in iterations, we can observe the following decomposition at iteration $t-1$:

$$\begin{aligned} \mathbf{w}_t - \mathbf{w}_* &= \mathbf{w}_{t-1} - \mathbf{w}_* - \eta(\text{clip}_{s_{t-1}}(\mathbf{x}_{t-1}(\text{ReLU}(\mathbf{x}_{t-1}^\top \mathbf{w}_{t-1}) - y_{t-1})) + 2s_{t-1} f \mathbf{g}_{t-1}) \\ &= \mathbf{w}_{t-1} - \mathbf{w}_* - \eta(\mathbf{x}_{t-1}(\text{ReLU}(\mathbf{x}_{t-1}^\top \mathbf{w}_{t-1}) - y_{t-1}) + 2s_{t-1} f \mathbf{g}_{t-1}) \\ &= \mathbf{w}_{t-1} - \mathbf{w}_* - \eta \mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0] \cdot \mathbf{x}_t \mathbf{x}_t^\top \mathbf{w}_{t-1} + \eta \mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_* > 0] \cdot \mathbf{x}_t \mathbf{x}_t^\top \mathbf{w}_* \\ &\quad + \eta z_t \mathbf{x}_t + 2\eta s_{t-1} f \mathbf{g}_{t-1} \\ &= \mathbf{w}_{t-1} - \mathbf{w}_* - \eta \mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0] \cdot \mathbf{x}_t \mathbf{x}_t^\top (\mathbf{w}_{t-1} - \mathbf{w}_*) \\ &\quad + \eta (\mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0]) \cdot \mathbf{x}_t \mathbf{x}_t^\top \mathbf{w}_* + \eta z_t \mathbf{x}_t - 2\eta s_{t-1} f \mathbf{g}_{t-1} \\ &= (\mathbf{I} - \eta \mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0] \mathbf{x}_t \mathbf{x}_t^\top) (\mathbf{w}_{t-1} - \mathbf{w}_*) \\ &\quad + \eta (\mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0]) \mathbf{x}_t \mathbf{x}_t^\top \mathbf{w}_* + \eta z_t \mathbf{x}_t - 2\eta s_{t-1} f \mathbf{g}_{t-1}. \end{aligned} \tag{10}$$

We introduce the following notations for clarity

$$\mathbf{P}_t := \mathbf{I} - \eta \mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0] \mathbf{x}_t \mathbf{x}_t^\top, \quad \mathbf{u}_t = (\mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0]) \mathbf{x}_t \mathbf{x}_t^\top \mathbf{w}_*, \quad \mathbf{v}_t = z_t \mathbf{x}_t - 2\Gamma f \mathbf{g}_{t-1}.$$

Then the expected inner product w.r.t \mathbf{H} can be reformulated as follows.

$$\begin{aligned} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{w}_t - \mathbf{w}^*\|_{\mathbf{H}}^2] &= \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{P}_t(\mathbf{w}_{t-1} - \mathbf{w}_*) + \eta \mathbf{u}_{t-1} + \eta \mathbf{v}_{t-1}\|_{\mathbf{H}}^2] = \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{P}_t(\mathbf{w}_{t-1} - \mathbf{w}_*) + \eta \mathbf{u}_{t-1} + \eta \mathbf{v}_{t-1}\|_{\mathbf{H}}^2] \\ &= \underbrace{\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{P}_t(\mathbf{w}_{t-1} - \mathbf{w}_*)\|_{\mathbf{H}}^2]}_{\text{(quadratic term 1)}} + \underbrace{\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\eta \mathbf{u}_{t-1}\|_{\mathbf{H}}^2]}_{\text{(quadratic term 1)}} + \underbrace{\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\eta \mathbf{v}_{t-1}\|_{\mathbf{H}}^2]}_{\text{(quadratic term 1)}} \\ &\quad + 2 \underbrace{\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\mathbf{u}_{t-1}^\top \mathbf{H}(\mathbf{P}_t(\mathbf{w}_{t-1} - \mathbf{w}_*))] + \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\mathbf{v}_{t-1}^\top \mathbf{H}(\mathbf{P}_t(\mathbf{w}_{t-1} - \mathbf{w}_*))] + \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\mathbf{u}_{t-1}^\top \mathbf{H} \mathbf{v}_{t-1}]}_{\text{(crossing term)}}. \end{aligned} \tag{11}$$

where the cross terms involving z and \mathbf{g}_t have zero expectation, attributable to the fact that $\mathbb{E}[z \mid \mathbf{x}_t] = 0$ and $\mathbb{E}[\mathbf{g}_t \mid \mathbf{x}_t] = 0$.

For the second quadratic term in Equation (11), we observe the following

$$(\mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0])^2 = \mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0, \mathbf{x}_t^\top \mathbf{w}_* < 0] + \mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} < 0, \mathbf{x}_t^\top \mathbf{w}_* > 0]$$

Then, we have

$$\begin{aligned} &\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\text{(quadratic term 2)}] \\ &= \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [(\mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0])^2 (\mathbf{x}_t \mathbf{x}_t^\top \mathbf{w}_*)^\top (\mathbf{x}_t \mathbf{x}_t^\top \mathbf{w}_*)] \\ &= \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [(\mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0, \mathbf{x}_t^\top \mathbf{w}_* < 0] + \mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} < 0, \mathbf{x}_t^\top \mathbf{w}_* > 0]) \cdot (\mathbf{w}_*^\top \mathbf{x}_t)^2 \cdot \mathbf{x}_t^\top \mathbf{x}_t] \\ &= 2 \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0, \mathbf{x}_t^\top \mathbf{w}_* < 0] \cdot (\mathbf{w}_*^\top \mathbf{x}_t)^2 \cdot \mathbf{x}_t^\top \mathbf{x}_t], \end{aligned}$$

where the last equation follows from Assumption 3. Similarly, for the crossing terms in Equation (11), we have

$$\begin{aligned} &\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\text{(crossing term)}] \\ &= \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [(\mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0])^2 \cdot (\mathbf{x}_t \mathbf{x}_t^\top \mathbf{w}_*)^\top (\mathbf{I} - \eta \mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0] \cdot \mathbf{x}_t \mathbf{x}_t^\top) (\mathbf{w}_{t-1} - \mathbf{w}_*)] \\ &= 2 \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0, \mathbf{x}_t^\top \mathbf{w}_* < 0] \cdot (\mathbf{x}_t \mathbf{x}_t^\top \mathbf{w}_*)^\top (\mathbf{I} - \eta \mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0] \cdot \mathbf{x}_t \mathbf{x}_t^\top) (\mathbf{w}_{t-1} - \mathbf{w}_*)] \\ &= 2 \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0, \mathbf{x}_t^\top \mathbf{w}_* < 0] \cdot (\mathbf{x}_t \mathbf{x}_t^\top \mathbf{w}_*)^\top (\mathbf{w}_{t-1} - \mathbf{w}_* - \eta \mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0] \cdot \mathbf{x}_t \mathbf{x}_t^\top (\mathbf{w}_{t-1} - \mathbf{w}_*))] \\ &= 2 \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0, \mathbf{x}_t^\top \mathbf{w}_* < 0] \cdot (\mathbf{w}_*^\top \mathbf{x}_t \mathbf{x}_t^\top (\mathbf{w}_{t-1} - \mathbf{w}_*) - \eta \mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0] \cdot \mathbf{w}_*^\top \mathbf{x}_t \mathbf{x}_t^\top \mathbf{x}_t \mathbf{x}_t^\top (\mathbf{w}_{t-1} - \mathbf{w}_*))] \\ &= 2 \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0, \mathbf{x}_t^\top \mathbf{w}_* < 0] \cdot (1 - \eta \mathbf{x}_t^\top \mathbf{x}_t) \cdot \mathbf{w}_*^\top \mathbf{x}_t \cdot \mathbf{x}_t^\top (\mathbf{w}_{t-1} - \mathbf{w}_*)]. \end{aligned}$$

By applying the indicator function and considering $\eta \leq 1/R_x^2$, we deduce that

$$(1 - \eta \mathbf{x}_t^\top \mathbf{x}_t) \cdot \mathbf{x}_t^\top (\mathbf{w}_{t-1} - \mathbf{w}_*) \geq 0 \quad \text{and} \quad \mathbf{w}_*^\top \mathbf{x}_t \leq 0 \Rightarrow \mathbb{E}(\text{crossing term}) \leq 0$$

By invoking Assumption 3, it indicates that

$$\mathbb{E}(\mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0] \cdot \mathbf{x}_t \mathbf{x}_t^\top) = \frac{1}{2} \mathbf{H}.$$

Moreover, if $\eta \leq \frac{1}{2R_x^2}$, it holds that

$$\begin{aligned} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\mathbf{P}_{T-1}^\top \mathbf{P}_{T-1}] &= \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [(\mathbf{I} - \eta \mathbb{1}[\mathbf{x}_{T-1}^\top \mathbf{w}_{T-2} > 0] \mathbf{x}_{T-1} \mathbf{x}_{T-1}^\top) (\mathbf{I} - \eta \mathbb{1}[\mathbf{x}_{T-1}^\top \mathbf{w}_{T-2} > 0] \mathbf{x}_{T-1} \mathbf{x}_{T-1}^\top)^\top] \\ &= \mathbf{I} - \frac{\eta}{2} \mathbf{H} - \frac{\eta}{2} \mathbf{H} + \eta^2 \mathcal{M} \end{aligned}$$

$$\leq \mathbf{I} - \eta \mathbf{H} + \eta^2 R_x^2 \mathbf{H} \leq \mathbf{I} - \frac{\eta}{2} \mathbf{H}.$$

Combining the above results, the update of $\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{w}_t - \mathbf{w}^*\|_{\mathbf{H}}^2]$ holds that

$$\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{w}_t - \mathbf{w}^*\|_{\mathbf{H}}^2] \leq (1 - \frac{\eta\mu}{2}) \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{w}_{t-1} - \mathbf{w}^*\|_{\mathbf{H}}^2] + \sigma^2 \eta^2 \text{Tr}(\mathbf{H}) + \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [4\eta^2 s_{t-1}^2 f^2] \text{Tr}(\mathbf{H}).$$

Considering that the Equation (9) and the adaptive clipping algorithm, we have

$$s_t \leq R_x C_2 \log^{2a} N (\sqrt{\|\mathbf{H}\|} \|\mathbf{w}_t - \mathbf{w}_*\| + \|\mathbf{w}_*\| + \sigma + \Delta).$$

As a results, the update of $\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{w}_t - \mathbf{w}^*\|_{\mathbf{H}}^2]$ can be reformulated as

$$\begin{aligned} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{w}_t - \mathbf{w}^*\|_{\mathbf{H}}^2] &\leq (1 - \frac{\eta\mu}{2}) \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{w}_{t-1} - \mathbf{w}^*\|_{\mathbf{H}}^2] \\ &\quad + \sigma^2 \eta^2 \text{Tr}(\mathbf{H}) + 16R_x^2 C_2^2 \log^{4a} N \eta^2 f^2 \text{Tr}(\mathbf{H}) (\kappa \|\mathbf{w}_t - \mathbf{w}^*\|_{\mathbf{H}}^2 + \|\mathbf{w}_*\|^2 + \sigma^2 + \Delta^2) \\ &= (1 - (\frac{\eta\mu}{2} - 16\eta^2 f^2 C_2^2 R_x^2 \kappa \log^{4a} N \text{Tr}(\mathbf{H}))) \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{w}_{t-1} - \mathbf{w}^*\|_{\mathbf{H}}^2] \\ &\quad + \eta^2 \sigma^2 \text{Tr}(\mathbf{H}) + 16\eta^2 f^2 C_2^2 R_x^2 \log^{4a} N (\sigma^2 + \Delta^2 + \|\mathbf{w}_*\|^2) \text{Tr}(\mathbf{H}), \end{aligned}$$

where $\Delta = \frac{\|\mathbf{w}^*\|_{\mathbf{H} + \sigma}}{N^{100}}$ and we use the fact $\mathbf{I}\mu \preceq \mathbf{H} \implies \|\mathbf{H}\| \|\mathbf{w}_{t-1} - \mathbf{w}^*\|^2 \leq \kappa \|\mathbf{w}_{t-1} - \mathbf{w}^*\|_{\mathbf{H}}^2$.

If $\frac{\eta\mu}{4} \geq 16\eta^2 f^2 C_2^2 R_x^2 \kappa \log^{4a} N \text{Tr}(\mathbf{H})$, it means the step size η satisfies that $\eta \leq \frac{\mu}{64f^2 C_2^2 R_x^2 \kappa \log^{4a} N \text{Tr}(\mathbf{H})}$, therefore it holds that

$$\begin{aligned} &\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{w}_t - \mathbf{w}^*\|_{\mathbf{H}}^2] \\ &\leq (1 - \frac{\eta\mu}{4}) \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{w}_{t-1} - \mathbf{w}^*\|_{\mathbf{H}}^2] + \eta^2 \sigma^2 \text{Tr}(\mathbf{H}) + 16\eta^2 f^2 C_2^2 R_x^2 \log^{4a} N (\sigma^2 + \Delta^2 + \|\mathbf{w}_*\|^2) \text{Tr}(\mathbf{H}) \\ &\leq (1 - \eta\mu/4)^t \|\mathbf{w}_0 - \mathbf{w}^*\|_{\mathbf{H}}^2 + \frac{2}{\eta\mu} (\eta^2 \sigma^2 \text{Tr}(\mathbf{H}) + 16\eta^2 f^2 C_2^2 R_x^2 \log^{4a} N (\sigma^2 + \Delta^2 + \|\mathbf{w}_*\|^2) \text{Tr}(\mathbf{H})) \\ &\leq e^{-\eta\mu t/4} \|\mathbf{w}^*\|_{\mathbf{H}}^2 + \frac{2}{\eta\mu} (\eta^2 \sigma^2 \text{Tr}(\mathbf{H}) + 16\eta^2 f^2 C_2^2 R_x^2 \log^{4a} N (\sigma^2 + \Delta^2 + \|\mathbf{w}_*\|^2) \text{Tr}(\mathbf{H})). \end{aligned}$$

Since $s_t \leq R_x C_2 \log^{2a} N (\sqrt{\|\mathbf{H}\|} \|\mathbf{w}_t - \mathbf{w}^*\| + \|\mathbf{w}_*\| + \sigma + \Delta)$, the bound on $\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [s_t^2]$ will be

$$\begin{aligned} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [s_t^2] &\leq 4C_2^2 R_x^2 \log^{4a} N (\kappa \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{w}_t - \mathbf{w}^*\|_{\mathbf{H}}^2] + \sigma^2 + \|\mathbf{w}_*\|^2 + \Delta^2) \\ &\leq 4C_2^2 R_x^2 \log^{4a} N (\kappa (e^{-\eta\mu/4t} \|\mathbf{w}^*\|_{\mathbf{H}}^2 + \frac{2\eta\sigma^2}{\mu} \text{Tr}(\mathbf{H})) \\ &\quad + \frac{32\eta\alpha^2}{\mu} C_2^2 R_x^2 \log^{4a} N (\sigma^2 + \|\mathbf{w}_*\|^2 + \Delta^2) \text{Tr}(\mathbf{H})) + \sigma^2 + \|\mathbf{w}_*\|^2 + \Delta^2, \end{aligned}$$

which is decreasing with t (w.p. $\geq 1 - \frac{1}{\text{Poly}(N)}$).

Thus, if we define Γ s.t.

$$\begin{aligned} \Gamma^2 &= \max\{\text{Upper-Bound}(\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [s_0^2]), \dots, \text{Upper-Bound}(\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [s_T^2])\} \\ &= \text{Upper-Bound}(\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [s_0^2]) \\ &= 4C_2^2 R_x^2 \log^{4a} N (\kappa \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{w}_0 - \mathbf{w}^*\|_{\mathbf{H}}^2] + \sigma^2 + \|\mathbf{w}_*\|^2 + \Delta^2). \end{aligned}$$

□

Lemma 14 (Generic bounds on the DP-GLMtron iterates Wu et al. [2023]). *Suppose that Assumption 3 holds. Considering the DP-GLMtron algorithm, we have the following recursion:*

- $\mathbf{A}_t \preceq \mathbf{A}_{t-1} - \frac{\eta}{2}(\mathbf{H}\mathbf{A}_{t-1} + \mathbf{A}_{t-1}\mathbf{H}) + \eta^2\mathcal{M} \circ \mathbf{A}_{t-1} + \eta^2\sigma^2\mathbf{H} + 4\eta^2\Gamma^2f^2\mathbf{I}$
- $\mathbf{A}_t \succeq \mathbf{A}_{t-1} - \frac{\eta}{2}(\mathbf{H}\mathbf{A}_{t-1} + \mathbf{A}_{t-1}\mathbf{H}) + \frac{\eta^2}{4}\mathcal{M} \circ \mathbf{A}_{t-1} + \eta^2\sigma^2\mathbf{H} + 4\eta^2\Gamma^2f^2\mathbf{I}$

where $\mathbf{A}_t := \mathbb{E}(\mathbf{w}_t - \mathbf{w}_*)(\mathbf{w}_t - \mathbf{w}_*)^\top$, $t \geq 0$

Now consider the recursion of \mathbf{A}_t given in Lemma 14. Note that \mathbf{A}_t is related to \mathbf{A}_{t-1} through a linear operator, therefore \mathbf{A}_t can be understood as the sum of two iterates, i.e., $\mathbf{A}_t := \mathbf{B}_t + \mathbf{C}_t$, where

$$\begin{cases} \mathbf{B}_t \preceq (\mathcal{I} - \frac{\eta}{2} \cdot \mathcal{T}(2\eta)) \circ \mathbf{B}_{t-1}; \\ \mathbf{B}_0 = (\mathbf{w}_0 - \mathbf{w}_*)^{\otimes 2} \end{cases} \quad \begin{cases} \mathbf{C}_t \preceq (\mathcal{I} - \frac{\eta}{2} \cdot \mathcal{T}(2\eta)) \circ \mathbf{C}_{t-1} + \eta^2\sigma^2\mathbf{H} + 4\eta^2\Gamma^2f^2\mathbf{I}; \\ \mathbf{C}_0 = 0 \end{cases} \quad (12)$$

and

$$\begin{cases} \mathbf{B}_t \succeq (\mathcal{I} - \frac{\eta}{2} \cdot \mathcal{T}(\frac{\eta}{2})) \circ \mathbf{B}_{t-1}; \\ \mathbf{B}_0 = (\mathbf{w}_0 - \mathbf{w}_*)^{\otimes 2} \end{cases} \quad \begin{cases} \mathbf{C}_t \succeq (\mathcal{I} - \frac{\eta}{2} \cdot \mathcal{T}(\frac{\eta}{2})) \circ \mathbf{C}_{t-1} + \eta^2\sigma^2\mathbf{H} + 4\eta^2\Gamma^2f^2\mathbf{I}; \\ \mathbf{C}_0 = 0 \end{cases} \quad (13)$$

where

$$\begin{cases} (\mathcal{I} - \frac{\eta}{2} \cdot \mathcal{T}(2\eta)) \circ \mathbf{A}_{t-1} := \mathbf{A}_{t-1} - \frac{\eta}{2}(\mathbf{H}\mathbf{A}_{t-1} + \mathbf{A}_{t-1}\mathbf{H}) + \eta^2\mathcal{M} \circ \mathbf{A}_{t-1}; \\ (\mathcal{I} - \frac{\eta}{2} \cdot \mathcal{T}(\frac{\eta}{2})) \circ \mathbf{A}_{t-1} := \mathbf{A}_{t-1} - \frac{\eta}{2}(\mathbf{H}\mathbf{A}_{t-1} + \mathbf{A}_{t-1}\mathbf{H}) + \frac{\eta^2}{4}\mathcal{M} \circ \mathbf{A}_{t-1} \end{cases}$$

Besides, since our DP-GLM-tron is run with constant stepsize η and outputs the average of the iterates:

$$\bar{\mathbf{w}}_N := \frac{1}{N} \sum_{t=0}^{N-1} \mathbf{w}_t. \quad (14)$$

Then, the following lemma holds:

Lemma 15. *Suppose that Assumption 3 hold. For $\bar{\mathbf{w}}_N$ defined in Equation (14), we have that*

$$\begin{aligned} \mathbb{E}(\mathbf{H}, (\bar{\mathbf{w}}_N - \mathbf{w}_*)^{\otimes 2}) &\leq \sum_{t=0}^{N-1} \sum_{k=t}^{N-1} \frac{1}{\eta N^2} \langle (\mathbf{I} - \frac{\eta}{2}\mathbf{H})^{k-t}\mathbf{H}, \mathbf{A}_t \rangle, \\ \mathbb{E}(\mathbf{H}, (\bar{\mathbf{w}}_N - \mathbf{w}_*)^{\otimes 2}) &\geq \sum_{t=0}^{N-1} \sum_{k=t}^{N-1} \frac{1}{2\eta N^2} \langle (\mathbf{I} - \frac{\eta}{2}\mathbf{H})^{k-t}\mathbf{H}, \mathbf{A}_t \rangle. \end{aligned}$$

Proof.

$$\begin{aligned} \mathbb{E}[\mathbf{w}_t - \mathbf{w}_* \mid \mathbf{w}_{t-1}] &= \mathbb{E}[(\mathbf{I} - \eta \mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0])\mathbf{x}_t \mathbf{x}_t^\top (\mathbf{w}_{t-1} - \mathbf{w}_*) \mid \mathbf{w}_{t-1}] + 2\eta \Gamma f \mathbb{E}[\mathbf{g}_{t-1} \mid \mathbf{w}_{t-1}] \\ &\quad + \eta \cdot \mathbb{E}[(\mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0])\mathbf{x}_t \mathbf{x}_t^\top \mathbf{w}_* \mid \mathbf{w}_{t-1}] + \eta \mathbb{E}[z_t \mathbf{x}_t \mid \mathbf{w}_{t-1}] \\ &= \mathbb{E}[(\mathbf{I} - \eta \mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0])\mathbf{x}_t \mathbf{x}_t^\top (\mathbf{w}_{t-1} - \mathbf{w}_*) \mid \mathbf{w}_{t-1}] \\ &= (\mathbf{I} - \frac{\eta}{2}\mathbf{H})(\mathbf{w}_{t-1} - \mathbf{w}_*) \end{aligned}$$

The remaining proof simply follows Zou et al. [2021]. □

From the decomposition presented in Equation (12) and Equation (13), we know that $\sum_{t=0}^N \mathbf{A}_t = \sum_{t=0}^N \mathbf{B}_t + \sum_{t=0}^N \mathbf{C}_t$. With this foundation, we can now bound the bias and variance terms separately.

Variance error

For $t = 0$ we have $\mathbf{C}_0 = 0 \preceq \frac{\eta\sigma^2}{1-\eta R_x^2}\mathbf{I} + \frac{4\eta\Gamma^2f^2}{1-\eta R_x^2}\mathbf{H}^{-1}$.

We then assume that $\mathbf{C}_{t-1} \preceq \frac{\eta\sigma^2}{1-\eta R_x^2}\mathbf{I} + \frac{4\eta\Gamma^2f^2}{1-\eta R_x^2}\mathbf{H}^{-1}$, and exam \mathbf{C}_t based on Equation (12):

$$\begin{aligned} \mathbf{C}_t &\preceq (\mathcal{I} - \eta \cdot \mathcal{T}(\eta)) \circ \mathbf{C}_{t-1} + \eta^2\sigma^2\mathbf{H} + 4\eta^2\Gamma^2f^2\mathbf{I} \\ &= \mathbf{C}_{t-1} - \frac{\eta}{2}(\mathbf{H}\mathbf{C}_{t-1} + \mathbf{C}_{t-1}\mathbf{H}) + \eta^2\mathcal{M} \circ \mathbf{C}_{t-1} + \eta^2\sigma^2\mathbf{H} + 4\eta^2\Gamma^2f^2\mathbf{I} \end{aligned}$$

$$\begin{aligned}
&\preceq \frac{\eta\sigma^2}{1-\eta R_x^2}\mathbf{I} + \frac{4\eta\Gamma^2 f^2}{1-\eta R_x^2}\mathbf{H}^{-1} - \eta\left(\frac{\eta\sigma^2}{1-\eta R_x^2}\mathbf{H} + \frac{4\eta\Gamma^2 f^2}{1-\eta R_x^2}\mathbf{I}\right) \\
&+ \eta^2 R_x^2\left(\frac{\eta\sigma^2}{1-\eta R_x^2}\mathbf{H} + \frac{4\eta\Gamma^2 f^2}{1-\eta R_x^2}\mathbf{I}\right) + \eta^2\sigma^2\mathbf{H} + 4\eta^2\Gamma^2 f^2\mathbf{I} \\
&\preceq \frac{\eta\sigma^2}{1-\eta R_x^2}\mathbf{I} + \frac{4\eta\Gamma^2 f^2}{1-\eta R_x^2}\mathbf{H}^{-1}.
\end{aligned}$$

For the simplicity, we define $\Sigma := \sigma^2\mathbf{H} + 4\Gamma^2 f^2\mathbf{I}$. By the definitions of \mathcal{T} and $\widetilde{\mathcal{T}}$, we have:

$$\begin{aligned}
\mathbf{C}_t &= (\mathcal{I} - \frac{\eta}{2} \cdot \mathcal{T}(2\eta)) \circ \mathbf{C}_{t-1} + \eta^2 \Sigma \\
&= (\mathcal{I} - \frac{\eta}{2} \cdot \widetilde{\mathcal{T}}(2\eta)) \circ \mathbf{C}_{t-1} + \eta^2 (\mathcal{M} - \frac{1}{4}\widetilde{\mathcal{M}}) \circ \mathbf{C}_{t-1} + \eta^2 \Sigma \\
&\preceq (\mathcal{I} - \frac{\eta}{2} \cdot \widetilde{\mathcal{T}}(2\eta)) \circ \mathbf{C}_{t-1} + \eta^2 \mathcal{M} \circ \mathbf{C}_{t-1} + \eta^2 \Sigma,
\end{aligned}$$

where the last inequality is due to the fact that $\widetilde{\mathcal{M}}$ is a PSD mapping. Then by the iteration of variance, we have for all $t \geq 0$,

$$\mathcal{M} \circ \mathbf{C}_t \preceq \mathcal{M} \circ \left(\frac{\eta\sigma^2}{1-\eta R_x^2}\mathbf{I} + \frac{4\eta\Gamma^2 f^2}{1-\eta R_x^2}\mathbf{H}^{-1} \right) \preceq \frac{\eta\sigma^2 R_x^2}{1-\eta R_x^2}\mathbf{H} + \frac{4\eta\Gamma^2 f^2 R_x^2}{1-\eta R_x^2}\mathbf{I}.$$

Substituting the above into the previous result, we obtain

$$\begin{aligned}
\mathbf{C}_t &\preceq (\mathcal{I} - \frac{\eta}{2} \cdot \widetilde{\mathcal{T}}(2\eta)) \circ \mathbf{C}_{t-1} + \frac{\eta^3 R_x^2}{1-\eta R_x^2} \cdot (\sigma^2\mathbf{H} + 4\Gamma^2 f^2\mathbf{I}) + \eta^2 \Sigma \\
&= (\mathcal{I} - \frac{\eta}{2} \cdot \widetilde{\mathcal{T}}(2\eta)) \circ \mathbf{C}_{t-1} + \frac{\eta^3 R_x^2}{1-\eta R_x^2} \cdot (\sigma^2\mathbf{H} + 4\Gamma^2 f^2\mathbf{I}) + \eta^2 (\sigma^2\mathbf{H} + 4\Gamma^2 f^2\mathbf{I}) \\
&= (\mathcal{I} - \frac{\eta}{2} \cdot \widetilde{\mathcal{T}}(2\eta)) \circ \mathbf{C}_{t-1} + \frac{\eta^2}{1-\eta R_x^2} \cdot (\sigma^2\mathbf{H} + 4\Gamma^2 f^2\mathbf{I}).
\end{aligned}$$

It follows

$$\begin{aligned}
\mathbf{C}_t &\preceq \frac{\eta^2}{1-\eta R_x^2} \cdot \sum_{k=0}^{t-1} (\mathcal{I} - \frac{\eta}{2} \cdot \widetilde{\mathcal{T}}(2\eta))^k \circ (\sigma^2\mathbf{H} + 4\Gamma^2 f^2\mathbf{I}) \\
&= \frac{\eta^2}{1-\eta R_x^2} \cdot \sum_{k=0}^{t-1} (\mathbf{I} - \frac{\eta}{2}\mathbf{H})^k (\sigma^2\mathbf{H} + 4\Gamma^2 f^2\mathbf{I}) (\mathbf{I} - \frac{\eta}{2}\mathbf{H})^k \\
&\preceq \frac{\eta^2}{1-\eta R_x^2} \cdot \sum_{k=0}^{t-1} (\mathbf{I} - \frac{\eta}{2}\mathbf{H})^k (\sigma^2\mathbf{H} + 4\Gamma^2 f^2\mathbf{I}) \\
&= \frac{\eta\sigma^2}{1-\eta R_x^2} \cdot (\mathbf{I} - (\mathbf{I} - \frac{\eta}{2}\mathbf{H})^t) + \frac{4\eta\Gamma^2 f^2}{1-\eta R_x^2} \cdot (\mathbf{I} - (\mathbf{I} - \frac{\eta}{2}\mathbf{H})^t) \cdot \mathbf{H}^{-1}.
\end{aligned}$$

Consequently, the variance error can be represented as follows, in accordance with Lemma 15

$$\begin{aligned}
\text{variance error} &\leq \frac{1}{\eta N^2} \langle \mathbf{I} - (\mathbf{I} - \frac{\eta}{2}\mathbf{H})^N, \sum_{t=0}^N \mathbf{C}_t \rangle \\
&\leq \frac{1}{\eta N^2} \langle \mathbf{I} - (\mathbf{I} - \frac{\eta}{2}\mathbf{H})^N, \sum_{t=0}^N \frac{\eta\sigma^2}{1-\eta R_x^2} \cdot (\mathbf{I} - (\mathbf{I} - \frac{\eta}{2}\mathbf{H})^t) \rangle \\
&\quad + \frac{1}{\eta N^2} \langle \mathbf{I} - (\mathbf{I} - \frac{\eta}{2}\mathbf{H})^N, \sum_{t=0}^N \frac{4\eta\Gamma^2 f^2}{1-\eta R_x^2} \cdot (\mathbf{I} - (\mathbf{I} - \frac{\eta}{2}\mathbf{H})^t) \cdot \mathbf{H}^{-1} \rangle \\
&\leq \frac{\sigma^2}{(1-\eta R_x^2)N} \langle \mathbf{I} - (\mathbf{I} - \frac{\eta}{2}\mathbf{H})^N, (\mathbf{I} - (\mathbf{I} - \frac{\eta}{2}\mathbf{H})^N) \rangle
\end{aligned}$$

$$+ \frac{4\Gamma^2 f^2}{(1 - \eta R_x^2)N} \langle \mathbf{I} - (\mathbf{I} - \frac{\eta}{2}\mathbf{H})^N, (\mathbf{I} - (\mathbf{I} - \frac{\eta}{2}\mathbf{H})^N) \cdot \mathbf{H}^{-1} \rangle.$$

Therefore, by integrating the Γ and f , the variance error follows that

$$\text{variance error} \lesssim \frac{d\sigma^2}{N} + \frac{\Gamma^2 f^2}{N} \cdot \text{tr}(\mathbf{H}^{-1}) \lesssim \frac{d\sigma^2}{N} + \frac{d^2 \log^2(N/\delta)}{N^2 \varepsilon^2} \cdot C_2^2 \kappa^2 (\sigma^2 + \|\mathbf{w}_*\|_{\mathbf{H}}^2 + \Delta^2).$$

Bias error

Now we consider the bias error, which depends on the initial error regardless of noise. According to Lemma 15, the bias error of average iterate follows that

$$\text{bias error} \leq \sum_{t=0}^{N-1} \sum_{k=t}^{N-1} \frac{1}{\eta N^2} \langle (\mathbf{I} - \frac{\eta}{2}\mathbf{H})^{k-t} \mathbf{H}, \mathbf{B}_t \rangle \leq \frac{1}{\eta N^2} \langle \mathbf{I} - (\mathbf{I} - \frac{\eta}{2}\mathbf{H})^N, \sum_{t=0}^N \mathbf{B}_t \rangle \leq \sum_{t=0}^N \frac{1}{\eta N^2} \text{tr}(\mathbf{B}_t).$$

Considering the recursion of \mathbf{B}_t , we have $\mathbf{B}_t \preceq \mathbf{B}_{t-1} - \frac{\eta}{2}(\mathbf{H}\mathbf{B}_{t-1} + \mathbf{B}_{t-1}\mathbf{H}) + \eta^2 \mathcal{M} \circ \mathbf{B}_{t-1}$, which indicates the recursion of \mathbf{B}_t follows that

$$\begin{aligned} \mathbf{B}_t &\preceq \mathbf{B}_{t-1} - \eta \mathbf{H} \mathbf{B}_{t-1} + \eta^2 R_x^2 \mathbf{H} \mathbf{B}_{t-1} \\ &\preceq (\mathbf{I} - \frac{\eta}{2}\mathbf{H}) \mathbf{B}_{t-1} \end{aligned}$$

The last inequality derives from the choice of step size. Consequently, the bias error will be

$$\text{bias error} \leq \sum_{t=0}^N \frac{1}{\eta N^2} \text{tr}(\mathbf{B}_t) \leq \sum_{t=0}^N \frac{1}{\eta N^2} (1 - \frac{\eta \mu}{2})^t \text{tr}(\mathbf{B}_0) \leq \frac{1}{\eta N} \|\mathbf{w}_*\|^2 \lesssim \frac{d \log^2(N/\delta)}{N^2 \varepsilon^2} \cdot C_2^2 \kappa^2 \|\mathbf{w}_*\|^2.$$

Combining the previous variance error, we complete the proof.

D DP-MBGLMTRON

For DP-MBGLMtron algorithm, we perform the following update

$$\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - \eta \left(\frac{1}{b} \sum_{i=0}^{b-1} \text{clip}_{s_t}(\mathbf{x}_{\tau+m+i}(\text{ReLU}(\mathbf{x}_{\tau+m+i}^\top \mathbf{w}_t) - y_{\tau+m+i})) + f \cdot \frac{2s_t}{b} \cdot \mathbf{g}_t \right).$$

D.1 PRIVACY GUARANTEE

Lemma 16. *Algorithm DP - mini-batch-GLMtron with noise multiplier f satisfies $\frac{1}{f^2}$ -zCDP, and correspondingly satisfies (ε, δ) -differential privacy when we set the noise multiplier $f \geq \frac{2\sqrt{\log(1/\delta)+\varepsilon}}{\varepsilon}$. Furthermore, if $\varepsilon \leq \log(1/\delta)$, then $f \geq \frac{\sqrt{8\log(1/\delta)}}{\varepsilon}$ suffices to ensure (ε, δ) -differential privacy.*

We first show the step of gradient estimation is $\frac{1}{2f^2}$ -zCDP.

Notice the update of c has sensitivity one and the variance of DP noise is $\lceil \log_2(B/\Delta) \rceil f^2$, hence, each step is $\frac{1}{2\lceil \log_2(B/\Delta) \rceil f^2}$ -zCDP.

If we take at most $\lceil \log_2(B/\Delta) \rceil f^2$ operations, we will have the aggregated privacy accumulation $\frac{1}{2f^2}$, which completes the privacy guarantee of gradient estimation.

Now we turn our attention to the update of \mathbf{w} and consider the step without the Gaussian noise.

$$\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - \frac{\eta}{b} \sum_{i=0}^{b-1} \text{clip}_{s_t}(\mathbf{x}_{t,i}(\text{ReLU}(\mathbf{x}_{t,i}^\top \mathbf{w}_t) - y_{t,i})).$$

where $\text{clip}_\Gamma(\nu) = \nu \cdot \max\{1, \frac{\Gamma}{\|\nu\|_2}\}$. Therefore, the local L_2 sensitivity of the \mathbf{w}_{t+1} due to a sample difference in the t -th batch is $\Delta_2 = \frac{2\eta s_t}{b}$. Meanwhile, we know the variance of DP noise is $\frac{2\eta s_t f}{b}$, the above step is $\frac{1}{2f^2}$ -zCDP since $\frac{\Delta_2^2}{2 \cdot \frac{4\eta^2 s_t^2 f^2}{b^2}} = \frac{1}{2f^2}$.

According to the previous results and composition theorem, we know each iteration step is $\frac{1}{f^2}$ -zCDP. In our algorithm, every individual data point, denoted as (\mathbf{x}_i, y_i) , where i is an index belonging to the set of all indices N , is included in precisely one mini-batch, which indicates the algorithm traverses the complete dataset exactly once, thereby ensuring that each data point is processed in a single iteration. Hence, according to the parallel composition of zCDP, DP-mini-batch-FLMtron is $\frac{1}{f^2}$ -zCDP.

Recall that ρ -zCDP implies a $(\mu, \mu\rho)$ -RDP. We aim to optimize for any $\mu \geq 1$ and verify that the noise scaler f prescribed in the theorem satisfies (ε, δ) -Differential Privacy.

It is noted that $(\mu, \mu\rho)$ -RDP implies (ε, δ) -Approximate Privacy where $\varepsilon = \mu\rho + \frac{\log(1/\delta)}{\mu-1}$ for all $\mu > 1$. The minimum value of ε , denoted as ε_{\min} , which equals $\rho + 2\sqrt{\rho \log(1/\delta)}$, is obtained when the derivative of ε with respect to μ is zero, yielding $\mu = 1 + \sqrt{\log(1/\delta)/\rho}$.

For a given ε , we seek to minimize f (which scales as $1/\sqrt{\rho}$), such that the computed maximum allowable ρ ensures that $\varepsilon_{\min}(\rho) \leq \varepsilon$. Since $\varepsilon_{\min}(\rho)$ is a monotonically increasing function of f and forms a second-order polynomial in $\sqrt{\rho}$ with its vertex corresponding to the maximum at $\varepsilon_{\min}(\rho) = \varepsilon$, we obtain the following relation:

$$\frac{1}{f^2} = (\sqrt{\log(1/\delta)} + \varepsilon - \sqrt{\log(1/\delta)})^2 = \frac{\varepsilon^2}{(\sqrt{\log(1/\delta)} + \varepsilon + \sqrt{\log(1/\delta)})^2}$$

As the derived f satisfied (ε, δ) -DP, it is deduced that $f \geq \frac{2\sqrt{\log(1/\delta)+\varepsilon}}{\varepsilon}$, which ensures the algorithm's compliance with (ε, δ) -Differential Privacy.

D.2 UTILITY GUARANTEE

Similar to DP-GLMtron, we also provide several auxiliary results that will be used in our utility analysis.

Lemma 17. *If $\eta \leq \frac{b}{R_x^2 + (b-1)\|\mathbf{H}\|}$, $\tau(t) = t \cdot (b+s)$ and $\mathbf{P}_t := (\mathbf{I} - \frac{\eta}{b} \sum_{i=0}^{b-1} \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} \geq 0] \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top)$, then $\forall t$, $\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}}[\mathbf{P}_t^\top \mathbf{P}_t] \preceq \mathbf{I} - \eta \mathbf{H}$.*

Proof. Note that

$$\begin{aligned} & \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}}[\mathbf{P}_t^\top \mathbf{P}_t] \\ &= \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}}[(\mathbf{I} - \frac{\eta}{b} \sum_{i=0}^{b-1} \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} \geq 0] \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top)^\top (\mathbf{I} - \frac{\eta}{b} \sum_{i=0}^{b-1} \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} \geq 0] \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top)] \\ &= \mathbf{I} - \frac{2\eta}{b} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}}[\sum_{i=0}^{b-1} \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} \geq 0] \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top] \\ &+ \frac{\eta^2}{b^2} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}}[\sum_{i=0}^{b-1} \sum_{j=0}^{b-1} \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} \geq 0] \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \cdot \mathbb{1}[\mathbf{x}_{\tau(t)+m+j}^\top \mathbf{w}_{t-1} \geq 0] \mathbf{x}_{\tau(t)+m+j} \mathbf{x}_{\tau(t)+m+j}^\top] \end{aligned}$$

where we know $\mathbb{E}(\mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0] \cdot \mathbf{x}_t \mathbf{x}_t^\top \otimes \mathbf{x}_t \mathbf{x}_t^\top) \preceq \mathbb{E}(\mathbf{x}_t \mathbf{x}_t^\top \otimes \mathbf{x}_t \mathbf{x}_t^\top) = \mathcal{M}$, it indicates that

$$\begin{aligned} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}}[\mathbf{P}_t^\top \mathbf{P}_t] &\preceq \mathbf{I} - \eta \mathbf{H} + \frac{\eta^2}{b^2} (b\mathcal{M} + b(b-1)\mathbf{H}\mathbf{H}) \quad (*) \\ &\preceq \mathbf{I} - \eta \mathbf{H} + \frac{\eta^2}{b^2} (bR_x^2 \mathbf{H} + b(b-1)\|\mathbf{H}\|\mathbf{H}) \\ &= \mathbf{I} - \eta \mathbf{H} (1 - \frac{\eta}{b} (R_x^2 + (b-1)\|\mathbf{H}\|)). \end{aligned}$$

With the assumption of stepsize $\eta \leq \frac{1}{(R_x^2 + (b-1)\|\mathbf{H}\|)}$, we complete the proof. \square

Lemma 18. If $\mathbf{v}_t = \frac{1}{b} \sum_{i=0}^{b-1} z_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i} - \frac{2sf}{b} \mathbf{g}_t$, and $\tau(t) = t \cdot (b + s)$, then $\forall t$

$$\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}}[\mathbf{v}_t \mathbf{v}_t^\top] = \frac{1}{b} \Sigma + \frac{4f^2}{b^2} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}}[s_t^2] \mathbf{I},$$

where $\Sigma := \sum_{i=0}^{b-1} \sum_{j=0}^{b-1} (z_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i})(z_{\tau(t)+m+j} \mathbf{x}_{\tau(t)+m+j})^\top$.

Proof. Note that

$$\begin{aligned} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}}[\mathbf{v}_t \mathbf{v}_t^\top] &= \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}}\left[\left(\frac{1}{b} \sum_{i=0}^{b-1} z_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i} - \frac{2sf}{b} \mathbf{g}_t\right) \left(\frac{1}{b} \sum_{i=0}^{b-1} z_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i} - \frac{2sf}{b} \mathbf{g}_t\right)^\top\right] \\ &= \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}}\left[\left(\frac{1}{b} \sum_{i=0}^{b-1} z_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}\right) \left(\frac{1}{b} \sum_{i=0}^{b-1} z_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}\right)^\top\right] + \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}}\left[\frac{4s_t^2 f^2}{b^2} \mathbf{g}_t \mathbf{g}_t^\top\right] \\ &= \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}}\left[\frac{1}{b^2} \sum_{i=0}^{b-1} \sum_{j=0}^{b-1} (z_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i})(z_{\tau(t)+m+j} \mathbf{x}_{\tau(t)+m+j})^\top\right] + \frac{4f^2}{b^2} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}}[s_t^2] \mathbf{I}, \end{aligned}$$

where we have utilized the fact that $\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}}[(z_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i})(z_{\tau(t)+m+j} \mathbf{x}_{\tau(t)+m+j})^\top] = \mathbf{0}$ for $i \neq j$ that stems from the independence of samples and the fact that $\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}}[\mathbf{g}_j] = \mathbf{0}$ has been sampled independently at each step. \square

Considering no clipping, the t -th update is given by:

$$\mathbf{w}_{t+1} = \mathbf{w}_t - \frac{\eta}{b} \sum_{i=0}^{b-1} \mathbf{x}_{\tau(t)+m+i} (\text{ReLU}(\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_t) - y_{\tau(t)+m+i}) - \frac{2\eta s_t f}{b} \mathbf{g}_t,$$

where $\tau(t) = t \cdot (b + s)$. Hence, we could derive the following

$$\begin{aligned} \mathbf{w}_{t+1} - \mathbf{w}_* &= (\mathbf{I} - \frac{\eta}{b} \sum_{i=0}^{b-1} \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_t > 0] \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top) (\mathbf{w}_t - \mathbf{w}_*) \\ &+ \frac{\eta}{b} \sum_{i=0}^{b-1} (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_t > 0]) \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* + \frac{\eta}{b} \sum_{i=0}^{b-1} z_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i} - \frac{2\eta \Gamma \alpha}{b} \mathbf{g}_t \\ &:= \mathbf{P}_t (\mathbf{w}_{t-1} - \mathbf{w}_*) + \eta \mathbf{u}_t + \eta \mathbf{v}_t, \end{aligned}$$

where we denote

$$\begin{aligned} \mathbf{P}_t &= (\mathbf{I} - \frac{\eta}{b} \sum_{i=0}^{b-1} \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0] \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top) \\ \mathbf{u}_t &= \frac{1}{b} \sum_{i=0}^{b-1} (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0]) \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* \\ \mathbf{v}_t &= \frac{1}{b} \sum_{i=0}^{b-1} z_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i} - \frac{2\eta \Gamma \alpha}{b} \mathbf{g}_t \end{aligned}$$

Let us consider the expected inner product w.r.t \mathbf{H} :

$$\begin{aligned} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}}[\|\mathbf{w}_t - \mathbf{w}_*\|_{\mathbf{H}}^2] &= \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}}[\|\mathbf{P}_t (\mathbf{w}_{t-1} - \mathbf{w}_*) + \eta \mathbf{u}_{t-1} + \eta \mathbf{v}_{t-1}\|_{\mathbf{H}}^2] \\ &= \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}}[\|\mathbf{P}_t (\mathbf{w}_{t-1} - \mathbf{w}_*)\|_{\mathbf{H}}^2] + \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}}[\|\eta \mathbf{u}_{t-1}\|_{\mathbf{H}}^2] + \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}}[\|\eta \mathbf{v}_{t-1}\|_{\mathbf{H}}^2] \end{aligned}$$

$$+ 2 \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\mathbf{u}_{t-1} \mathbf{H}(\mathbf{P}_t(\mathbf{w}_{t-1} - \mathbf{w}_*))] + 2 \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\mathbf{v}_{t-1} \mathbf{H}(\mathbf{P}_t(\mathbf{w}_{t-1} - \mathbf{w}_*))] + 2 \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\mathbf{u}_{t-1} \mathbf{H} \mathbf{v}_{t-1}].$$

Notice that $\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\mathbf{u}_t] = 0$ and is independent, thus it holds that

$$\begin{aligned} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{w}_t - \mathbf{w}^*\|_{\mathbf{H}}^2] &= \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{P}_t(\mathbf{w}_{t-1} - \mathbf{w}_*)\|_{\mathbf{H}}^2] + \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\eta \mathbf{u}_{t-1}\|_{\mathbf{H}}^2] + \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\eta \mathbf{v}_{t-1}\|_{\mathbf{H}}^2] \\ &\quad + \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\eta \mathbf{u}_{t-1}^\top \mathbf{H}(\mathbf{P}_t(\mathbf{w}_{t-1} - \mathbf{w}_*))] + \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\eta (\mathbf{P}_t(\mathbf{w}_{t-1} - \mathbf{w}_*))^\top \mathbf{H} \mathbf{u}_{t-1}]. \end{aligned} \quad (15)$$

Recall that $\mathbf{u}_t = \frac{1}{b} \sum_{i=0}^{b-1} (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0]) \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_*$, it follows

$$\begin{aligned} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\eta \mathbf{u}_{t-1}\|_{\mathbf{H}}^2] &= \frac{\eta^2}{b^2} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [(\sum_{i=0}^{b-1} (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0]) \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_*)^\top \mathbf{H} \\ &\quad \cdot (\sum_{i=0}^{b-1} (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0]) \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_*)]. \end{aligned}$$

According to Lemma 10 (where each x is independent and symmetric), the following conditions hold when $i \neq j$:

$$\begin{aligned} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [(\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0]) \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \\ \cdot (\mathbb{1}[\mathbf{x}_{\tau(t)+m+j}^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_{\tau(t)+m+j}^\top \mathbf{w}_{t-1} > 0]) \mathbf{x}_{\tau(t)+m+j} \mathbf{x}_{\tau(t)+m+j}^\top] = 0. \end{aligned}$$

It implies that

$$\begin{aligned} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\eta \mathbf{u}_{t-1}\|_{\mathbf{H}}^2] &= \frac{\eta^2}{b^2} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\sum_{i=0}^{b-1} (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0])^2 \\ &\quad \cdot (\mathbf{w}_*^\top \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{H} \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_*)] \\ &= \frac{\eta^2}{b^2} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\sum_{i=0}^{b-1} (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0, \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* < 0] + \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} < 0, \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0]) \\ &\quad \cdot (\mathbf{w}_*^\top \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{H} \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_*)] \\ &= \frac{2\eta^2}{b^2} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\sum_{i=0}^{b-1} (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0, \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* < 0]) \\ &\quad \cdot (\mathbf{w}_*^\top \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{H} \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_*)], \end{aligned}$$

where we use the following fact in the second equality

$$\begin{aligned} &(\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0])^2 \\ &= (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0, \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* < 0] + \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} < 0, \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0]). \end{aligned}$$

Now we move on to the crossing term in Equation (15).

$$\begin{aligned} &\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\eta \mathbf{u}_{t-1}^\top \mathbf{H}(\mathbf{P}_t(\mathbf{w}_{t-1} - \mathbf{w}_*))] + \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\eta (\mathbf{P}_t(\mathbf{w}_{t-1} - \mathbf{w}_*))^\top \mathbf{H} \mathbf{u}_{t-1}] \\ &= \eta \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [(\frac{1}{b} \sum_{i=0}^{b-1} (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0]) \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_*)^\top \mathbf{H} \\ &\quad \cdot ((\mathbf{I} - \frac{\eta}{b} \sum_{i=0}^{b-1} \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0] \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top) (\mathbf{w}_{t-1} - \mathbf{w}_*)) \\ &\quad + ((\mathbf{I} - \frac{\eta}{b} \sum_{i=0}^{b-1} \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0] \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top) (\mathbf{w}_{t-1} - \mathbf{w}_*))^\top \mathbf{H} \end{aligned}$$

$$\cdot \left(\frac{1}{b} \sum_{i=0}^{b-1} (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0]) \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* \right).$$

Similarly, for any $i \neq j$, Lemma 10 holds that

$$\begin{aligned} & \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [(\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0]) \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \\ & \cdot \mathbb{1}[\mathbf{x}_{\tau(t)+m+j}^\top \mathbf{w}_{t-1} > 0] \mathbf{x}_{\tau(t)+m+j} \mathbf{x}_{\tau(t)+m+j}^\top] = 0. \end{aligned}$$

Therefore, the crossing term can be represented as

$$\begin{aligned} & = \eta \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} \left[\left(\frac{1}{b} \sum_{i=0}^{b-1} (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0]) \right) \right. \\ & \quad \cdot (\mathbf{w}_*^\top \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{H}(\mathbf{w}_{t-1} - \mathbf{w}_*) + (\mathbf{w}_{t-1} - \mathbf{w}_*)^\top \mathbf{H} \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_*) \\ & \quad - \frac{\eta}{b^2} \left(\left(\sum_{i=0}^{b-1} (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0]) \cdot (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0]) \right) \right. \\ & \quad \cdot 2(\mathbf{w}_*^\top \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{H} \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top (\mathbf{w}_{t-1} - \mathbf{w}_*)) \left. \right) \end{aligned}$$

Notice that

$$\begin{aligned} & - \left(\sum_{i=0}^{b-1} (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0]) \cdot (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0]) \right) \\ & = \sum_{i=0}^{b-1} (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0] - \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0] \cdot \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0]) \\ & = \sum_{i=0}^{b-1} \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0, \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0]. \end{aligned}$$

Combining the Lemma 10, the crossing term holds that

$$\begin{aligned} \text{crossing term} & = \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} \left[\frac{2\eta^2}{b^2} \sum_{i=0}^{b-1} (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0, \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0]) \right. \\ & \quad \cdot (\mathbf{w}_*^\top \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{H} \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top (\mathbf{w}_{t-1} - \mathbf{w}_*)) \left. \right]. \end{aligned}$$

Let us add $\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\eta \mathbf{u}_{t-1}\|_{\mathbf{H}}^2]$ and the crossing term together

$$\begin{aligned} & \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\eta \mathbf{u}_{t-1}\|_{\mathbf{H}}^2] + \text{crossing term} \\ & = \frac{2\eta^2}{b^2} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} \left[\sum_{i=0}^{b-1} (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0, \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* < 0]) \cdot (\mathbf{w}_*^\top \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{H} \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1}) \right]. \end{aligned}$$

It is clear that $\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\eta \mathbf{u}_{t-1}\|_{\mathbf{H}}^2] + \text{crossing term} \leq 0$. Moreover, according to Lemma 17 and Lemma 18, we have that

$$\begin{aligned} & \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{P}_t(\mathbf{w}_{t-1} - \mathbf{w}_*)\|_{\mathbf{H}}^2] + \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\eta \mathbf{v}_{t-1}\|_{\mathbf{H}}^2] \\ & \leq (1 - \eta\mu) \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{w}_{t-1} - \mathbf{w}_*\|_{\mathbf{H}}^2] + \frac{\eta^2}{b} \text{Tr}(\mathbf{H}\Sigma) - 0 + \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} \left[\eta^2 \frac{4s_{t-1}^2 f^2}{b^2} \right] \text{Tr}(\mathbf{H}). \end{aligned}$$

Therefore, the update of $\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{w}_t - \mathbf{w}_*\|_{\mathbf{H}}^2]$ in Equation (15) holds that

$$\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{w}_t - \mathbf{w}_*\|_{\mathbf{H}}^2] \leq (1 - \eta\mu) \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{w}_{t-1} - \mathbf{w}_*\|_{\mathbf{H}}^2] + \frac{\eta^2}{b} \text{Tr}(\mathbf{H}\Sigma) - 0 + \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} \left[\eta^2 \frac{4s_{t-1}^2 f^2}{b^2} \right] \text{Tr}(\mathbf{H}).$$

Considering the adaptive clipping algorithm, we have

$$s_t \leq R_x C_2 \log^{2a} N (\sqrt{\|\mathbf{H}\|} \|\mathbf{w}_t - \mathbf{w}_*\| + \|\mathbf{w}_*\| + \sigma + \Delta).$$

Similar to one sample case, the recursion of $\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{w}_t - \mathbf{w}^*\|_{\mathbf{H}}^2]$ will be

$$\begin{aligned} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{w}_t - \mathbf{w}^*\|_{\mathbf{H}}^2] &\leq (1 - \eta\mu) \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{w}_{t-1} - \mathbf{w}^*\|_{\mathbf{H}}^2] \\ &\quad + \frac{\eta^2 \sigma^2}{b} \text{Tr}(\mathbf{H}) + 16R_x^2 C_2^2 \log^{4a} N \eta^2 \frac{f^2}{b^2} \text{Tr}(\mathbf{H}) (\kappa \|\mathbf{w}_t - \mathbf{w}_*\|_{\mathbf{H}}^2 + \|\mathbf{w}_*\|^2 + \sigma^2 + \Delta^2) \\ &= (1 - (\eta\mu - 16\eta^2 \frac{f^2}{b^2} C_2^2 R_x^2 \kappa \log^{4a} N \text{Tr}(\mathbf{H}))) \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{w}_{t-1} - \mathbf{w}^*\|_{\mathbf{H}}^2] \\ &\quad + \frac{\eta^2 \sigma^2}{b} \text{Tr}(\mathbf{H}) + 16\eta^2 \frac{f^2}{b^2} C_2^2 R_x^2 \log^{4a} N (\sigma^2 + \Delta^2 + \|\mathbf{w}_*\|^2) \text{Tr}(\mathbf{H}). \end{aligned}$$

Notice that $T \cdot (b + m) = N$, thus if we have $\frac{\eta\mu}{2} \geq 16\eta^2 \frac{f^2}{b^2} C_2^2 R_x^2 \kappa \log^{4a} N \text{Tr}(\mathbf{H})$, it equals to

$$\left(\frac{N}{T} - m\right)^2 \geq \frac{32\eta f^2 C_2^2 R_x^2 \kappa \log^{4a} N \text{Tr}(\mathbf{H})}{\mu},$$

which implies that

$$\begin{aligned} &\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{w}_t - \mathbf{w}^*\|_{\mathbf{H}}^2] \\ &\leq \left(1 - \frac{\eta\mu}{2}\right) \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{w}_{t-1} - \mathbf{w}^*\|_{\mathbf{H}}^2] \frac{\eta^2 \sigma^2}{b} \text{Tr}(\mathbf{H}) + 16\eta^2 \frac{f^2}{b^2} C_2^2 R_x^2 \log^{4a} N (\sigma^2 + \Delta^2 + \|\mathbf{w}_*\|^2) \text{Tr}(\mathbf{H}) \\ &\leq (1 - \eta\mu/2)^t \|\mathbf{w}_0 - \mathbf{w}^*\|_{\mathbf{H}}^2 + \frac{2}{\eta\mu} \left(\frac{\eta^2 \sigma^2}{b} \text{Tr}(\mathbf{H}) + 16\eta^2 \frac{f^2}{b^2} C_2^2 R_x^2 \log^{4a} N (\sigma^2 + \Delta^2 + \|\mathbf{w}_*\|^2) \text{Tr}(\mathbf{H})\right) \\ &\leq e^{-\eta\mu t/2} \|\mathbf{w}^*\|_{\mathbf{H}}^2 + \frac{2}{\eta\mu} \left(\frac{\eta^2 \sigma^2}{b} \text{Tr}(\mathbf{H}) + 16\eta^2 \frac{f^2}{b^2} C_2^2 R_x^2 \log^{4a} N (\sigma^2 + \Delta^2 + \|\mathbf{w}_*\|^2) \text{Tr}(\mathbf{H})\right). \end{aligned}$$

Substituting the above to the $\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [s_t^2]$, similarly, we will have the following results

$$\begin{aligned} \Gamma^2 &= \max\{\text{Upper-Bound}(\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [s_0^2]), \dots, \text{Upper-Bound}(\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [s_T^2])\} \\ &= \text{Upper-Bound}(\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [s_0^2]) \\ &= 4C_2^2 R_x^2 \log^{4a} N (\kappa \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\|\mathbf{w}_0 - \mathbf{w}^*\|_{\mathbf{H}}^2] + \sigma^2 + \|\mathbf{w}_*\|^2 + \Delta^2). \end{aligned}$$

Before presenting the utility guarantee, we first need to redefine certain notations and properties.

We denote the recursion:

$$\begin{aligned} (\mathcal{I} - \mathcal{T}(\eta, b, \mathbf{H})) \circ \mathbf{A}_{t-1} &= \mathbf{A}_{t-1} - \frac{\eta}{2} (\mathbf{H}\mathbf{A}_{t-1} + \mathbf{A}_{t-1}\mathbf{H}) + \frac{\eta^2}{b} \left(\frac{1}{2}\mathcal{M} + (b-1)\frac{1}{4}\mathbf{H}^2\right) \circ \mathbf{A}_{t-1} \\ (\mathcal{I} - \tilde{\mathcal{T}}(\eta, b, \mathbf{H})) \circ \mathbf{A}_{t-1} &= \mathbf{A}_{t-1} - \frac{\eta}{2} (\mathbf{H}\mathbf{A}_{t-1} + \mathbf{A}_{t-1}\mathbf{H}) + \frac{\eta^2}{4} \mathbf{H}^2 \circ \mathbf{A}_t = (\mathbf{I} - \frac{\eta}{2}\mathbf{H})\mathbf{A}_t(\mathbf{I} - \frac{\eta}{2}\mathbf{H}), \end{aligned} \tag{16}$$

where $\mathcal{I} \circ \mathbf{A} = \mathbf{A}$, $\mathcal{M} \circ \mathbf{A} = \mathbb{E}[(\mathbf{x}^\top \mathbf{A} \mathbf{x}) \mathbf{x} \mathbf{x}^\top]$ and $\tilde{\mathcal{M}} \circ \mathbf{A} = \mathbf{H} \mathbf{A} \mathbf{H}$ for a symmetric matrix \mathbf{A} . For simplicity, we will use $(\mathcal{I} - \mathcal{T})$ and $(\mathcal{I} - \tilde{\mathcal{T}})$ in place of the complete notation.

It can be readily understood that the following properties are satisfied:

Lemma 19 (Zou et al. [2021]). *An operator \mathcal{O} , when defined on symmetric matrices, is termed a Positive Semi-Definite (PSD) mapping, if $\mathbf{A} \succeq 0$ implies $\mathcal{O} \circ \mathbf{A} \succeq 0$. Consequently, we have:*

1. \mathcal{M} and $\tilde{\mathcal{M}}$ are both PSD mappings.

2. $\mathcal{M} - \widetilde{\mathcal{M}}$ and $\widetilde{\mathcal{T}} - \mathcal{T}$ are both PSD mappings.
3. $\mathcal{I} - \eta\mathcal{T}$ and $\mathcal{I} - \eta\widetilde{\mathcal{T}}$ are both PSD mappings.
4. If $0 < \eta < 1/\lambda_1$, then $\widetilde{\mathcal{T}}^{-1}$ exists, and is a PSD mapping.
5. If $0 < \eta < 1/(\alpha \text{tr}(\mathbf{H}))$, then $\mathcal{T}^{-1} \circ \mathbf{A}$ exists for PSD matrix \mathbf{A} , and \mathcal{T}^{-1} is a PSD mapping.

Proof. The subsequent proofs are summarized from Jain et al. [2018], Zou et al. [2021], and are included herein for the sake of completeness.

1. For any PSD matrix $\mathbf{A} \succeq 0$, by definition, we have

$$\begin{aligned}\mathcal{M} \circ \mathbf{A} &= \mathbb{E}[\mathbf{xx}^\top \mathbf{A} \mathbf{xx}^\top] \succeq 0, \\ \widetilde{\mathcal{M}} \circ \mathbf{A} &= \mathbf{H} \mathbf{A} \mathbf{H} \succeq 0.\end{aligned}$$

2. For any PSD matrix $\mathbf{A} \succeq 0$,

$$(\mathcal{M} - \widetilde{\mathcal{M}}) \circ \mathbf{A} = \mathbb{E}[\mathbf{xx}^\top \mathbf{A} \mathbf{xx}^\top] - \mathbf{H} \mathbf{A} \mathbf{H} = \mathbb{E}[(\mathbf{xx}^\top - \mathbf{H}) \mathbf{A} (\mathbf{xx}^\top - \mathbf{H})] \succeq 0.$$

Also, we have $\widetilde{\mathcal{T}} - \mathcal{T} = \frac{\eta^2}{2b} \mathcal{M} - \frac{\eta^2}{4b} \widetilde{\mathcal{M}} \succeq 0$, which indicates $\mathcal{M} - \widetilde{\mathcal{M}}$ and $\widetilde{\mathcal{T}} - \mathcal{T}$ are both PSD mappings.

3. For any PSD matrix $\mathbf{A} \succeq 0$, we have

$$\begin{aligned}(\mathcal{I} - \eta\mathcal{T}) \circ \mathbf{A} &= (\mathbf{I} - \frac{\eta}{2} \mathbf{H}) \mathbf{A} (\mathbf{I} - \frac{\eta}{2} \mathbf{H}) + \frac{\eta^2}{2b} \mathcal{M} - \frac{\eta^2}{4b} \widetilde{\mathcal{M}} \succeq 0 \\ (\mathcal{I} - \eta\widetilde{\mathcal{T}}) \circ \mathbf{A} &= (\mathbf{I} - \frac{\eta}{2} \mathbf{H}) \mathbf{A} (\mathbf{I} - \frac{\eta}{2} \mathbf{H}) \succeq 0.\end{aligned}$$

4. The proof adheres to Lemma B.1 in Zou et al. [2021].
5. For any finite PSD matrix \mathbf{A} , we have:

$$\mathcal{T}^{-1} \circ \mathbf{A} = \eta \sum_{t=0}^{\infty} (\mathcal{I} - \eta\mathcal{T})^t \circ \mathbf{A}.$$

It is evident that if the right-hand side exists, it must be PSD, owing to the fact that $\mathcal{I} - \eta\mathcal{T}$ is a PSD mapping. Demonstrating that the trace of $\sum_{t=0}^{\infty} (\mathcal{I} - \eta\mathcal{T})^t \circ \mathbf{A}$ is finite would suffice to establish the conclusion.

$$\text{tr}(\sum_{t=0}^{\infty} (\mathcal{I} - \eta\mathcal{T})^t \circ \mathbf{A}) = \sum_{t=0}^{\infty} \text{tr}((\mathcal{I} - \eta\mathcal{T})^t \circ \mathbf{A}) = \sum_{t=0}^{\infty} \text{tr}(\mathbf{A}_t)$$

By Equation (16), we have:

$$\begin{aligned}\text{tr}(\mathbf{A}_t) &= \text{tr}(\mathbf{A}_{t-1}) - \eta \text{tr}(\mathbf{H} \mathbf{A}_{t-1}) + \frac{\eta^2}{2b} \text{tr}(\mathbb{E}[\mathbf{xx}^\top \mathbf{A}_{t-1} \mathbf{xx}^\top]) + \frac{b-1}{2} \text{tr}(\mathbf{H} \mathbf{A}_{t-1} \mathbf{H}) \\ &\leq \text{tr}(\mathbf{A}_{t-1}) - \eta \text{tr}(\mathbf{H} \mathbf{A}_{t-1}) + \frac{\eta^2}{2b} \text{tr}(\mathbf{A}_{t-1} \alpha \text{tr}(\mathbf{H}) \mathbf{H}) + \frac{b-1}{2} \text{tr}(\mathbf{H} \mathbf{A}_{t-1} \mathbf{H}) \\ &\leq \text{tr}(\mathbf{A}_{t-1}) - \eta(1 - \frac{\eta \alpha \text{tr}(\mathbf{H})}{2b} - \frac{\eta(b-1) \text{tr}(\mathbf{H})}{4b}) \text{tr}(\mathbf{H} \mathbf{A}_{t-1}) \\ &\leq \text{tr}(\mathbf{A}_{t-1}) - \frac{\eta}{2} \text{tr}(\mathbf{H} \mathbf{A}_{t-1}) \\ &\leq (1 - \frac{\eta}{2} \lambda_d) \text{tr}(\mathbf{A}_{t-1}),\end{aligned}$$

where we use $\eta \leq \frac{2b}{2\alpha \text{tr}(\mathbf{H}) + (b-1) \text{tr}(\mathbf{H})}$ in the penultimate inequality.

Hence, we have $\sum_{t=0}^{\infty} \text{tr}(\mathbf{A}_t) \leq \frac{2 \text{tr}(\mathbf{A})}{\eta \lambda_d} < \infty$, which complete the proofs.

□

Now we are ready to provide the evolution of \mathbf{A}_t .

Consider the gradient norm not exceeding the clipping norm:

$$\begin{aligned}
\mathbf{w}_t - \mathbf{w}_* &= \mathbf{w}_{t-1} - \frac{\eta}{b} \sum_{i=1}^b (\mathbf{x}_{\tau(t)+m+i}^\top (\text{ReLU}(\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_t) - y_{t,i})) - \frac{2\eta\Gamma f}{b} \mathbf{g}_t - \mathbf{w}_* \\
&= \mathbf{w}_{t-1} - \frac{\eta}{b} \sum_{i=1}^b (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0] \cdot \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} \\
&\quad - \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0] \cdot \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_*) + \frac{\eta}{b} \sum_{i=1}^b z_t \mathbf{x}_{\tau(t)+m+i} - \frac{2\eta\Gamma f}{b} \mathbf{g}_t \\
&= (\mathbf{I} - \frac{\eta}{b} \sum_{i=1}^b \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0] \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top) (\mathbf{w}_{t-1} - \mathbf{w}_*) \\
&\quad + \frac{\eta}{b} \sum_{i=1}^b (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0]) \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* + \frac{\eta}{b} \sum_{i=1}^b z_t \mathbf{x}_{\tau(t)+m+i} - \frac{2\eta\Gamma f}{b} \mathbf{g}_t.
\end{aligned}$$

Let's consider the expected outer product:

$$\begin{aligned}
\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} (\mathbf{w}_t - \mathbf{w}_*)^{\otimes 2} &= [(\text{quadratic term 1}) + (\text{quadratic term 2}) + (\text{quadratic term 3}) \\
&\quad + (\text{crossing term 1}) + (\text{crossing term 2})],
\end{aligned} \tag{17}$$

where

$$\begin{aligned}
(\text{quadratic term 1}) &= \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} (\mathbf{I} - \frac{\eta}{b} \sum_{i=1}^b \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0] \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top)^{\otimes 2} \circ (\mathbf{w}_{t-1} - \mathbf{w}_*)^{\otimes 2} \\
(\text{quadratic term 2}) &= (\frac{\eta}{b})^2 \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} \sum_{i=1}^b (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0]) \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* \\
&\quad \cdot \sum_{j=1}^b (\mathbb{1}[\mathbf{x}_{\tau(t)+m+j}^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_{\tau(t)+m+j}^\top \mathbf{w}_{t-1} > 0]) \mathbf{w}_*^\top \mathbf{x}_{\tau(t)+m+j} \mathbf{x}_{\tau(t)+m+j}^\top \\
(\text{quadratic term 3}) &= \frac{\eta^2}{b^2} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} \{ (\sum_{i=1}^b z_t \mathbf{x}_{\tau(t)+m+i} - 2\Gamma f \mathbf{g}_t) (\sum_{j=1}^b z_t \mathbf{x}_{\tau(t)+m+j} - 2\Gamma f \mathbf{g}_t)^\top \\
(\text{crossing term 1}) &= \frac{\eta}{b} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} \sum_{i=1}^b (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0]) \\
&\quad \cdot \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* (\mathbf{w}_{t-1} - \mathbf{w}_*)^\top (\mathbf{I} - \frac{\eta}{b} \sum_{j=1}^b \mathbb{1}[\mathbf{x}_{\tau(t)+m+j}^\top \mathbf{w}_{t-1} > 0] \mathbf{x}_{\tau(t)+m+j} \mathbf{x}_{\tau(t)+m+j}^\top) \\
(\text{crossing term 2}) &= \frac{\eta}{b} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} (\mathbf{I} - \frac{\eta}{b} \sum_{i=1}^b \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0] \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top) (\mathbf{w}_{t-1} - \mathbf{w}_*)^\top \\
&\quad \cdot \frac{\eta}{b} \sum_{j=1}^b (\mathbb{1}[\mathbf{x}_{\tau(t)+m+j}^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_{\tau(t)+m+j}^\top \mathbf{w}_{t-1} > 0]) \mathbf{x}_{\tau(t)+m+j} \mathbf{x}_{\tau(t)+m+j}^\top \mathbf{w}_*.
\end{aligned} \tag{18}$$

We will consider the above separately.

According to Assumption 3 (where each x is independent and symmetric), the following conditions hold when $i \neq j$:

$$\begin{aligned}
& \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0]) \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \\
& \cdot (\mathbb{1}[\mathbf{x}_{\tau(t)+m+j}^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_{\tau(t)+m+j}^\top \mathbf{w}_{t-1} > 0]) \mathbf{x}_{\tau(t)+m+j} \mathbf{x}_{\tau(t)+m+j}^\top = 0 \\
& \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0]) \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \\
& \cdot \mathbb{1}[\mathbf{x}_{\tau(t)+m+j}^\top \mathbf{w}_{t-1} > 0] \mathbf{x}_{\tau(t)+m+j} \mathbf{x}_{\tau(t)+m+j}^\top = 0.
\end{aligned} \tag{19}$$

Substituting the above into the quadratic term 2 and the crossing terms, we will obtain

$$\begin{aligned}
& \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} (\text{quadratic term 2}) \\
& = \left(\frac{\eta}{b}\right)^2 \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} \left(\sum_{i=1}^b (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0])^2 \cdot (\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_*)^2 \cdot \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \right) \\
& = \left(\frac{\eta}{b}\right)^2 \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} \left(\sum_{i=1}^b (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0, \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* < 0] + \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} < 0, \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0]) \right. \\
& \quad \cdot (\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_*)^2 \cdot \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \Big) \\
& = 2 \left(\frac{\eta}{b}\right)^2 \cdot \mathbb{E} \left(\sum_{i=1}^b \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0, \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* < 0] \cdot (\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_*)^2 \cdot \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \right),
\end{aligned} \tag{20}$$

where the first equality comes from

$$(\mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0])^2 = \mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0, \mathbf{x}_t^\top \mathbf{w}_* < 0] + \mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} < 0, \mathbf{x}_t^\top \mathbf{w}_* > 0].$$

For the crossing terms, we know

$$\begin{aligned}
& (\text{crossing term 1}) + (\text{crossing term 2}) \\
& = \frac{2\eta}{b} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} \left[\sum_{i=1}^b (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0]) \cdot (\mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* (\mathbf{w}_{t-1} - \mathbf{w}_*)^\top \right. \\
& \quad \left. + (\mathbf{w}_{t-1} - \mathbf{w}_*) \mathbf{w}_*^\top \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top) \right] \\
& - \frac{2\eta^2}{b^2} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} \left[\sum_{i=1}^b (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0]) \cdot \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0] \right. \\
& \quad \left. \cdot \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* \cdot \mathbf{x}_{\tau(t)+m+i}^\top (\mathbf{w}_{t-1} - \mathbf{w}_*) \cdot \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \right] \\
& = \frac{2\eta^2}{b^2} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} \left[\sum_{i=1}^b \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0, \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* < 0] \cdot \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* \cdot \mathbf{x}_{\tau(t)+m+i}^\top (\mathbf{w}_{t-1} - \mathbf{w}_*) \cdot \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \right].
\end{aligned} \tag{21}$$

Now, we turn our attention to the first quadratic term

$$\begin{aligned}
(\text{quadratic term 1}) & = \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} (\mathbf{w}_{t-1} - \mathbf{w}_*)^{\otimes 2} - \frac{\eta}{b} \left(\sum_{i=1}^b \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0] \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \right) \circ (\mathbf{w}_{t-1} - \mathbf{w}_*)^{\otimes 2} \\
& - \frac{\eta}{b} (\mathbf{w}_{t-1} - \mathbf{w}_*)^{\otimes 2} \circ \left(\sum_{i=1}^b \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0] \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \right) \\
& + \frac{\eta^2}{b^2} \left(\sum_{i=1}^b \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0] \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \right) \\
& \cdot \left(\sum_{j=1}^b \mathbb{1}[\mathbf{x}_{\tau(t)+m+j}^\top \mathbf{w}_{t-1} > 0] \mathbf{x}_{\tau(t)+m+j} \mathbf{x}_{\tau(t)+m+j}^\top \right) \circ (\mathbf{w}_{t-1} - \mathbf{w}_*)^{\otimes 2}
\end{aligned}$$

$$\begin{aligned}
&= \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} (\mathbf{w}_{t-1} - \mathbf{w}_*)^{\otimes 2} - \frac{\eta}{2} \mathbf{H} \circ (\mathbf{w}_{t-1} - \mathbf{w}_*)^{\otimes 2} - \frac{\eta}{2} (\mathbf{w}_{t-1} - \mathbf{w}_*)^{\otimes 2} \circ \mathbf{H} \\
&+ \frac{\eta^2}{b} \left(\frac{1}{2} \mathcal{M} + (b-1) \frac{1}{4} \mathbf{H}^2 \right) (\mathbf{w}_{t-1} - \mathbf{w}_*)^{\otimes 2}.
\end{aligned}$$

Applying Equation (21)) and Equation (20) to the expected outer product, we have the following recursion

$$\begin{aligned}
\mathbf{A}_t &= \mathbf{A}_{t-1} - \frac{\eta}{2} (\mathbf{H} \mathbf{A}_{t-1} + \mathbf{A}_{t-1} \mathbf{H}) + \frac{\eta^2}{b} \left(\frac{1}{2} \mathcal{M} + (b-1) \frac{1}{4} \mathbf{H}^2 \right) \circ \mathbf{A}_{t-1} + \frac{\eta^2}{b} \sigma^2 \mathbf{H} + \frac{4\eta^2 \Gamma^2 f^2}{b^2} \mathbf{I} \\
&+ \frac{2\eta^2}{b^2} \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} \left[\sum_{i=1}^b \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0, \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* < 0] \cdot \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* \cdot \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} \cdot \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{x}_{\tau(t)+m+i}^\top \right] \\
&\preceq \mathbf{A}_{t-1} - \frac{\eta}{2} (\mathbf{H} \mathbf{A}_{t-1} + \mathbf{A}_{t-1} \mathbf{H}) + \frac{\eta^2}{b} \left(\frac{1}{2} \mathcal{M} + (b-1) \frac{1}{4} \mathbf{H}^2 \right) \circ \mathbf{A}_{t-1} + \frac{\eta^2}{b} \sigma^2 \mathbf{H} + \frac{4\eta^2 \Gamma^2 f^2}{b^2} \mathbf{I}.
\end{aligned}$$

The indicator function shows $\mathbb{1}[\mathbf{x}_t^\top \mathbf{w}_{t-1} > 0, \mathbf{x}_t^\top \mathbf{w}_* < 0] \cdot \mathbf{x}_t^\top \mathbf{w}_* \cdot \mathbf{x}_t^\top \mathbf{w}_{t-1} \leq 0$ in the last inequation.

Consequently, analogous to the one-sample case, we can decompose \mathbf{A}_t as follows:

$$\begin{cases} \mathbf{B}_t \preceq (\mathbf{I} - \eta \mathcal{T}(\eta, b, \mathbf{H})) \circ \mathbf{B}_{t-1}; \\ \mathbf{B}_0 = (\mathbf{w}_0 - \mathbf{w}_*)^{\otimes 2} \end{cases} \quad \begin{cases} \mathbf{C}_t \preceq (\mathbf{I} - \eta \mathcal{T}(\eta, b, \mathbf{H})) \circ \mathbf{C}_{t-1} + \frac{\eta^2}{b} \sigma^2 \mathbf{H} + \frac{4\eta^2 \Gamma^2 f^2}{b^2} \mathbf{I}; \\ \mathbf{C}_0 = 0 \end{cases} \quad (22)$$

Lemma 20. Suppose that Assumption 2 and Assumption 3 hold. For $\overline{\mathbf{w}}_N$ defined by previously, we have that

$$\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} \langle \mathbf{H}, (\overline{\mathbf{w}}_{s+1, T} - \mathbf{w}_*)^{\otimes 2} \rangle \leq \frac{1}{N^2} \cdot \sum_{t=s+1}^{s+T} \sum_{k=t}^{s+T} \langle (\mathbf{I} - \frac{\eta}{2} \mathbf{H})^{k-t} \mathbf{H}, \mathbf{A}_t \rangle.$$

Proof. We first focus on the expectation of $\mathbf{w}_t - \mathbf{w}_*$

$$\begin{aligned}
\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\mathbf{w}_t - \mathbf{w}_* \mid \mathbf{w}_{t-1}] &= \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} \left[\left(\mathbf{I} - \frac{\eta}{b} \sum_{i=1}^b \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0] \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \right) (\mathbf{w}_{t-1} - \mathbf{w}_*) \mid \mathbf{w}_{t-1} \right] \\
&+ \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} \left[\frac{\eta}{b} \sum_{i=1}^b (\mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* > 0] - \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0]) \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_* \mid \mathbf{w}_{t-1} \right] \\
&+ \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} \left[\frac{\eta}{b} \sum_{i=1}^b \varepsilon_t \mathbf{x}_{\tau(t)+m+i} - \frac{2\eta s_t f}{b} \mathbf{g}_t \mid \mathbf{w}_{t-1} \right] \\
&= \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} \left[\left(\mathbf{I} - \frac{\eta}{b} \sum_{i=1}^b \mathbb{1}[\mathbf{x}_{\tau(t)+m+i}^\top \mathbf{w}_{t-1} > 0] \mathbf{x}_{\tau(t)+m+i} \mathbf{x}_{\tau(t)+m+i}^\top \right) (\mathbf{w}_{t-1} - \mathbf{w}_*) \mid \mathbf{w}_{t-1} \right] \\
&= (\mathbf{I} - \frac{\eta}{2} \mathbf{H}) (\mathbf{w}_{t-1} - \mathbf{w}_*).
\end{aligned}$$

Applying the aforementioned recursively, we deduce that, for $t \geq s$, $\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [\mathbf{w}_t - \mathbf{w}_* \mid \mathbf{w}_s] = (\mathbf{I} - \frac{\eta}{2} \mathbf{H})^{t-s} (\mathbf{w}_s - \mathbf{w}_*)$, which also implies that

$$\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}} [(\mathbf{w}_t - \mathbf{w}_*) \otimes (\mathbf{w}_s - \mathbf{w}_*)] = (\mathbf{I} - \frac{\eta}{2} \mathbf{H})^{t-s} \cdot \mathbb{E}(\mathbf{w}_s - \mathbf{w}_*)^{\otimes 2} = (\mathbf{I} - \frac{\eta}{2} \mathbf{H})^{t-s} \cdot \mathbf{A}_s.$$

Then, we consider the tail-averaged mini-batch SGD algorithm and we denote $\bar{\mathbf{w}}_{s+1,T} - \mathbf{w}_* = \frac{1}{T} \sum_{t=s+1}^{s+T} \mathbf{w}_t - \mathbf{w}_*$:

$$\begin{aligned}
\mathbb{E}_{(\mathbf{x},y) \sim \mathcal{D}}[(\bar{\mathbf{w}}_{s+1,T} - \mathbf{w}_*)^{\otimes 2}] &= \frac{1}{T^2} \sum_{t=s+1}^{s+T} \sum_{k=s+1}^{s+T} \mathbb{E}_{(\mathbf{x},y) \sim \mathcal{D}}[(\mathbf{w}_t - \mathbf{w}_*) \otimes (\mathbf{w}_k - \mathbf{w}_*)] \\
&= \frac{1}{T^2} \cdot \left(\sum_{t \geq k} \mathbb{E}_{(\mathbf{x},y) \sim \mathcal{D}}[(\mathbf{w}_t - \mathbf{w}_*) \otimes (\mathbf{w}_k - \mathbf{w}_*)] + \sum_{t \leq k} \mathbb{E}[(\mathbf{w}_t - \mathbf{w}_*) \otimes (\mathbf{w}_k - \mathbf{w}_*)] \right) \\
&\preceq \frac{1}{T^2} \cdot \left(\sum_{t \geq k} \mathbb{E}_{(\mathbf{x},y) \sim \mathcal{D}}[(\mathbf{w}_t - \mathbf{w}_*) \otimes (\mathbf{w}_k - \mathbf{w}_*)] + \sum_{t \leq k} \mathbb{E}[(\mathbf{w}_t - \mathbf{w}_*) \otimes (\mathbf{w}_k - \mathbf{w}_*)] \right) \\
&= \frac{1}{T^2} \cdot \left(\sum_{t \geq k} (\mathbf{I} - \frac{\eta}{2} \mathbf{H})^{t-k} \mathbf{A}_k + \sum_{t \leq k} \mathbf{A}_t (\mathbf{I} - \frac{\eta}{2} \mathbf{H})^{k-t} \right) \\
&= \frac{1}{T^2} \cdot \sum_{t \leq k} (\mathbf{A}_t (\mathbf{I} - \frac{\eta}{2} \mathbf{H})^{k-t} + (\mathbf{I} - \frac{\eta}{2} \mathbf{H})^{k-t} \mathbf{A}_t) \\
&= \frac{1}{T^2} \cdot \sum_{t=s+1}^{s+T} \sum_{k=t}^{s+T} (\mathbf{A}_t (\mathbf{I} - \frac{\eta}{2} \mathbf{H})^{k-t} + (\mathbf{I} - \frac{\eta}{2} \mathbf{H})^{k-t} \mathbf{A}_t).
\end{aligned} \tag{23}$$

Then we consider the excess risk of tail-averaged mini-batch:

$$\begin{aligned}
\mathbb{E}_{(\mathbf{x},y) \sim \mathcal{D}} \langle \mathbf{H}, (\bar{\mathbf{w}}_{s+1,T} - \mathbf{w}_*)^{\otimes 2} \rangle &\leq \langle \mathbf{H}, \frac{1}{T^2} \cdot \sum_{t=s+1}^{s+T} \sum_{k=t}^{s+T} (\mathbf{A}_t (\mathbf{I} - \frac{\eta}{2} \mathbf{H})^{k-t} + (\mathbf{I} - \frac{\eta}{2} \mathbf{H})^{k-t} \mathbf{A}_t) \rangle \\
&= \frac{1}{T^2} \cdot \sum_{t=s+1}^{s+T} \sum_{k=t}^{s+T} \langle \mathbf{H}, \mathbf{A}_t (\mathbf{I} - \frac{\eta}{2} \mathbf{H})^{k-t} \rangle + \frac{1}{T^2} \cdot \sum_{t=s+1}^{s+T} \sum_{k=t}^{s+T} \langle \mathbf{H}, (\mathbf{I} - \frac{\eta}{2} \mathbf{H})^{k-t} \mathbf{A}_t \rangle \\
&\leq \frac{1}{\eta T^2} \cdot \sum_{t=s+1}^{s+T} \langle \mathbf{I} - (\mathbf{I} - \frac{\eta}{2} \mathbf{H})^T, \mathbf{A}_t \rangle.
\end{aligned}$$

□

Variance Error

Lemma 21. Suppose Assumptions hold. Suppose $\eta < \min\{\frac{1}{R_x^2}, \frac{4b}{2R_x^2 + (b-1)\|\mathbf{H}\|_2}\}$. Then for every t we have

$$\mathbf{C}_t \preceq \frac{4\eta\sigma^2}{4b - 2\eta R_x^2 - \eta(b-1)\|\mathbf{H}\|_2} \mathbf{I} + \frac{16\eta\Gamma^2 f^2}{b(4b - 2\eta R_x^2 - \eta(b-1)\|\mathbf{H}\|_2)} \mathbf{H}^{-1}.$$

Proof. We proceed with induction.

For $t = 0$ we have $\mathbf{C}_0 = 0 \preceq \frac{4\eta\sigma^2}{4b - 2\eta R_x^2 - \eta(b-1)\|\mathbf{H}\|_2} \mathbf{I} + \frac{16\eta\Gamma^2 f^2}{b(4b - 2\eta R_x^2 - \eta(b-1)\|\mathbf{H}\|_2)} \mathbf{H}^{-1}$.

We then assume that \mathbf{C}_{t-1} holds for Lemma 21, and exam \mathbf{C}_t based on Equation (22)

$$\begin{aligned}
\mathbf{C}_t &\preceq (\mathbf{I} - \eta\mathcal{T}(\eta, b, \mathbf{H})) \circ \mathbf{C}_{t-1} + \frac{\eta^2}{b} \sigma^2 \mathbf{H} + \frac{4\eta^2 \Gamma^2 f^2}{b^2} \mathbf{I} \\
&= \mathbf{C}_{t-1} - \frac{\eta}{2} (\mathbf{H} \mathbf{C}_{t-1} + \mathbf{C}_{t-1} \mathbf{H}) + \frac{\eta^2}{b} (\frac{1}{2} \mathcal{M} + (b-1) \frac{1}{4} \mathbf{H}^2) \circ \mathbf{C}_{t-1} + \frac{\eta^2 \sigma^2}{b} \mathbf{H} + \frac{4\eta^2 \Gamma^2 f^2}{b^2} \mathbf{I} \\
&\preceq \frac{4\eta\sigma^2}{4b - 2\eta R_x^2 - \eta(b-1)\|\mathbf{H}\|_2} \mathbf{I} + \frac{16\eta\Gamma^2 f^2}{b(4b - 2\eta R_x^2 - \eta(b-1)\|\mathbf{H}\|_2)} \mathbf{H}^{-1} \\
&\quad - \eta \left(\frac{4\eta\sigma^2}{4b - 2\eta R_x^2 - \eta(b-1)\|\mathbf{H}\|_2} \mathbf{H} + \frac{16\eta\Gamma^2 f^2}{b(4b - 2\eta R_x^2 - \eta(b-1)\|\mathbf{H}\|_2)} \mathbf{I} \right) \\
&\quad + \left(\frac{\eta^2 R_x^2}{2b} + \frac{\eta^2 (b-1)\|\mathbf{H}\|_2}{4b} \right) \left(\frac{4\eta\sigma^2}{4b - 2\eta R_x^2 - \eta(b-1)\|\mathbf{H}\|_2} \mathbf{H} + \frac{16\eta\Gamma^2 f^2}{b(4b - 2\eta R_x^2 - \eta(b-1)\|\mathbf{H}\|_2)} \mathbf{I} \right)
\end{aligned}$$

$$\begin{aligned}
& + \frac{\eta^2 \sigma^2}{b} \mathbf{H} + \frac{4\eta^2 \Gamma^2 f^2}{b^2} \mathbf{I} \\
& \preceq \frac{4\eta \sigma^2}{4b - 2\eta R_x^2 - \eta(b-1)\|\mathbf{H}\|_2} \mathbf{I} + \frac{16\eta \Gamma^2 f^2}{b(4b - 2\eta R_x^2 - \eta(b-1)\|\mathbf{H}\|_2)} \mathbf{H}^{-1}.
\end{aligned}$$

□

For the simplicity, we define $\Sigma := \frac{\sigma^2}{b} \mathbf{H} + \frac{4\Gamma^2 f^2}{b^2} \mathbf{I}$ and $\mu_b := (4b - 2\eta R_x^2 - \eta(b-1)\|\mathbf{H}\|_2)$. By the definitions of \mathcal{T} and $\tilde{\mathcal{T}}$, we have

$$\begin{aligned}
\mathbf{C}_t &= (\mathbf{I} - \eta \mathcal{T}(\eta, b, \mathbf{H})) \circ \mathbf{C}_{t-1} + \eta^2 \Sigma \\
&= (\mathbf{I} - \frac{\eta}{2} \mathbf{H}) \mathbf{C}_{t-1} (\mathbf{I} - \frac{\eta}{2} \mathbf{H}) + (\frac{\eta^2}{2b} \mathcal{M} - \frac{\eta^2}{4b} \mathbf{H} \mathbf{H}) \circ \mathbf{C}_{t-1} + \eta^2 \Sigma \\
&\preceq (\mathbf{I} - \frac{\eta}{2} \mathbf{H}) \mathbf{C}_{t-1} (\mathbf{I} - \frac{\eta}{2} \mathbf{H}) + \frac{\eta^2 R_x^2}{2b} \mathbf{H} \circ \mathbf{C}_{t-1} + \eta^2 \Sigma.
\end{aligned} \tag{24}$$

Then by Lemma 21, we have for all $t \geq 0$,

$$\mathbf{H} \circ \mathbf{C}_t \preceq \mathbf{H} \circ (\frac{4\eta \sigma^2}{\mu_b} \mathbf{I} + \frac{16\eta \Gamma^2 f^2}{b\mu_b} \mathbf{H}^{-1}) \preceq \frac{4\eta \sigma^2}{\mu_b} \mathbf{H} + \frac{16\eta \Gamma^2 f^2}{b\mu_b} \mathbf{I}.$$

Substituting the above into Equation (24), it holds that

$$\begin{aligned}
\mathbf{C}_t &\preceq (\mathbf{I} - \frac{\eta}{2} \mathbf{H}) \mathbf{C}_{t-1} (\mathbf{I} - \frac{\eta}{2} \mathbf{H}) + \frac{\eta^2 R_x^2}{2b} (\frac{4\eta \sigma^2}{\mu_b} \mathbf{H} + \frac{16\eta \Gamma^2 f^2}{b\mu_b} \mathbf{I}) + \eta^2 (\frac{\sigma^2}{b} \mathbf{H} + \frac{4\Gamma^2 f^2}{b^2} \mathbf{I}) \\
&\preceq (\mathbf{I} - \frac{\eta}{2} \mathbf{H}) \mathbf{C}_{t-1} (\mathbf{I} - \frac{\eta}{2} \mathbf{H}) + \frac{\eta^2 \sigma^2 (\mu_b - 2\eta R_x^2)}{b\mu_b} \mathbf{H} + \frac{4\eta^2 \Gamma^2 f^2 (\mu_b - 2\eta R_x^2)}{b^2 \mu_b} \mathbf{I},
\end{aligned}$$

which implies that

$$\begin{aligned}
\mathbf{C}_t &\preceq \frac{\eta^2 (\mu_b - 2\eta R_x^2)}{b\mu_b} \cdot \sum_{k=0}^{t-1} (\mathbf{I} - \frac{\eta}{2} \mathbf{H})^k (\sigma^2 \mathbf{H} + \frac{4\Gamma^2 f^2}{b} \mathbf{I}) (\mathbf{I} - \frac{\eta}{2} \mathbf{H})^k \\
&\preceq \frac{\eta^2 (\mu_b - 2\eta R_x^2)}{b\mu_b} \cdot \sum_{k=0}^{t-1} (\mathbf{I} - \frac{\eta}{2} \mathbf{H})^k (\sigma^2 \mathbf{H} + \frac{4\Gamma^2 f^2}{b} \mathbf{I}) \\
&= \frac{\eta \sigma^2 (\mu_b - 2\eta R_x^2)}{b\mu_b} \cdot (\mathbf{I} - (\mathbf{I} - \frac{\eta}{2} \mathbf{H})^t) + \frac{4\eta \Gamma^2 f^2 (\mu_b - 2\eta R_x^2)}{b^2 \mu_b} \cdot (\mathbf{I} - (\mathbf{I} - \frac{\eta}{2} \mathbf{H})^t) \cdot \mathbf{H}^{-1}.
\end{aligned}$$

Consequently, the variance error can be represented as follows, in accordance with Lemma 20:

$$\begin{aligned}
\text{variance error} &= \frac{1}{\eta T^2} \cdot \sum_{t=s+1}^{s+T} \langle \mathbf{I} - (\mathbf{I} - \frac{\eta}{2} \mathbf{H})^T, \mathbf{C}_t \rangle \\
&\leq \frac{1}{\eta T^2} \cdot \frac{\eta \sigma^2 (\mu_b - 2\eta R_x^2)}{b\mu_b} \cdot \sum_{t=s+1}^{s+T} \langle \mathbf{I} - (\mathbf{I} - \frac{\eta}{2} \mathbf{H})^T, (\mathbf{I} - (\mathbf{I} - \frac{\eta}{2} \mathbf{H})^t) \rangle \\
&\quad + \frac{1}{\eta T^2} \cdot \frac{4\eta \Gamma^2 f^2 (\mu_b - 2\eta R_x^2)}{b^2 \mu_b} \cdot \sum_{t=s+1}^{s+T} \langle \mathbf{I} - (\mathbf{I} - \frac{\eta}{2} \mathbf{H})^T, (\mathbf{I} - (\mathbf{I} - \frac{\eta}{2} \mathbf{H})^t) \mathbf{H}^{-1} \rangle \\
&\lesssim \frac{\sigma^2 d}{Tb} + \frac{\Gamma^2 f^2 \text{tr}(\mathbf{H})}{Tb^2} \lesssim \frac{\sigma^2 d}{N} + \frac{d^2 \log N \log(1/\delta)}{N^2 \varepsilon^2} \cdot C_2^2 \kappa^2 (\sigma^2 + \|\mathbf{w}_*\|_{\mathbf{H}}^2 + \Delta^2).
\end{aligned} \tag{25}$$

Bias Error According to Lemma 20, the bias error of tail average iterate follows that

$$\text{bias error} \leq \frac{1}{\eta T^2} \cdot \sum_{t=s+1}^{s+T} \langle \mathbf{I} - (\mathbf{I} - \frac{\eta}{2} \mathbf{H})^T, \mathbf{B}_t \rangle \leq \sum_{t=s+1}^{s+T} \frac{1}{\eta T^2} \text{tr}(\mathbf{B}_t).$$

Considering the recursion of \mathbf{B}_t , we have

$$\begin{aligned}\mathbf{B}_t &\preceq \mathbf{B}_{t-1} - \frac{\eta}{2}(\mathbf{H}\mathbf{B}_{t-1} + \mathbf{B}_{t-1}\mathbf{H}) + \frac{\eta^2}{b}(\frac{1}{2}\mathcal{M} + (b-1)\frac{1}{4}\mathbf{H}^2) \circ \mathbf{B}_{t-1} \\ &\preceq \mathbf{B}_{t-1} - \eta\mathbf{H}\mathbf{B}_{t-1} + \frac{\eta^2}{b}(R_x^2 + (b-1)\|\mathbf{H}\|_2)\mathbf{B}_{t-1} \\ &\preceq (\mathbf{I} - \frac{\eta}{2}\mathbf{H})\mathbf{B}_{t-1}.\end{aligned}$$

The last inequality derives from the choice of step size. Consequently, the bias error will be

$$\text{bias error} \leq \sum_{t=0}^N \frac{1}{\eta N^2} \text{tr}(\mathbf{B}_t) \leq \sum_{t=0}^N \frac{1}{\eta N^2} (1 - \frac{\eta\mu}{2})^t \text{tr}(\mathbf{B}_0) \lesssim \frac{1}{\eta N} \|\mathbf{w}_*\|_{\mathbf{H}}^2 \lesssim \frac{d \log^2(N/\delta)}{N^2 \varepsilon^2} \cdot C_2^2 \kappa^2 \|\mathbf{w}_*\|_{\mathbf{H}}^2.$$

Combining the previous variance error, we complete the proof.

E LOWER BOUND

Proof of Lemma 6. We will denote $T_i = \mathcal{T}_{\mathbf{w}}((\mathbf{x}_i, y_i), M(D))$ and $T'_i = \mathcal{T}_{\mathbf{w}}((\mathbf{x}_i, y_i), M(D'_i))$. Since we have $y - \text{ReLU}(\mathbf{w}^\top \mathbf{x}) = z$ and $\mathbf{x} \cdot \mathbb{1}(\mathbf{w}^\top \mathbf{x} > 0)$, and $M(D'_i) - \mathbf{w}$ are independent, we have

$$\mathbb{E}[T'_i] = \mathbb{E}(y - \text{ReLU}(\mathbf{w}^\top \mathbf{x}))\mathbb{E}\langle M(D'_i) - \mathbf{w}, \mathbf{x} \cdot \mathbb{1}(\mathbf{w}^\top \mathbf{x} > 0) \rangle = 0. \quad (26)$$

Moreover, we have

$$\mathbb{E}[T'_i] \leq \sqrt{\mathbb{E}[T'^2_i]} \leq \sigma \sqrt{\mathbb{E}\|M(D'_i) - \mathbf{w}\|_{\Sigma_{\mathbf{x}}}^2} = \sigma \sqrt{\mathbb{E}\|M(D) - \mathbf{w}\|_{\Sigma_{\mathbf{x}}}^2}. \quad (27)$$

For the second part, we have

$$\sum_{i \in [n]} \mathbb{E}[T_i] = \sum_{j=1}^d \mathbb{E}M(D)_j \sum_{i=1}^N (y_i - \text{ReLU}(\mathbf{w}^\top \mathbf{x}_i)) \mathbb{1}(\mathbf{w}^\top \mathbf{x} > 0) \mathbf{x}_{i,j}$$

For each j we have

$$\mathbb{E}M(D)_j \sum_{i=1}^N (y_i - \text{ReLU}(\mathbf{w}^\top \mathbf{x}_i)) \mathbb{1}(\mathbf{w}^\top \mathbf{x} > 0) \mathbf{x}_{i,j} = \sigma^2 \mathbb{E}M(D)_j \frac{\partial \log f_{\mathbf{w}}(Y|X)}{\partial \mathbf{w}_j} = \sigma^2 \frac{\partial}{\partial \mathbf{w}_j} \mathbb{E}_{Y,X|\mathbf{w}} M(D)_j.$$

Thus we have

$$\sum_{i \in [n]} \mathbb{E}[T_i] = \sum_{j=1}^d \sigma^2 \frac{\partial}{\partial \mathbf{w}_j} \mathbb{E}_{Y,X|\mathbf{w}} M(D)_j.$$

Recall the following Stein's lemma:

Lemma 22. *Let Z be distributed according to some density $p(z)$ that is continuously differentiable w.r.t. z and let h be a differentiable function such that $\mathbb{E}|h'(Z)| < \infty$. We have*

$$\mathbb{E}[h'(Z)] = \mathbb{E}\left[-\frac{h(Z)p'(Z)}{p(Z)}\right].$$

Consider the following prior distribution π for \mathbf{w} : let v_1, \dots, v_d be i.i.d. sampled from the truncated $\mathcal{N}(0, 1)$ with truncation at -1 and 1 , and let $\mathbf{w}_j = \frac{v_j}{\sqrt{d}}$ thus $\mathbf{w} \in \mathcal{W}$. Denote $g_j(\mathbf{w}) = \mathbb{E}_{Y,X|\mathbf{w}} M(D)_j$. For each $j \in [d]$ by using the above lemma we have

$$\mathbb{E}_{\pi} \frac{\partial}{\partial \mathbf{w}_j} g_j(\mathbf{w}) = \mathbb{E}_{\pi} \frac{\partial}{\partial \mathbf{w}_j} \mathbb{E}(g_j(\mathbf{w})|\mathbf{w}_j) \geq \mathbb{E}_{\pi} \left(\frac{-\mathbf{w}_j \pi'_j(\mathbf{w}_j)}{\pi_j(\mathbf{w}_j)} - \mathbb{E}(|g_j(\mathbf{w}) - \mathbf{w}_j| | \frac{\pi'_j(\mathbf{w}_j)}{\pi_j(\mathbf{w}_j)}) \right).$$

Since π_j is a truncated normal distribution, we can easily get $\frac{\pi'_j(\mathbf{w}_j)}{\pi_j(\mathbf{w}_j)} = -d\mathbf{w}_j$. Therefore, it holds that

$$\begin{aligned} \mathbb{E}_\pi \sum_{j=1}^d \left(\frac{-\mathbf{w}_j \pi'_j(\mathbf{w}_j)}{\pi_j(\mathbf{w}_j)} - \mathbb{E}(|g_j(\mathbf{w}) - \mathbf{w}_j| \frac{\pi'_j(\mathbf{w}_j)}{\pi_j(\mathbf{w}_j)}) \right) &= \mathbb{E}_\pi [d \sum_{j=1}^d \mathbf{w}_j^2] - \sum_{j=1}^d \mathbb{E}_\pi (|g_j(\mathbf{w}) - \mathbf{w}_j| d |\mathbf{w}_j|) \\ &\geq d \{ \mathbb{E}_\pi [\sum_{j=1}^d \beta_j^2] - \sqrt{\mathbb{E}_\pi \mathbb{E}_{Y,X|\mathbf{w}} \|M(D) - \mathbf{w}\|_2^2} \sqrt{\mathbb{E}_\pi \sum_{j=1}^d \mathbf{w}_j^2} \}. \end{aligned}$$

As $\mathbb{E}_\pi [\sum_{j=1}^d \mathbf{w}_j^2] = \mathcal{W}(1)$, in total we have

$$\sum_{i \in [n]} \mathbb{E}[T_i] \geq O(\sigma^2 d \{1 - \sqrt{\mathbb{E}_\pi \mathbb{E}_{Y,X|\mathbf{w}} \|M(D) - \mathbf{w}\|_2^2}\})$$

We have the proof under the assumption that $\mathbb{E}_\pi \mathbb{E}_{Y,X|\mathbf{w}} \|M(D) - \mathbf{w}\|_2^2 = o(1)$. \square

Proof. We first prove the following lemma, whose proof is the same as the proof of Lemma B.2 in Cai et al. [2021].

Lemma 23. *For all $i \in [n]$, if M is (ε, δ) -DP then for every $T > 0$*

$$\mathbb{E}[T_i] \leq \mathbb{E}[T'_i] + 2\varepsilon \mathbb{E}[|T'_i|] + 2\delta T + \int_T^\infty \mathbb{P}(|A_i| \geq t). \quad (28)$$

By the above lemma, we have

$$\mathbb{E}_{Y,X|\mathbf{w}} \sum_{i=1}^N T_i \leq 2n\varepsilon\sigma \sqrt{\mathbb{E}_{Y,X|\mathbf{w}} \|M(D) - \mathbf{w}\|_{\Sigma_x}^2} + 2n\delta T + n \int_T^\infty \mathbb{P}(|T_i| \geq t). \quad (29)$$

For the last term, we have

$$\begin{aligned} \mathbb{P}(|T_i| \geq t) &= \mathbb{P}(|(y_i - \text{ReLU}(\mathbf{w}^\top \mathbf{x}_i))| |\langle M(D) - \mathbf{w}, x_i \mathbb{1}(w^\top x_i > 0) \rangle| > t) \\ &\leq \mathbb{P}(|(y_i - \text{ReLU}(\mathbf{w}^\top \mathbf{x}_i))| |\langle M(D) - \mathbf{w}, \mathbf{x}_i \rangle| > t) \\ &\leq \mathbb{P}(|(y_i - \text{ReLU}(\mathbf{w}^\top \mathbf{x}_i))| \sqrt{d} > t) \\ &\leq 2 \exp(-\frac{t^2}{2d\sigma^2}). \end{aligned}$$

Choosing $T = \sqrt{2}\sigma \sqrt{d \log(1/\delta)}$ we have

$$\begin{aligned} O(\sigma^2 d) &\leq \mathbb{E}_{Y,X|\mathbf{w}} \sum_{i=1}^N T_i \\ &\leq 2n\varepsilon\sigma \sqrt{\mathbb{E}_{Y,X|\mathbf{w}} \|M(D) - \mathbf{w}\|_{\Sigma_x}^2} + O(\sigma n \delta \sqrt{d \log(1/\delta)}). \end{aligned}$$

Thus we have the result when $\delta \leq N^{-(1+u)}$ for large enough u .

Next we will show that $\mathcal{L}(M(D)) - \mathcal{L}(\mathbf{w}_*) \geq \frac{1}{4} \|M(D) - \mathbf{w}_*\|_{\Sigma_x}^2$. Specifically, we will show for any \mathbf{w} , $\mathcal{L}(\mathbf{w}) - \mathcal{L}(\mathbf{w}_*) \geq \frac{1}{4} \|\mathbf{w} - \mathbf{w}_*\|_{\Sigma_x}^2$. Under the well-specified condition, we can easily see that

$$\begin{aligned} \mathcal{L}(\mathbf{w}) - \mathcal{L}(\mathbf{w}_*) &= \mathbb{E}[\text{ReLU}(\mathbf{x}^\top \mathbf{w}) - \text{ReLU}(\mathbf{x}^\top \mathbf{w}_*)]^2 \\ &= \mathbb{E}(\mathbf{x}^\top \mathbf{w} \cdot \mathbb{1}[\mathbf{x}^\top \mathbf{w} > 0] - \mathbf{x}^\top \mathbf{w}_* \cdot \mathbb{1}[\mathbf{x}^\top \mathbf{w}_* > 0])^2 \\ &= \mathbb{E}[\mathbf{w}^\top \mathbf{x} \mathbf{x}^\top \mathbf{w} \cdot \mathbb{1}[\mathbf{x}^\top \mathbf{w} > 0]] + \mathbb{E}[\mathbf{w}_*^\top \mathbf{x} \mathbf{x}^\top \mathbf{w}_* \cdot \mathbb{1}[\mathbf{x}^\top \mathbf{w}_* > 0]] \\ &\quad - 2\mathbb{E}[\mathbf{w}^\top \mathbf{x} \mathbf{x}^\top \mathbf{w}_* \cdot \mathbb{1}[\mathbf{x}^\top \mathbf{w} > 0, \mathbf{x}^\top \mathbf{w}_* > 0]]. \end{aligned}$$

According to Assumption 4, it further implied that

$$\mathbb{E}(\text{ReLU}(\mathbf{x}^\top \mathbf{w}) - \text{ReLU}(\mathbf{x}^\top \mathbf{w}_*))^2$$

$$\begin{aligned}
&= \mathbb{E}[\mathbf{w}^\top \mathbf{x} \mathbf{x}^\top \mathbf{w} \cdot \mathbb{1}[\mathbf{x}^\top \mathbf{w} < 0]] + \mathbb{E}[\mathbf{w}_*^\top \mathbf{x} \mathbf{x}^\top \mathbf{w}_* \cdot \mathbb{1}[\mathbf{x}^\top \mathbf{w}_* < 0]] \\
&- 2\mathbb{E}[\mathbf{w}^\top \mathbf{x} \mathbf{x}^\top \mathbf{w}_* \cdot \mathbb{1}[\mathbf{x}^\top \mathbf{w} < 0, \mathbf{x}^\top \mathbf{w}_* < 0]] \\
&= \mathbb{E}(\text{ReLU}(-\mathbf{x}^\top \mathbf{w}) - \text{ReLU}(-\mathbf{x}^\top \mathbf{w}_*))^2.
\end{aligned}$$

Moreover, we have:

$$\begin{aligned}
(\mathbf{x}^\top \mathbf{w} - \mathbf{x}^\top \mathbf{w}_*)^2 &= (\mathbf{x}^\top \mathbf{w} \mathbb{1}[\mathbf{x}^\top \mathbf{w} > 0] - \mathbf{x}^\top \mathbf{w}_* \mathbb{1}[\mathbf{x}^\top \mathbf{w}_* > 0] + \mathbf{x}^\top \mathbf{w} \mathbb{1}[\mathbf{x}^\top \mathbf{w} < 0] - \mathbf{x}^\top \mathbf{w}_* \mathbb{1}[\mathbf{x}^\top \mathbf{w}_* < 0])^2 \\
&\leq 2(\mathbf{x}^\top \mathbf{w} \mathbb{1}[\mathbf{x}^\top \mathbf{w} > 0] - \mathbf{x}^\top \mathbf{w}_* \mathbb{1}[\mathbf{x}^\top \mathbf{w}_* > 0])^2 + 2(\mathbf{x}^\top \mathbf{w} \mathbb{1}[\mathbf{x}^\top \mathbf{w} < 0] - \mathbf{x}^\top \mathbf{w}_* \mathbb{1}[\mathbf{x}^\top \mathbf{w}_* < 0])^2 \\
&= 2(\text{ReLU}(\mathbf{x}^\top \mathbf{w}) - \text{ReLU}(\mathbf{x}^\top \mathbf{w}_*))^2 + 2(\text{ReLU}(-\mathbf{x}^\top \mathbf{w}) - \text{ReLU}(-\mathbf{x}^\top \mathbf{w}_*))^2.
\end{aligned}$$

Then taking an expectation on both sides we obtain that

$$\begin{aligned}
\mathbb{E}(\mathbf{x}^\top \mathbf{w} - \mathbf{x}^\top \mathbf{w}_*)^2 &\leq 2\mathbb{E}(\text{ReLU}(\mathbf{x}^\top \mathbf{w}) - \text{ReLU}(\mathbf{x}^\top \mathbf{w}_*))^2 + 2\mathbb{E}(\text{ReLU}(-\mathbf{x}^\top \mathbf{w}) - \text{ReLU}(-\mathbf{x}^\top \mathbf{w}_*))^2 \\
&= 4\mathbb{E}(\text{ReLU}(\mathbf{x}^\top \mathbf{w}) - \text{ReLU}(\mathbf{x}^\top \mathbf{w}_*))^2.
\end{aligned}$$

The proof is completed. □